

```
In [2]: # Jessica Gallo & Asslidin Asliev
# CSC 732 Pattern Recognition & Neural Networks
# Final Project
# Part 1
# Image Classification with Convolutional Neural Networks
# Problem 1

# Created: 4/22/2020
# Last Modified: 4/27/2020
```

```
In [24]: # -----
# IMPORTS /
# -----

# A
import os, shutil # DOWNLOADING & SPLITTING DATA
# B
from keras import layers # INSTANTIATING A SMALL CONVNET
from keras import models # INSTANTIATING A SMALL CONVNET
from keras import optimizers # CONFIGURING THE MODEL FOR TRAINING
from keras.preprocessing.image import ImageDataGenerator # USING IMAGEDATAGENERATOR TO READ IMAGE
import matplotlib.pyplot as plt # DISPLAYING CURVES OF LOSS & ACCURACY DURING TRAINING
# C
from keras.preprocessing import image # DISPLAYING SOME RANDOMLY AUGMENTED TRAINING IMAGES
from keras.models import load_model # VISUALIZING INTERMEDIATE ACTIVATIONS
from keras.preprocessing import image # preprocessing a single image
import numpy as np # preprocessing a single image
import matplotlib.pyplot as plt # displaying the test picture
from keras import models # instantiating a model from an input tensor and a list of output tensors
import matplotlib.pyplot as plt # visualizing the fourth channel
from keras.applications import VGG16 # defining the loss tensor for filter visualization
from keras import backend as K # defining the loss tensor for filter visualization
import numpy as np # fetching numpy output values given numpy input values
```

```
In [4]: # =====
# A
# Download the original dataset and create a new one containing 3 subsets: |
# a training set with 1400 samples of each class, a validation set with |
# 400 samples of each class, and a test set with 200 samples for each class |
# (70%, 20%, 10%) |
# =====
```

```
In [16]: # -----
# DOWNLOADING & SPLITTING THE DATA /
# -----
```

```
# Path to original dataset
original_dataset_dir = './Downloads/dogs_vs_cats_original_dataset/train'

# Making a new smaller dataset
base_dir = './Downloads/dogs_vs_cats_small_dataset'
os.mkdir(base_dir)

# Making directories for training, validation & test
train_dir = os.path.join(base_dir, 'train')
os.mkdir(train_dir)
validation_dir = os.path.join(base_dir, 'validation')
os.mkdir(validation_dir)
test_dir = os.path.join(base_dir, 'test')
os.mkdir(test_dir)

# Directories with training cat and dog pictures
train_cats_dir = os.path.join(train_dir, 'cats')
os.mkdir(train_cats_dir)
train_dogs_dir = os.path.join(train_dir, 'dogs')
os.mkdir(train_dogs_dir)

# Directories with validation cat and dog pictures
validation_cats_dir = os.path.join(validation_dir, 'cats')
os.mkdir(validation_cats_dir)
validation_dogs_dir = os.path.join(validation_dir, 'dogs')
os.mkdir(validation_dogs_dir)

# Directories with test cat and dog pictures
test_cats_dir = os.path.join(test_dir, 'cats')
os.mkdir(test_cats_dir)
test_dogs_dir = os.path.join(test_dir, 'dogs')
os.mkdir(test_dogs_dir)

# Copies the first 1,000 cat images to train_cats_dir
fnames = ['cat.{}.jpg'.format(i) for i in range(1000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(train_cats_dir, fname)
    shutil.copyfile(src, dst)

# Copies the next 500 cat images to validation_cats_dir
fnames = ['cat.{}.jpg'.format(i) for i in range(1000, 1500)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(validation_cats_dir, fname)
    shutil.copyfile(src, dst)

# Copies the next 500 cat images to test_cats_dir
fnames = ['cat.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_cats_dir, fname)
    shutil.copyfile(src, dst)

# Copies the first 1,000 dog images to train_dogs_dir
fnames = ['dog.{}.jpg'.format(i) for i in range(1000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
```

```

dst = os.path.join(train_dogs_dir, fname)
shutil.copyfile(src, dst)

# Copies the next 500 dog images to validation_dogs_dir
fnames = ['dog.{}.jpg'.format(i) for i in range(1000, 1500)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(validation_dogs_dir, fname)
    shutil.copyfile(src, dst)

fnames = ['dog.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_dogs_dir, fname)
    shutil.copyfile(src, dst)

```

In [17]: # Check how many pictures are in each train, validation and test directories

```

print('Total training cat images:', len(os.listdir(train_cats_dir)))
print('Total training dog images:', len(os.listdir(train_dogs_dir)))
print('Total validation cat images:', len(os.listdir(validation_cats_dir)))
print('Total validation dog images:', len(os.listdir(validation_dogs_dir)))
print('Total test cat images:', len(os.listdir(test_cats_dir)))
print('Total test dog images:', len(os.listdir(test_dogs_dir)))

```

```

Total training cat images: 1000
Total training dog images: 1000
Total validation cat images: 500
Total validation dog images: 500
Total test cat images: 500
Total test dog images: 500

```

In [18]: # =====

```

# B
# Use Sequential model and build the topology of your CNN implementing |
# augmentation by including 4 convolutional layers Conv2D (with Relu |
# activation), 4 MaxPooling2D layers, a Dense layer of size 1 and a sigmoid |
# activation as the number of classes is 2. Use Adam optimizer. |
# =====

```

In [20]:

```
# -----
# INSTANTIATING A SMALL CONVNET /
# -----
```

model = models.Sequential()
tells keras to stack all layers sequentially
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
filter size[32]: the output dimension (i.e. number of output filters in convolution)
kernel_size[3,3]: specified height and weight of 2D convolution window
activation['relu']: activation function also called non-linearity to be out neural network
input_shape[150,,150,3]: resized image dimensions 150x150 with 3 channel
model.add(layers.MaxPooling2D((2, 2)))
MaxPooling: reduces spacial size of incoming features to help reduce amount
of parameters & computation in network, helping to reduce overfitting
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid')) *# Sigmoid function at the end because it's just*
Want the model to output a probability of how sure an image is a dog and not a cat, which means
we want a probability score where higher values means the classifier believes the image is a dog
Lower values means it is a cat. The sigmoid is perfect for this because it takes in a set of n
and returns a probability distribution in the range of 0 to 1

```
model.summary()
# previews the arrangement and parameter size of our convnet
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
<hr/>		
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
<hr/>		
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
<hr/>		
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
<hr/>		
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
<hr/>		
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
<hr/>		
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
<hr/>		
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
<hr/>		
flatten_1 (Flatten)	(None, 6272)	0
<hr/>		
dense_1 (Dense)	(None, 512)	3211776
<hr/>		
dense_2 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

```
In [22]: # -----
# CONFIGURING THE MODEL FOR TRAINING /
# -----
```

```
model.compile(loss='binary_crossentropy', optimizer=optimizers.Adam(lr=1e-4), metrics=['acc'])
# Adam optimizer with learning rate of 0.0001
# Loss['binary_crossentropy']: because its a binary classification & to specify a loss function
# Accuracy metric (acc) is a good metric to use to know if our model is doing well or not
```

```
In [25]: # -----
# USING THE IMAGEDATAGENERATOR TO READ IMAGES FROM DIRECTORIES /
# -----
```

```
# rescale images by 1/255
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1.255)

train_generator = train_datagen.flow_from_directory(train_dir, # target directory
                                                    target_size=(150, 150),
                                                    batch_size=20,
                                                    class_mode='binary')
# rotation, width_shift, height_shift, shear, zoom and horizontal_flip all are image augmentation
# tells ImageDataGen to randomly apply some transformation to image to help augment dataset & i

# because of binary_crossentropy Loss, you need binary labels

validation_generator = test_datagen.flow_from_directory(validation_dir,
                                                       target_size=(150, 150),
                                                       batch_size=20,
                                                       class_mode='binary')
```

Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

```
In [26]: # Samples from each batch
for data_batch, labels_batch in train_generator:
    print('Data batch shape:', data_batch.shape)
    print('Labels batch shape:', labels_batch.shape)
    break
```

Data batch shape: (20, 150, 150, 3)

Labels batch shape: (20,)

In [28]:

```
# -----  
# FITTING THE MODEL USING A BATCH GENERATOR |  
# -----
```

```
history = model.fit_generator(train_generator,  
                             steps_per_epoch=100,  
                             epochs=30,  
                             validation_data=validation_generator,  
                             validation_steps=50)
```

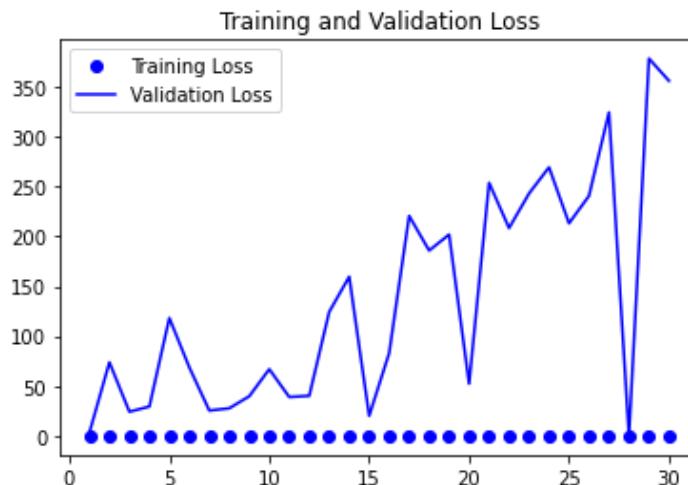
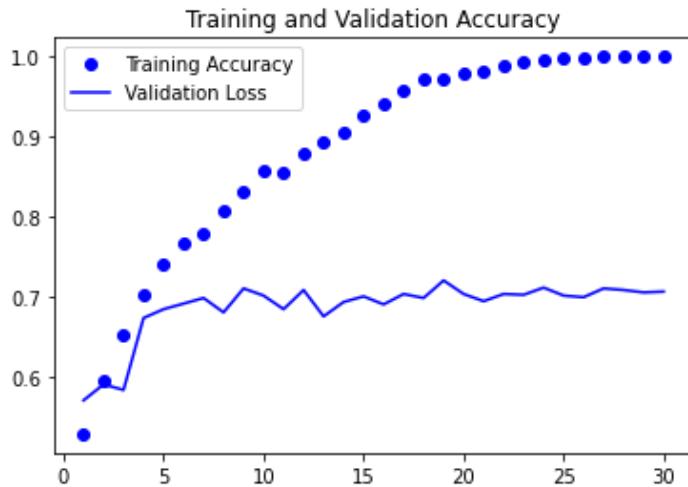
```
Epoch 1/30  
100/100 [=====] - 146s 1s/step - loss: 0.6931 - acc: 0.5280 - val_loss:  
s: 3.3309 - val_acc: 0.5700  
Epoch 2/30  
100/100 [=====] - 155s 2s/step - loss: 0.6700 - acc: 0.5945 - val_loss:  
s: 73.7493 - val_acc: 0.5900  
Epoch 3/30  
100/100 [=====] - 137s 1s/step - loss: 0.6312 - acc: 0.6510 - val_loss:  
s: 24.2847 - val_acc: 0.5830  
Epoch 4/30  
100/100 [=====] - 131s 1s/step - loss: 0.5686 - acc: 0.7030 - val_loss:  
s: 29.4228 - val_acc: 0.6730  
Epoch 5/30  
100/100 [=====] - 123s 1s/step - loss: 0.5168 - acc: 0.7405 - val_loss:  
s: 118.1477 - val_acc: 0.6840  
Epoch 6/30  
100/100 [=====] - 122s 1s/step - loss: 0.4926 - acc: 0.7675 - val_loss:  
s: 68.6237 - val_acc: 0.6910  
Epoch 7/30  
100/100 [=====] - 131s 1s/step - loss: 0.4549 - acc: 0.7790 - val_loss:  
s: 25.4330 - val_acc: 0.6980  
Epoch 8/30  
100/100 [=====] - 127s 1s/step - loss: 0.4292 - acc: 0.8075 - val_loss:  
s: 27.7126 - val_acc: 0.6800  
Epoch 9/30  
100/100 [=====] - 134s 1s/step - loss: 0.3883 - acc: 0.8320 - val_loss:  
s: 39.7610 - val_acc: 0.7100  
Epoch 10/30  
100/100 [=====] - 135s 1s/step - loss: 0.3466 - acc: 0.8565 - val_loss:  
s: 66.9264 - val_acc: 0.7010  
Epoch 11/30  
100/100 [=====] - 135s 1s/step - loss: 0.3371 - acc: 0.8555 - val_loss:  
s: 38.9676 - val_acc: 0.6840  
Epoch 12/30  
100/100 [=====] - 132s 1s/step - loss: 0.2917 - acc: 0.8785 - val_loss:  
s: 40.1939 - val_acc: 0.7080  
Epoch 13/30  
100/100 [=====] - 131s 1s/step - loss: 0.2665 - acc: 0.8930 - val_loss:  
s: 124.3160 - val_acc: 0.6750  
Epoch 14/30  
100/100 [=====] - 134s 1s/step - loss: 0.2409 - acc: 0.9040 - val_loss:  
s: 159.5469 - val_acc: 0.6930  
Epoch 15/30  
100/100 [=====] - 144s 1s/step - loss: 0.1995 - acc: 0.9270 - val_loss:  
s: 20.1422 - val_acc: 0.7000  
Epoch 16/30  
100/100 [=====] - 136s 1s/step - loss: 0.1764 - acc: 0.9410 - val_loss:  
s: 83.0314 - val_acc: 0.6900  
Epoch 17/30  
100/100 [=====] - 127s 1s/step - loss: 0.1416 - acc: 0.9580 - val_loss:  
s: 220.4671 - val_acc: 0.7030  
Epoch 18/30
```

```
100/100 [=====] - 122s 1s/step - loss: 0.1094 - acc: 0.9720 - val_loss:  
s: 185.7634 - val_acc: 0.6980  
Epoch 19/30  
100/100 [=====] - 127s 1s/step - loss: 0.1047 - acc: 0.9715 - val_loss:  
s: 201.8043 - val_acc: 0.7200  
Epoch 20/30  
100/100 [=====] - 111s 1s/step - loss: 0.0834 - acc: 0.9790 - val_loss:  
s: 52.2085 - val_acc: 0.7030  
Epoch 21/30  
100/100 [=====] - 105s 1s/step - loss: 0.0710 - acc: 0.9825 - val_loss:  
s: 253.6445 - val_acc: 0.6940  
Epoch 22/30  
100/100 [=====] - 103s 1s/step - loss: 0.0585 - acc: 0.9885 - val_loss:  
s: 208.2495 - val_acc: 0.7030  
Epoch 23/30  
100/100 [=====] - 106s 1s/step - loss: 0.0393 - acc: 0.9935 - val_loss:  
s: 242.9491 - val_acc: 0.7020  
Epoch 24/30  
100/100 [=====] - 104s 1s/step - loss: 0.0296 - acc: 0.9965 - val_loss:  
s: 269.1779 - val_acc: 0.7110  
Epoch 25/30  
100/100 [=====] - 103s 1s/step - loss: 0.0217 - acc: 0.9990 - val_loss:  
s: 213.0229 - val_acc: 0.7010  
Epoch 26/30  
100/100 [=====] - 103s 1s/step - loss: 0.0164 - acc: 0.9990 - val_loss:  
s: 240.5958 - val_acc: 0.6990  
Epoch 27/30  
100/100 [=====] - 105s 1s/step - loss: 0.0111 - acc: 1.0000 - val_loss:  
s: 324.1751 - val_acc: 0.7100  
Epoch 28/30  
100/100 [=====] - 104s 1s/step - loss: 0.0093 - acc: 1.0000 - val_loss:  
s: 0.0000e+00 - val_acc: 0.7080  
Epoch 29/30  
100/100 [=====] - 102s 1s/step - loss: 0.0063 - acc: 1.0000 - val_loss:  
s: 378.1235 - val_acc: 0.7050  
Epoch 30/30  
100/100 [=====] - 103s 1s/step - loss: 0.0053 - acc: 1.0000 - val_loss:  
s: 355.8469 - val_acc: 0.7060
```

```
In [29]: # -----  
# SAVE THE MODEL /  
# -----  
model.save('cats_and_dogs_small_1.h5')
```

In [30]:

```
# -----  
# DISPLAYING CURVES OF LOSS & ACCURACY DURING TRAINING |  
# -----  
  
acc = history.history['acc']  
val_acc = history.history['val_acc']  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs = range(1, len(acc) + 1)  
  
plt.plot(epochs, acc, 'bo', label='Training Accuracy')  
plt.plot(epochs, val_acc, 'b', label='Validation Loss')  
plt.title('Training and Validation Accuracy')  
plt.legend()  
  
plt.figure()  
  
plt.plot(epochs, loss, 'bo', label='Training Loss')  
plt.plot(epochs, val_loss, 'b', label='Validation Loss')  
plt.title('Training and Validation Loss')  
plt.legend()  
  
plt.show()  
  
# Plots show overfitting
```



```
In [31]: # ======  
# C  
# Display randomly augmented images(four images of two different cats and |  
# four images of two different dogs) |  
# ======
```

```
In [34]: # -----  
# SETTING UP A DATA AUGMENTATION CONFIGURATION VIA IMAGEDATAGENERATOR |  
# -----  
  
datagen = ImageDataGenerator(rotation_range=40,  
                             width_shift_range=0.2,  
                             height_shift_range=0.2,  
                             shear_range=0.2,  
                             zoom_range=0.2,  
                             horizontal_flip=True,  
                             fill_mode='nearest')
```

In [70]:

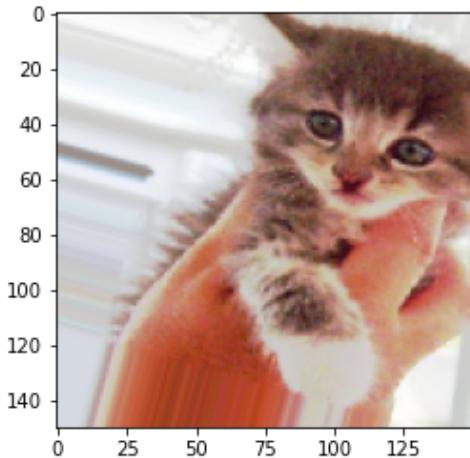
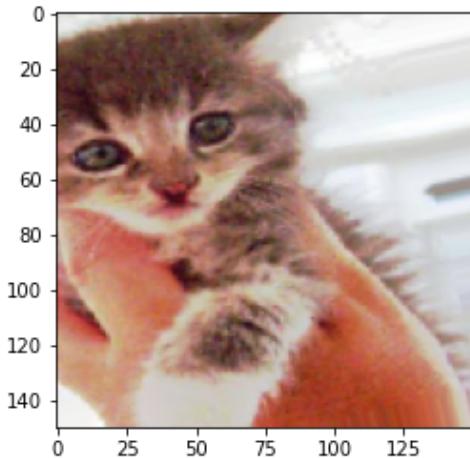
```
# -----
# DISPLAYING SOME RANDOMLY AUGMENTED TRAINING IMAGES |
# -----  
  
# 1st CAT  
# ~~~~~~  
fnames = [os.path.join(train_cats_dir, fname) for fname in os.listdir(train_cats_dir)]  
img_path = fnames[3] # chooses one image augment  
img = image.load_img(img_path, target_size=(150, 150)) # reads the image & resizes it  
  
x = image.img_to_array(img) # converts it to a numpy array with shape(150, 150, 3)  
  
x = x.reshape((1,) + x.shape) # reshapes it to (1, 150, 150, 3)  
  
# generates batches of randomly transformed images  
# loops indefinitely, so you need to break the loop at some point]  
i = 0  
for batch in datagen.flow(x, batch_size=1):  
    plt.figure(i)  
    imgplot = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 2 == 0:  
        break  
  
plt.show()  
  
# 2nd CAT  
# ~~~~~~  
  
img_path2 = fnames[2]  
img2 = image.load_img(img_path2, target_size=(150, 150))  
x2 = image.img_to_array(img2)  
x2 = x2.reshape((1,) + x2.shape)  
  
i = 0  
for batch in datagen.flow(x2, batch_size=1):  
    plt.figure(i)  
    imgplot2 = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 2 == 0:  
        break  
  
plt.show()  
  
# 1st DOG  
# ~~~~~~  
  
fnames2 = [os.path.join(train_dogs_dir, fname) for fname in os.listdir(train_dogs_dir)]  
img_path3 = fnames2[3]  
img3 = image.load_img(img_path3, target_size=(150, 150))  
x3 = image.img_to_array(img3)  
x3 = x3.reshape((1,) + x3.shape)  
  
i = 0  
for batch in datagen.flow(x3, batch_size=1):  
    plt.figure(i)  
    imgplot3 = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 2 == 0:  
        break  
  
plt.show()
```

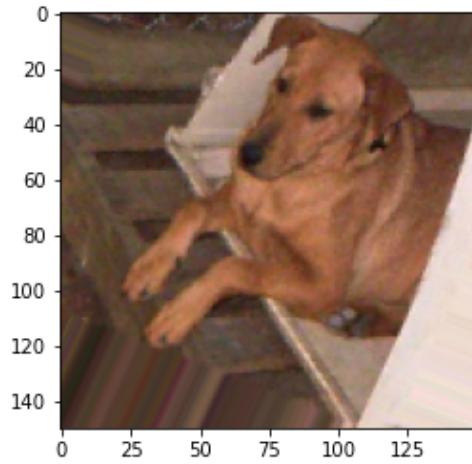
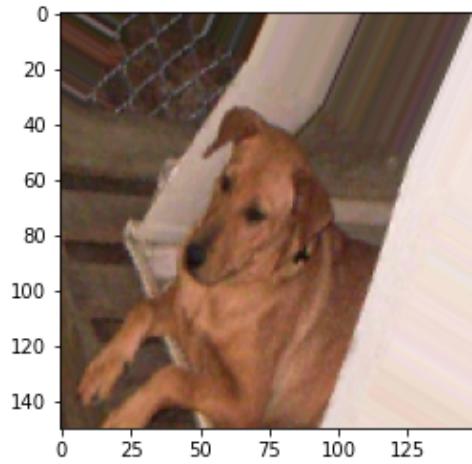
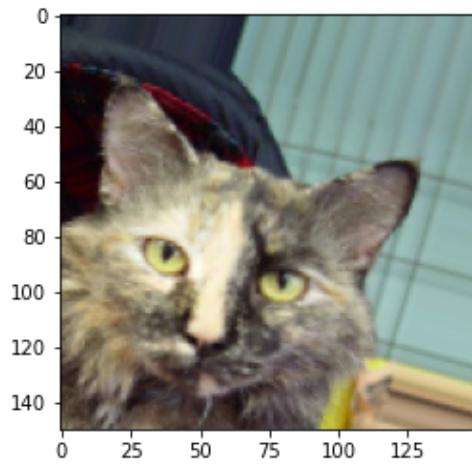
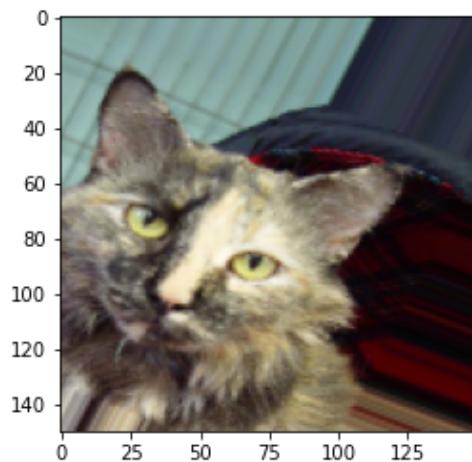
```
# 2nd DOG
# ~~~~~

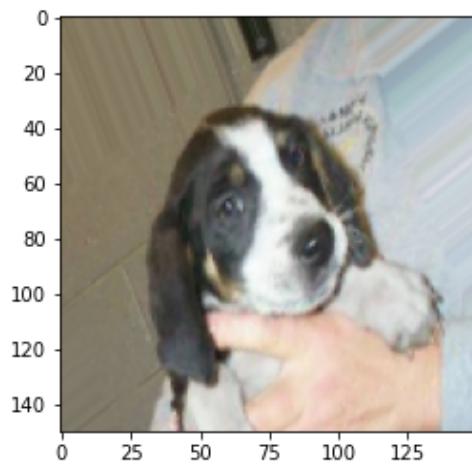
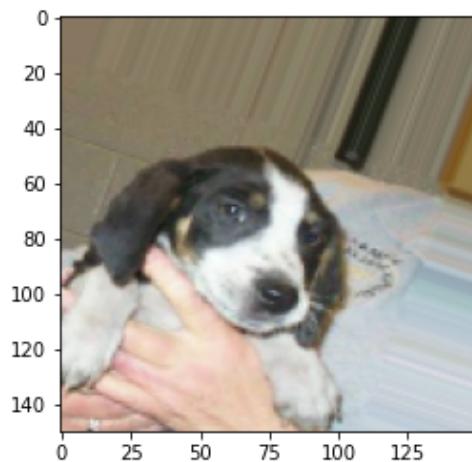
img_path4 = fnames2[2]
img4 = image.load_img(img_path4, target_size=(150, 150))
x4 = image.img_to_array(img4)
x4 = x4.reshape((1,) + x4.shape)

i = 0
for batch in datagen.flow(x4, batch_size=1):
    plt.figure(i)
    imgplot4 = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 2 == 0:
        break

plt.show()
```







```
In [71]: # ======  
# D  
# To avoid overfitting add Dropout Layer and do the training of your CNN  
# using data augmentation generators. Plot training and validation accuracy  
# as well as training and validation loss  
# ======
```

In [73]:

```
# -----
# DEFINING A NEW CONVNET THAT INCLUDES DROPOUT |
# -----  
  
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(128, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(128, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Flatten())  
model.add(layers.Dropout(0.5)) # Dropout for regularization  
# Dropout randomly drops some layers in a neural networks and then learns with the reduced network  
# The network learns to be independent and not reliable on a single layer  
# Helps with overfitting  
# .05 means to randomly drop half of the layers  
model.add(layers.Dense(512, activation='relu'))  
model.add(layers.Dense(1, activation='sigmoid'))  
  
model.compile(loss='binary_crossentropy', optimizer=optimizers.Adam(lr=1e-4), metrics=['acc'])  
  
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_9 (Conv2D)	(None, 148, 148, 32)	896
<hr/>		
max_pooling2d_9 (MaxPooling2D)	(None, 74, 74, 32)	0
<hr/>		
conv2d_10 (Conv2D)	(None, 72, 72, 64)	18496
<hr/>		
max_pooling2d_10 (MaxPooling2D)	(None, 36, 36, 64)	0
<hr/>		
conv2d_11 (Conv2D)	(None, 34, 34, 128)	73856
<hr/>		
max_pooling2d_11 (MaxPooling2D)	(None, 17, 17, 128)	0
<hr/>		
conv2d_12 (Conv2D)	(None, 15, 15, 128)	147584
<hr/>		
max_pooling2d_12 (MaxPooling2D)	(None, 7, 7, 128)	0
<hr/>		
flatten_3 (Flatten)	(None, 6272)	0
<hr/>		
dropout_2 (Dropout)	(None, 6272)	0
<hr/>		
dense_5 (Dense)	(None, 512)	3211776
<hr/>		
dense_6 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

In [77]:

```
# -----  
# TRAINING THE CONVNET USING DATA-AUGMENTATION GENERATORS |  
# -----  
  
train_datagen = ImageDataGenerator(rescale=1./255,  
                                    rotation_range=40,  
                                    width_shift_range=0.2,  
                                    height_shift_range=0.2,  
                                    shear_range=0.2,  
                                    zoom_range=0.2,  
                                    horizontal_flip=True)  
  
test_datagen = ImageDataGenerator(rescale=1./255) # not augmented  
  
train_generator = train_datagen.flow_from_directory(train_dir, # target directory  
                                                    target_size=(150, 150), # resizes all images  
                                                    batch_size=32,  
                                                    class_mode='binary')  
  
validation_generator = test_datagen.flow_from_directory(validation_dir,  
                                                       target_size=(150, 150),  
                                                       batch_size=32,  
                                                       class_mode='binary')  
  
history = model.fit_generator(train_generator,  
                             steps_per_epoch=100,  
                             epochs=100,  
                             validation_data=validation_generator,  
                             validation_steps=50)
```

```
Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.  
Epoch 1/100  
100/100 [=====] - 208s 2s/step - loss: 0.6914 - acc: 0.5303 - val_l  
oss: 12.1340 - val_acc: 0.5900  
Epoch 2/100  
100/100 [=====] - 214s 2s/step - loss: 0.6754 - acc: 0.5663 - val_l  
oss: 72.1128 - val_acc: 0.5490  
Epoch 3/100  
100/100 [=====] - 204s 2s/step - loss: 0.6642 - acc: 0.5891 - val_l  
oss: 92.5164 - val_acc: 0.5020  
Epoch 4/100  
100/100 [=====] - 236s 2s/step - loss: 0.6479 - acc: 0.6136 - val_l  
oss: 60.8748 - val_acc: 0.5930  
Epoch 5/100  
100/100 [=====] - 251s 3s/step - loss: 0.6263 - acc: 0.6487 - val_l  
oss: 208.2860 - val_acc: 0.5190  
Epoch 6/100  
100/100 [=====] - 232s 2s/step - loss: 0.6142 - acc: 0.6616 - val_l  
oss: 12.0749 - val_acc: 0.5910  
Epoch 7/100  
100/100 [=====] - 232s 2s/step - loss: 0.6074 - acc: 0.6731 - val_l  
oss: 184.3235 - val_acc: 0.5330  
Epoch 8/100  
100/100 [=====] - 234s 2s/step - loss: 0.5929 - acc: 0.6796 - val_l  
oss: 116.1794 - val_acc: 0.5250  
Epoch 9/100  
100/100 [=====] - 241s 2s/step - loss: 0.5839 - acc: 0.6903 - val_l  
oss: 180.8572 - val_acc: 0.5610  
Epoch 10/100  
100/100 [=====] - 244s 2s/step - loss: 0.5718 - acc: 0.6988 - val_l  
oss: 204.1322 - val_acc: 0.5720
```

```
Epoch 11/100
100/100 [=====] - 239s 2s/step - loss: 0.5576 - acc: 0.7139 - val_l
oss: 101.9894 - val_acc: 0.5770
Epoch 12/100
100/100 [=====] - 221s 2s/step - loss: 0.5536 - acc: 0.7132 - val_l
oss: 236.0468 - val_acc: 0.5870
Epoch 13/100
100/100 [=====] - 218s 2s/step - loss: 0.5402 - acc: 0.7236 - val_l
oss: 104.5838 - val_acc: 0.6140
Epoch 14/100
100/100 [=====] - 205s 2s/step - loss: 0.5355 - acc: 0.7290 - val_l
oss: 115.1791 - val_acc: 0.6170
Epoch 15/100
100/100 [=====] - 192s 2s/step - loss: 0.5255 - acc: 0.7367 - val_l
oss: 415.0259 - val_acc: 0.5460
Epoch 16/100
100/100 [=====] - 213s 2s/step - loss: 0.5311 - acc: 0.7311 - val_l
oss: 190.7686 - val_acc: 0.6040
Epoch 17/100
100/100 [=====] - 191s 2s/step - loss: 0.5214 - acc: 0.7522 - val_l
oss: 152.4332 - val_acc: 0.6310
Epoch 18/100
100/100 [=====] - 183s 2s/step - loss: 0.5121 - acc: 0.7538 - val_l
oss: 119.8108 - val_acc: 0.5900
Epoch 19/100
100/100 [=====] - 180s 2s/step - loss: 0.5076 - acc: 0.7472 - val_l
oss: 205.3783 - val_acc: 0.5870
Epoch 20/100
100/100 [=====] - 184s 2s/step - loss: 0.5061 - acc: 0.7516 - val_l
oss: 94.4848 - val_acc: 0.6060
Epoch 21/100
100/100 [=====] - 182s 2s/step - loss: 0.5134 - acc: 0.7412 - val_l
oss: 162.0881 - val_acc: 0.6170
Epoch 22/100
100/100 [=====] - 181s 2s/step - loss: 0.4962 - acc: 0.7595 - val_l
oss: 139.1903 - val_acc: 0.6270
Epoch 23/100
100/100 [=====] - 179s 2s/step - loss: 0.4880 - acc: 0.7579 - val_l
oss: 215.6177 - val_acc: 0.6290
Epoch 24/100
100/100 [=====] - 148s 1s/step - loss: 0.4962 - acc: 0.7569 - val_l
oss: 161.3710 - val_acc: 0.6530
Epoch 25/100
100/100 [=====] - 149s 1s/step - loss: 0.4821 - acc: 0.7711 - val_l
oss: 165.4077 - val_acc: 0.6550
Epoch 26/100
100/100 [=====] - 151s 2s/step - loss: 0.4832 - acc: 0.7707 - val_l
oss: 171.4778 - val_acc: 0.6110
Epoch 27/100
100/100 [=====] - 148s 1s/step - loss: 0.4808 - acc: 0.7667 - val_l
oss: 202.5868 - val_acc: 0.6310
Epoch 28/100
100/100 [=====] - 148s 1s/step - loss: 0.4744 - acc: 0.7660 - val_l
oss: 68.9730 - val_acc: 0.6430
Epoch 29/100
100/100 [=====] - 147s 1s/step - loss: 0.4675 - acc: 0.7819 - val_l
oss: 191.7742 - val_acc: 0.6520
Epoch 30/100
100/100 [=====] - 148s 1s/step - loss: 0.4564 - acc: 0.7827 - val_l
oss: 92.1834 - val_acc: 0.6570
Epoch 31/100
100/100 [=====] - 147s 1s/step - loss: 0.4661 - acc: 0.7721 - val_l
oss: 59.9744 - val_acc: 0.6180
```

```
Epoch 32/100
100/100 [=====] - 147s 1s/step - loss: 0.4567 - acc: 0.7838 - val_l
oss: 105.1589 - val_acc: 0.6540
Epoch 33/100
100/100 [=====] - 147s 1s/step - loss: 0.4429 - acc: 0.7852 - val_l
oss: 172.1588 - val_acc: 0.6900
Epoch 34/100
100/100 [=====] - 146s 1s/step - loss: 0.4440 - acc: 0.7904 - val_l
oss: 63.6146 - val_acc: 0.6850
Epoch 35/100
100/100 [=====] - 149s 1s/step - loss: 0.4425 - acc: 0.7942 - val_l
oss: 90.6968 - val_acc: 0.6540
Epoch 36/100
100/100 [=====] - 149s 1s/step - loss: 0.4572 - acc: 0.7830 - val_l
oss: 58.4663 - val_acc: 0.6690
Epoch 37/100
100/100 [=====] - 149s 1s/step - loss: 0.4398 - acc: 0.7949 - val_l
oss: 137.4117 - val_acc: 0.6900
Epoch 38/100
100/100 [=====] - 147s 1s/step - loss: 0.4249 - acc: 0.8097 - val_l
oss: 153.2112 - val_acc: 0.6280
Epoch 39/100
100/100 [=====] - 147s 1s/step - loss: 0.4328 - acc: 0.7999 - val_l
oss: 168.0579 - val_acc: 0.6780
Epoch 40/100
100/100 [=====] - 147s 1s/step - loss: 0.4167 - acc: 0.8131 - val_l
oss: 18.9001 - val_acc: 0.6980
Epoch 41/100
100/100 [=====] - 153s 2s/step - loss: 0.4301 - acc: 0.7993 - val_l
oss: 17.3868 - val_acc: 0.6720
Epoch 42/100
100/100 [=====] - 148s 1s/step - loss: 0.4419 - acc: 0.7989 - val_l
oss: 76.4298 - val_acc: 0.7230
Epoch 43/100
100/100 [=====] - 147s 1s/step - loss: 0.4135 - acc: 0.8094 - val_l
oss: 158.0459 - val_acc: 0.6900
Epoch 44/100
100/100 [=====] - 146s 1s/step - loss: 0.4071 - acc: 0.8093 - val_l
oss: 122.2766 - val_acc: 0.6940
Epoch 45/100
100/100 [=====] - 146s 1s/step - loss: 0.4244 - acc: 0.8056 - val_l
oss: 130.7091 - val_acc: 0.6670
Epoch 46/100
100/100 [=====] - 148s 1s/step - loss: 0.4046 - acc: 0.8112 - val_l
oss: 205.6271 - val_acc: 0.6480
Epoch 47/100
100/100 [=====] - 146s 1s/step - loss: 0.4079 - acc: 0.8109 - val_l
oss: 181.8083 - val_acc: 0.6620
Epoch 48/100
100/100 [=====] - 146s 1s/step - loss: 0.4030 - acc: 0.8191 - val_l
oss: 284.3588 - val_acc: 0.6620
Epoch 49/100
100/100 [=====] - 149s 1s/step - loss: 0.4030 - acc: 0.8147 - val_l
oss: 258.6402 - val_acc: 0.6960
Epoch 50/100
100/100 [=====] - 164s 2s/step - loss: 0.3971 - acc: 0.8201 - val_l
oss: 115.2727 - val_acc: 0.6870
Epoch 51/100
100/100 [=====] - 157s 2s/step - loss: 0.3999 - acc: 0.8223 - val_l
oss: 84.5245 - val_acc: 0.6820
Epoch 52/100
100/100 [=====] - 155s 2s/step - loss: 0.4048 - acc: 0.8150 - val_l
oss: 114.3616 - val_acc: 0.6610
```

```
Epoch 53/100
100/100 [=====] - 175s 2s/step - loss: 0.3971 - acc: 0.8179 - val_l
oss: 104.3057 - val_acc: 0.6850
Epoch 54/100
100/100 [=====] - 175s 2s/step - loss: 0.3869 - acc: 0.8239 - val_l
oss: 127.6097 - val_acc: 0.7250
Epoch 55/100
100/100 [=====] - 174s 2s/step - loss: 0.3823 - acc: 0.8229 - val_l
oss: 210.3990 - val_acc: 0.6840
Epoch 56/100
100/100 [=====] - 175s 2s/step - loss: 0.3905 - acc: 0.8182 - val_l
oss: 215.3125 - val_acc: 0.6510
Epoch 57/100
100/100 [=====] - 178s 2s/step - loss: 0.3681 - acc: 0.8364 - val_l
oss: 71.9983 - val_acc: 0.7350
Epoch 58/100
100/100 [=====] - 175s 2s/step - loss: 0.3961 - acc: 0.8188 - val_l
oss: 133.8043 - val_acc: 0.6890
Epoch 59/100
100/100 [=====] - 175s 2s/step - loss: 0.3826 - acc: 0.8273 - val_l
oss: 96.9491 - val_acc: 0.6490
Epoch 60/100
100/100 [=====] - 175s 2s/step - loss: 0.3806 - acc: 0.8268 - val_l
oss: 152.7578 - val_acc: 0.6920
Epoch 61/100
100/100 [=====] - 175s 2s/step - loss: 0.3715 - acc: 0.8354 - val_l
oss: 238.3345 - val_acc: 0.6600
Epoch 62/100
100/100 [=====] - 192s 2s/step - loss: 0.3533 - acc: 0.8430 - val_l
oss: 96.2947 - val_acc: 0.7270
Epoch 63/100
100/100 [=====] - 177s 2s/step - loss: 0.3738 - acc: 0.8261 - val_l
oss: 72.9868 - val_acc: 0.6750
Epoch 64/100
100/100 [=====] - 174s 2s/step - loss: 0.3612 - acc: 0.8368 - val_l
oss: 141.8358 - val_acc: 0.6920
Epoch 65/100
100/100 [=====] - 180s 2s/step - loss: 0.3730 - acc: 0.8279 - val_l
oss: 47.0057 - val_acc: 0.7350
Epoch 66/100
100/100 [=====] - 177s 2s/step - loss: 0.3601 - acc: 0.8400 - val_l
oss: 81.7871 - val_acc: 0.7010
Epoch 67/100
100/100 [=====] - 175s 2s/step - loss: 0.3669 - acc: 0.8352 - val_l
oss: 155.5197 - val_acc: 0.6780
Epoch 68/100
100/100 [=====] - 175s 2s/step - loss: 0.3540 - acc: 0.8401 - val_l
oss: 48.5071 - val_acc: 0.7420
Epoch 69/100
100/100 [=====] - 176s 2s/step - loss: 0.3582 - acc: 0.8346 - val_l
oss: 96.7944 - val_acc: 0.7160
Epoch 70/100
100/100 [=====] - 178s 2s/step - loss: 0.3483 - acc: 0.8455 - val_l
oss: 256.0091 - val_acc: 0.7060
Epoch 71/100
100/100 [=====] - 190s 2s/step - loss: 0.3460 - acc: 0.8477 - val_l
oss: 159.6587 - val_acc: 0.6340
Epoch 72/100
100/100 [=====] - 185s 2s/step - loss: 0.3522 - acc: 0.8447 - val_l
oss: 50.8508 - val_acc: 0.7090
Epoch 73/100
100/100 [=====] - 162s 2s/step - loss: 0.3487 - acc: 0.8438 - val_l
oss: 188.2545 - val_acc: 0.7340
```

```
Epoch 74/100
100/100 [=====] - 146s 1s/step - loss: 0.3312 - acc: 0.8592 - val_l
oss: 15.4837 - val_acc: 0.7570
Epoch 75/100
100/100 [=====] - 147s 1s/step - loss: 0.3485 - acc: 0.8461 - val_l
oss: 34.5536 - val_acc: 0.7470
Epoch 76/100
100/100 [=====] - 147s 1s/step - loss: 0.3557 - acc: 0.8390 - val_l
oss: 108.9379 - val_acc: 0.7070
Epoch 77/100
100/100 [=====] - 147s 1s/step - loss: 0.3244 - acc: 0.8589 - val_l
oss: 133.2246 - val_acc: 0.6970
Epoch 78/100
100/100 [=====] - 147s 1s/step - loss: 0.3203 - acc: 0.8634 - val_l
oss: 145.4416 - val_acc: 0.7060
Epoch 79/100
100/100 [=====] - 147s 1s/step - loss: 0.3144 - acc: 0.8684 - val_l
oss: 122.8492 - val_acc: 0.7390
Epoch 80/100
100/100 [=====] - 146s 1s/step - loss: 0.3188 - acc: 0.8640 - val_l
oss: 215.0446 - val_acc: 0.6970
Epoch 81/100
100/100 [=====] - 148s 1s/step - loss: 0.3054 - acc: 0.8668 - val_l
oss: 71.3096 - val_acc: 0.7350
Epoch 82/100
100/100 [=====] - 147s 1s/step - loss: 0.3253 - acc: 0.8599 - val_l
oss: 129.5654 - val_acc: 0.7050
Epoch 83/100
100/100 [=====] - 146s 1s/step - loss: 0.2984 - acc: 0.8674 - val_l
oss: 324.9236 - val_acc: 0.7040
Epoch 84/100
100/100 [=====] - 147s 1s/step - loss: 0.3062 - acc: 0.8700 - val_l
oss: 112.5832 - val_acc: 0.7430
Epoch 85/100
100/100 [=====] - 146s 1s/step - loss: 0.3175 - acc: 0.8665 - val_l
oss: 93.6838 - val_acc: 0.7030
Epoch 86/100
100/100 [=====] - 147s 1s/step - loss: 0.3088 - acc: 0.8662 - val_l
oss: 50.1084 - val_acc: 0.7160
Epoch 87/100
100/100 [=====] - 146s 1s/step - loss: 0.2973 - acc: 0.8722 - val_l
oss: 52.8702 - val_acc: 0.7280
Epoch 88/100
100/100 [=====] - 148s 1s/step - loss: 0.3028 - acc: 0.8678 - val_l
oss: 219.2124 - val_acc: 0.7100
Epoch 89/100
100/100 [=====] - 147s 1s/step - loss: 0.3018 - acc: 0.8668 - val_l
oss: 95.1776 - val_acc: 0.7310
Epoch 90/100
100/100 [=====] - 146s 1s/step - loss: 0.3020 - acc: 0.8643 - val_l
oss: 32.8336 - val_acc: 0.7340
Epoch 91/100
100/100 [=====] - 145s 1s/step - loss: 0.2982 - acc: 0.8722 - val_l
oss: 121.5771 - val_acc: 0.6840
Epoch 92/100
100/100 [=====] - 146s 1s/step - loss: 0.2824 - acc: 0.8769 - val_l
oss: 180.8146 - val_acc: 0.6500
Epoch 93/100
100/100 [=====] - 146s 1s/step - loss: 0.2897 - acc: 0.8763 - val_l
oss: 192.2625 - val_acc: 0.6440
Epoch 94/100
100/100 [=====] - 147s 1s/step - loss: 0.2986 - acc: 0.8684 - val_l
oss: 16.5741 - val_acc: 0.7430
```

```
Epoch 95/100
100/100 [=====] - 147s 1s/step - loss: 0.2857 - acc: 0.8794 - val_loss: 85.7954 - val_acc: 0.7370
Epoch 96/100
100/100 [=====] - 146s 1s/step - loss: 0.2799 - acc: 0.8816 - val_loss: 79.5412 - val_acc: 0.7260
Epoch 97/100
100/100 [=====] - 146s 1s/step - loss: 0.2907 - acc: 0.8722 - val_loss: 94.3148 - val_acc: 0.6740
Epoch 98/100
100/100 [=====] - 146s 1s/step - loss: 0.2875 - acc: 0.8797 - val_loss: 104.6750 - val_acc: 0.6980
Epoch 99/100
100/100 [=====] - 146s 1s/step - loss: 0.2678 - acc: 0.8844 - val_loss: 82.3507 - val_acc: 0.7180
Epoch 100/100
100/100 [=====] - 146s 1s/step - loss: 0.2760 - acc: 0.8781 - val_loss: 108.8086 - val_acc: 0.7270
```

```
In [78]: model.save('cats_and_dogs_small_2.h5')
```

```
In [79]: # -----
# DISPLAYING CURVES OF LOSS & ACCURACY DURING TRAINING |
# -----

acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

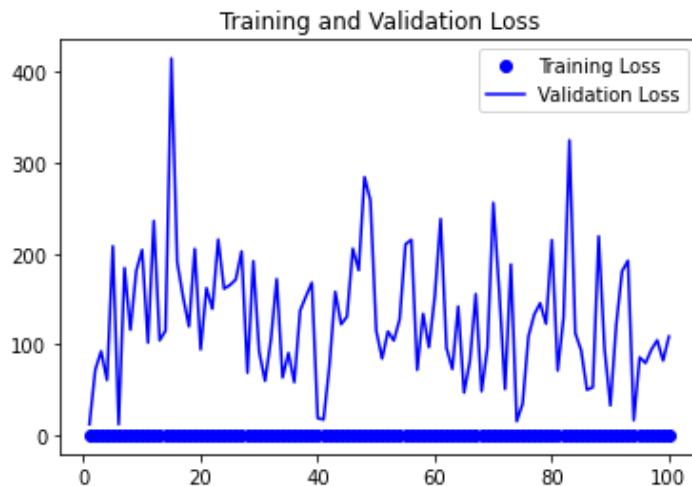
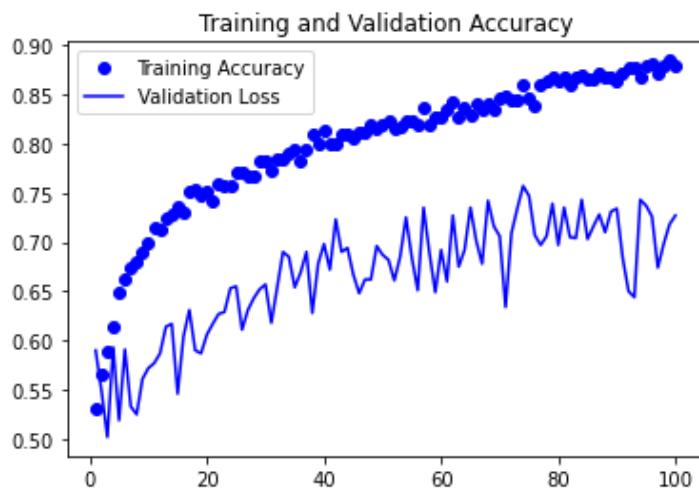
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training Accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation Loss')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training Loss')
plt.plot(epochs, val_loss, 'b', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.legend()

plt.show()
```



```
In [80]: # -----
# VISUALIZING INTERMEDIATE ACTIVATIONS /
# -----
```

```
model = load_model('cats_and_dogs_small_2.h5')
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_9 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_9 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_10 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_10 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_11 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_11 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_12 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_12 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_3 (Flatten)	(None, 6272)	0
dropout_2 (Dropout)	(None, 6272)	0
dense_5 (Dense)	(None, 512)	3211776
dense_6 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

```
In [86]: # -----
# PROCESSING A SINGLE IMAGE /
# -----
```

```
img_path = './Downloads/dogs_vs_cats_small_dataset/test/cats/cat.1700.jpg'

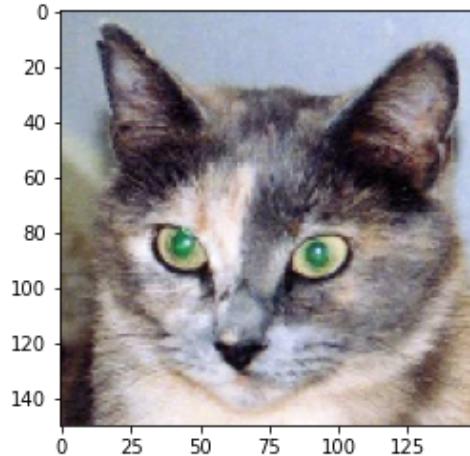
img = image.load_img(img_path, target_size=(150, 150))
img_tensor = image.img_to_array(img)
img_tensor = np.expand_dims(img_tensor, axis=0)
img_tensor /= 255. # the model was trained on inputs that were preprocessed this way

print(img_tensor.shape)
```

(1, 150, 150, 3)

```
In [87]: # -----  
# DISPLAYING THE TEST PICTURE |  
# -----
```

```
plt.imshow(img_tensor[0])  
plt.show()
```



```
In [90]: # Instantiaing a model from an input tensor and a list of output tensors
```

```
# extracts the outputs of the top 8 Layers  
layer_outputs = [layer.output for layer in model.layers[:8]]  
# creates a model that will return these outputs, given the model input  
activation_model = models.Model(inputs=model.input, outputs=layer_outputs)
```

```
In [91]: # Running the model in predict mode
```

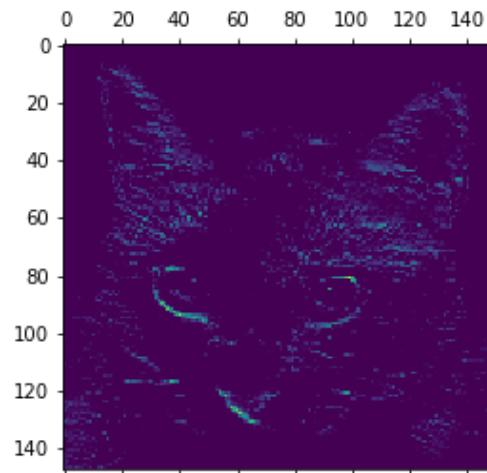
```
# returns a lsit of 5 numpy arrays: 1 array per Layer activations  
activations = activation_model.predict(img_tensor)  
  
first_layer_activation = activations[0]  
print(first_layer_activation.shape)
```

```
(1, 148, 148, 32)
```

```
In [92]: # Visualizing the fourth channel
```

```
plt.matshow(first_layer_activation[0, :, :, 4], cmap='viridis')
```

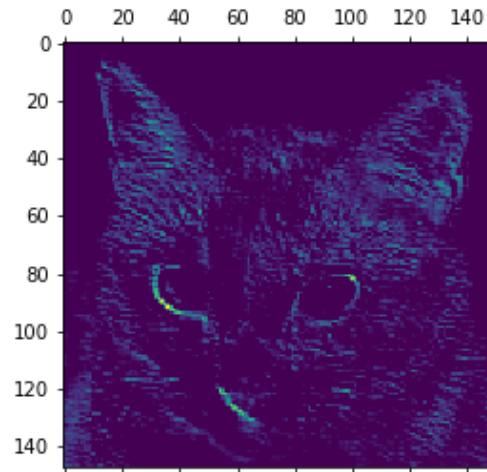
```
Out[92]: <matplotlib.image.AxesImage at 0x264a8197288>
```



```
In [93]: # Visualizing the seventh channel
```

```
plt.matshow(first_layer_activation [0, :, :, 7], cmap='viridis')
```

```
Out[93]: <matplotlib.image.AxesImage at 0x264a1c33448>
```



```
In [94]: # ======  
# E  
# Visualize every channel in every intermediate activations and explain why |  
# it is useful |  
# ======
```

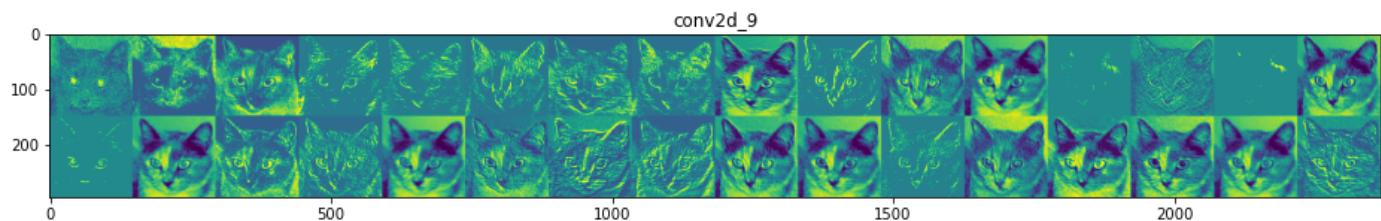
```
In [95]: # -----  
# EXPLAIN WHY IT IS USEFUL |  
# -----  
  
#"Intermediate activations are useful for understanding how successive  
# convnet layers transform their input, and for getting a first idea of the  
# meaning of individual convnet filters.  
# Visualizing intermediate activations consists of displaying the feature maps  
# that are output by various convolution and pooling layers in a network,  
# given a certain output. This gives a view into how an input is decomposed  
# into different filters learned by the network. Each channel encodes  
# relatively independent features, so the proper way to visualize these  
# feature maps is by independently plotting the contours of every channel as a  
# 2D image" - Francois Chollet "DEEP LEARNING with Python"
```

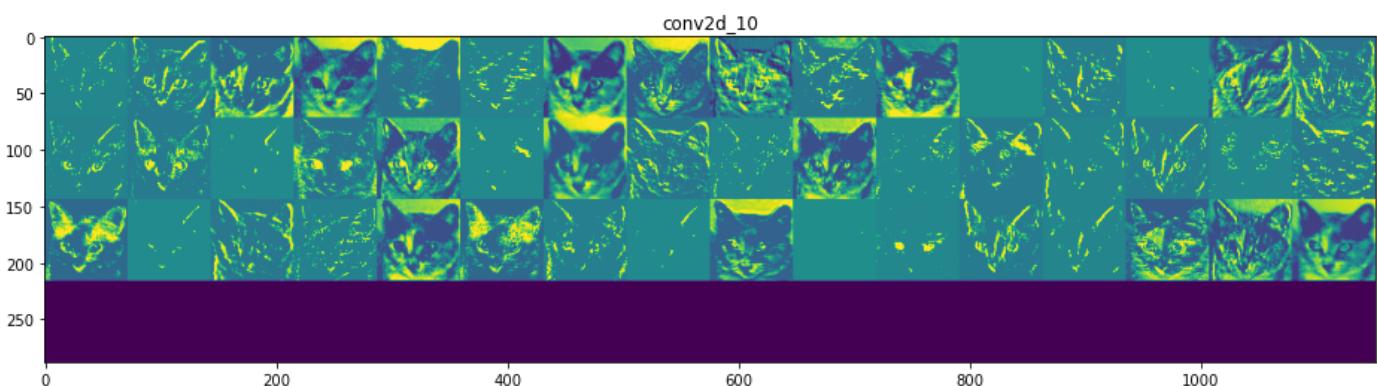
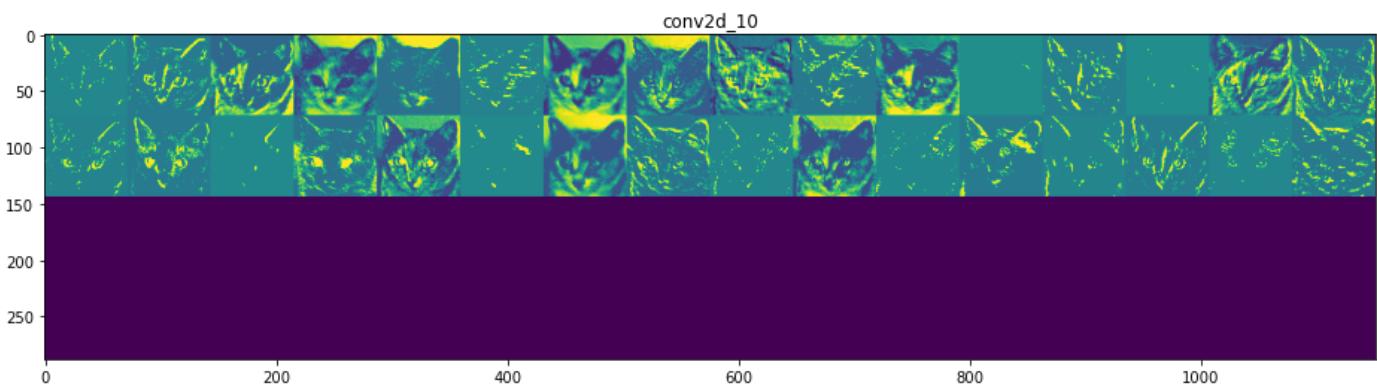
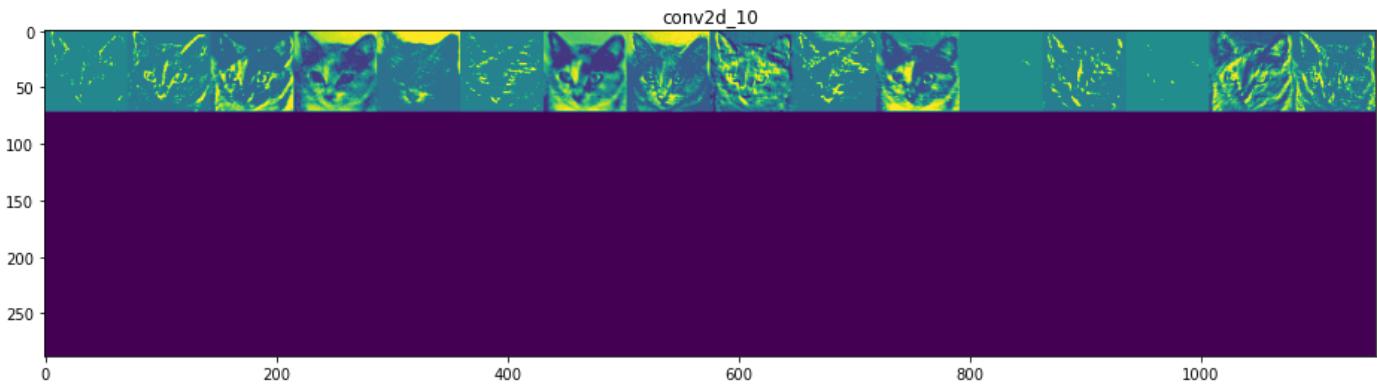
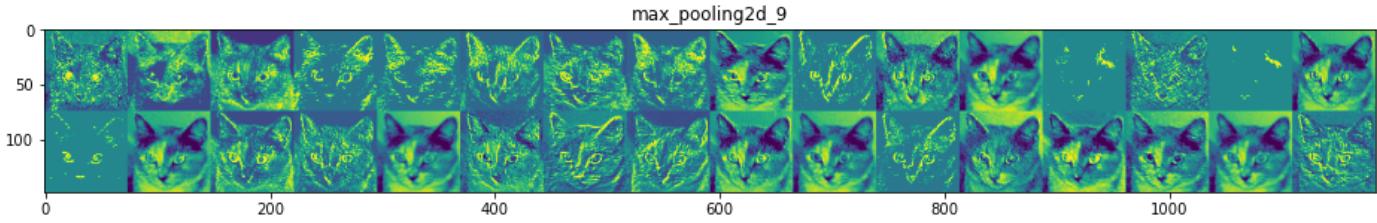
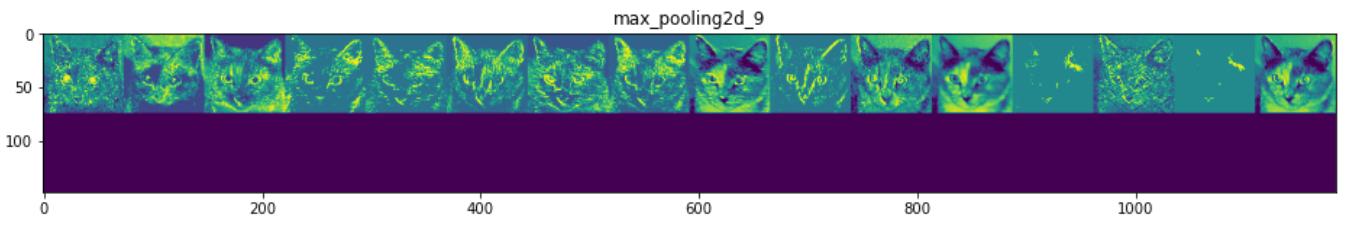
In [99]:

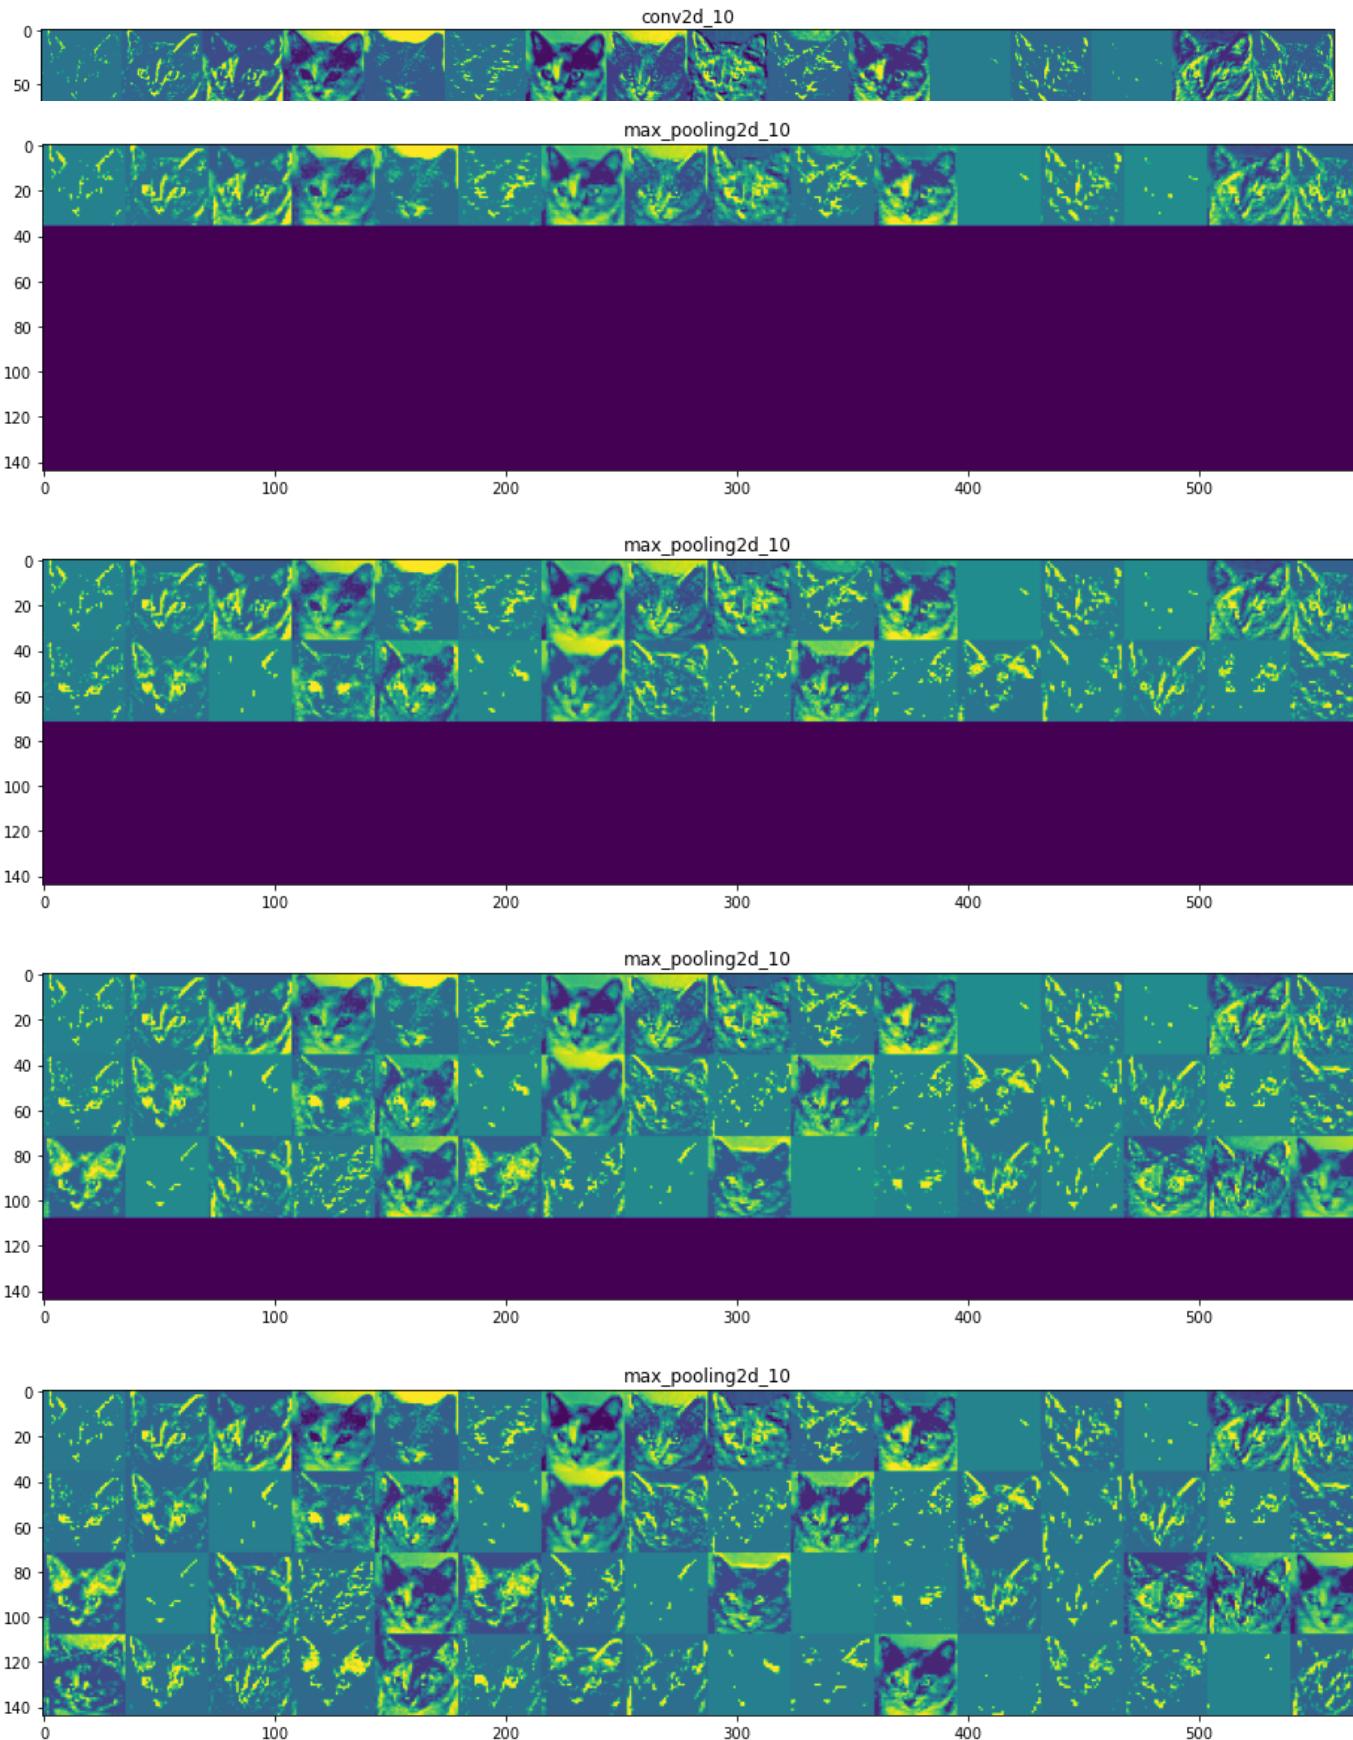
```
# -----#
# VISUALIZING EVERY CHANNEL IN EVERY INTERMEDIATE ACTIVATION |
# -----#  
  
# Names of the layers so you can have them as part of your plot  
layer_names = []  
for layer in model.layers[:8]:  
    layer_names.append(layer.name)  
  
images_per_row = 16 # displays the feature maps  
  
for layer_name, layer_activation in zip(layer_names, activations):  
    n_features = layer_activation.shape[-1] # number of features in the feature map  
  
    size = layer_activation.shape[1] # has shape (l, size, size, n_features)  
  
    n_cols = n_features // images_per_row # tiles the activation channels in this matrix  
    display_grid = np.zeros((size * n_cols, images_per_row * size))  
  
    for col in range(n_cols): # tiles each filter into a big horizontal grid  
        for row in range(images_per_row):  
            channel_image = layer_activation[0, :, :, col * images_per_row + row]  
            channel_image -= channel_image.mean() # post processes the feature to make it visual  
            channel_image /= channel_image.std()  
            channel_image *= 64  
            channel_image += 128  
            channel_image = np.clip(channel_image, 0, 255).astype('uint8')  
            display_grid[col * size : (col + 1) * size, # displays the grid  
                         row * size : (row + 1) * size] = channel_image  
  
    scale = 1./size  
    plt.figure(figsize=(scale * display_grid.shape[1],  
                     scale * display_grid.shape[0]))  
    plt.title(layer_name)  
    plt.grid(False)  
    plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

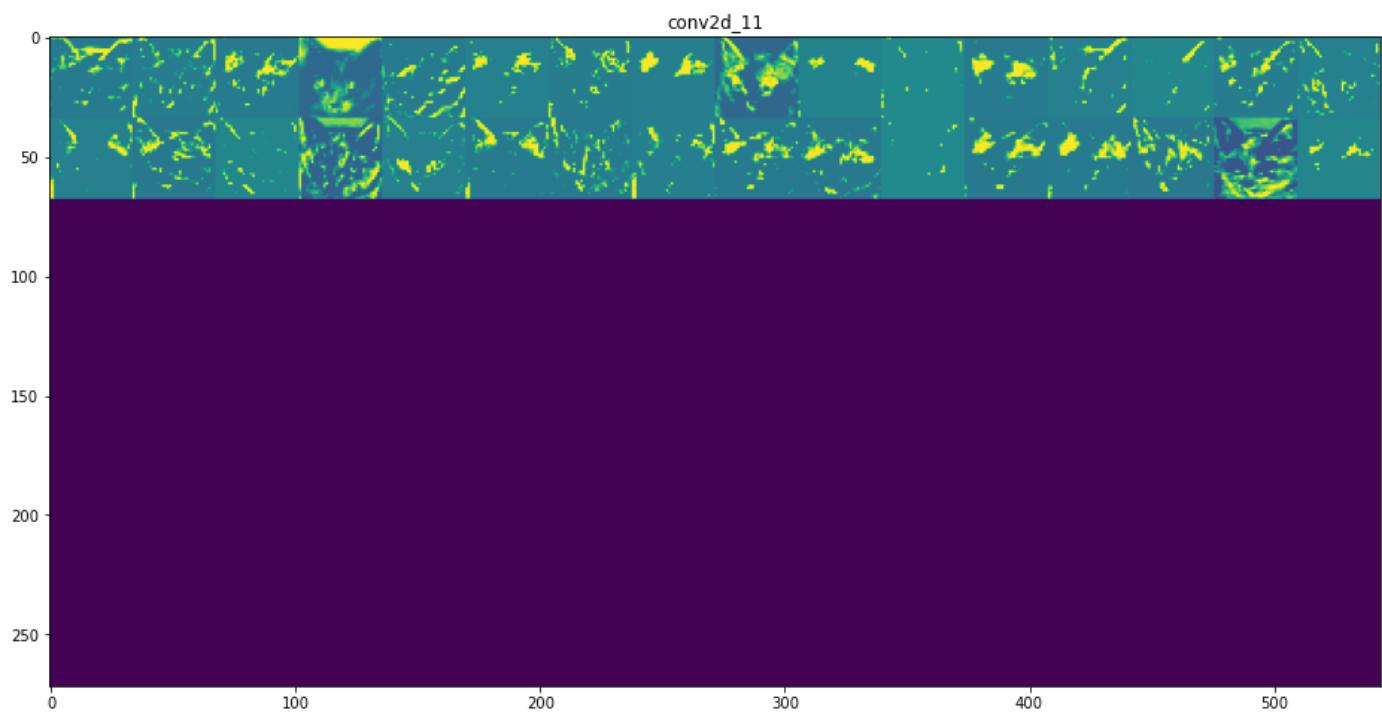
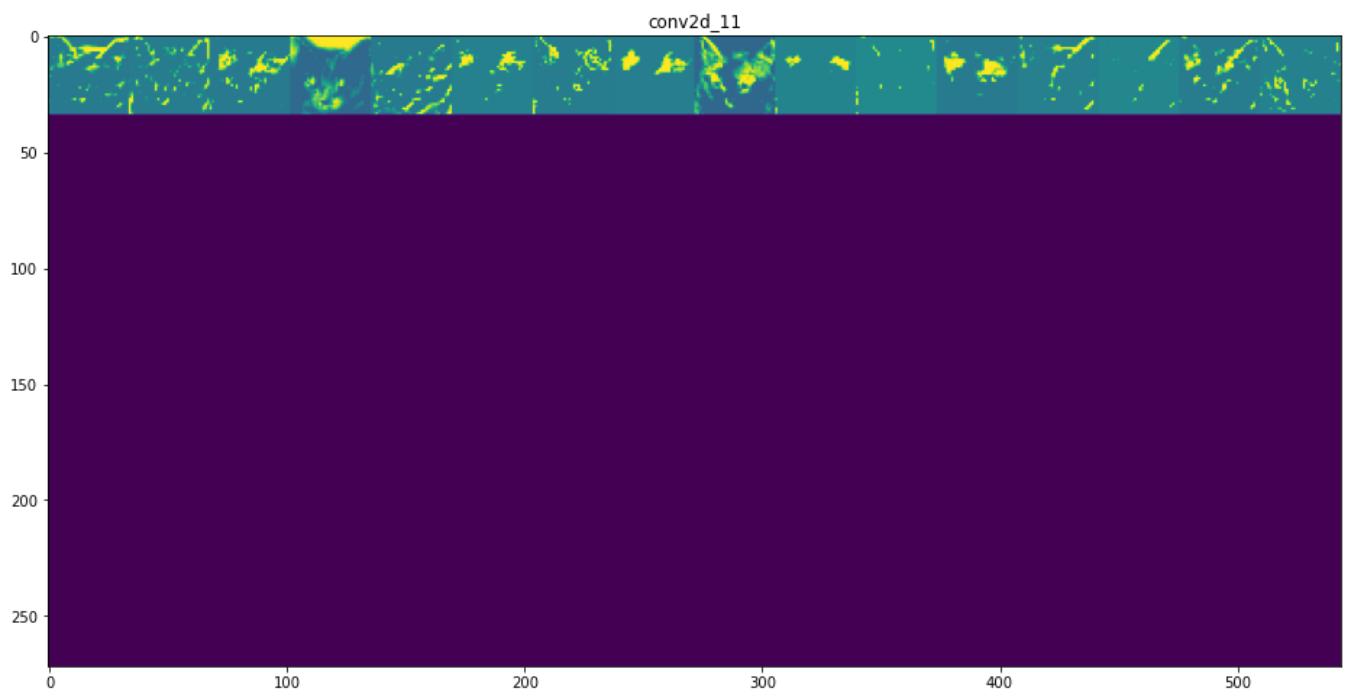
C:\Users\User\Anaconda3\lib\site-packages\ipykernel_launcher.py:24: RuntimeWarning: invalid value encountered in true_divide

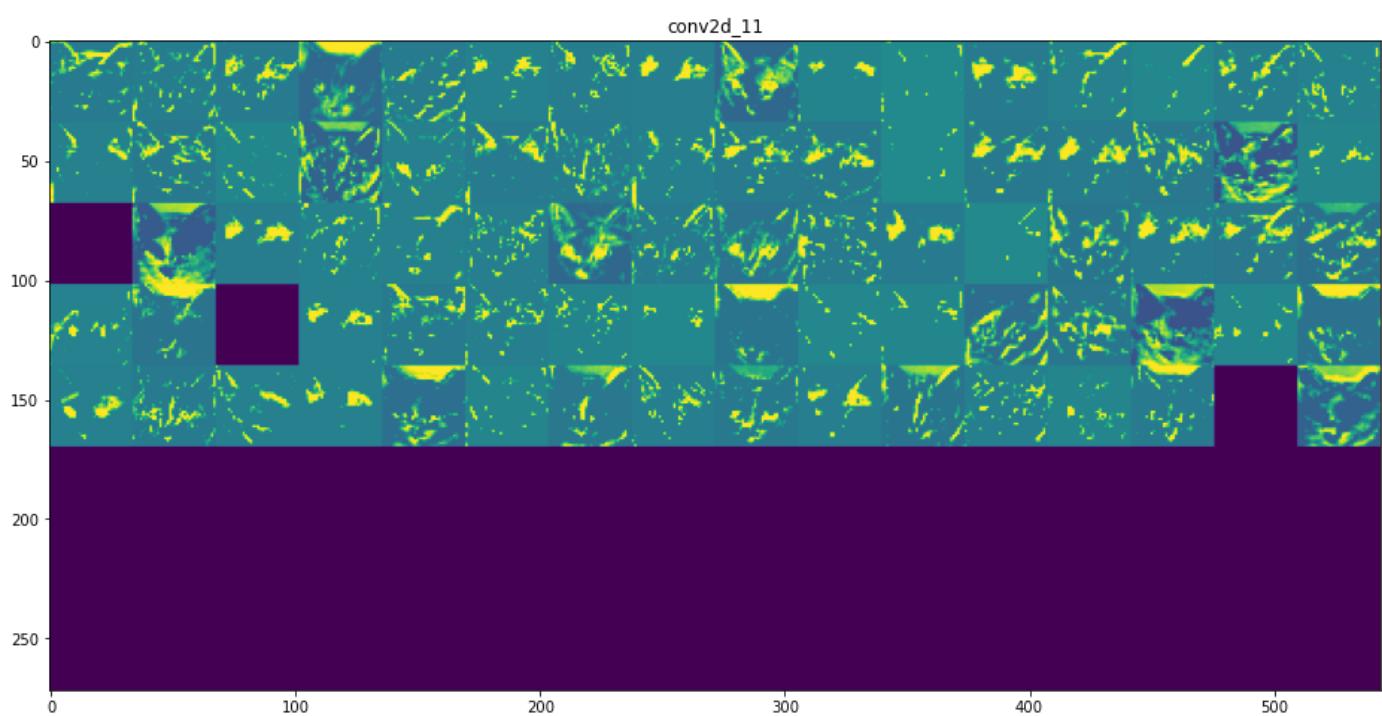
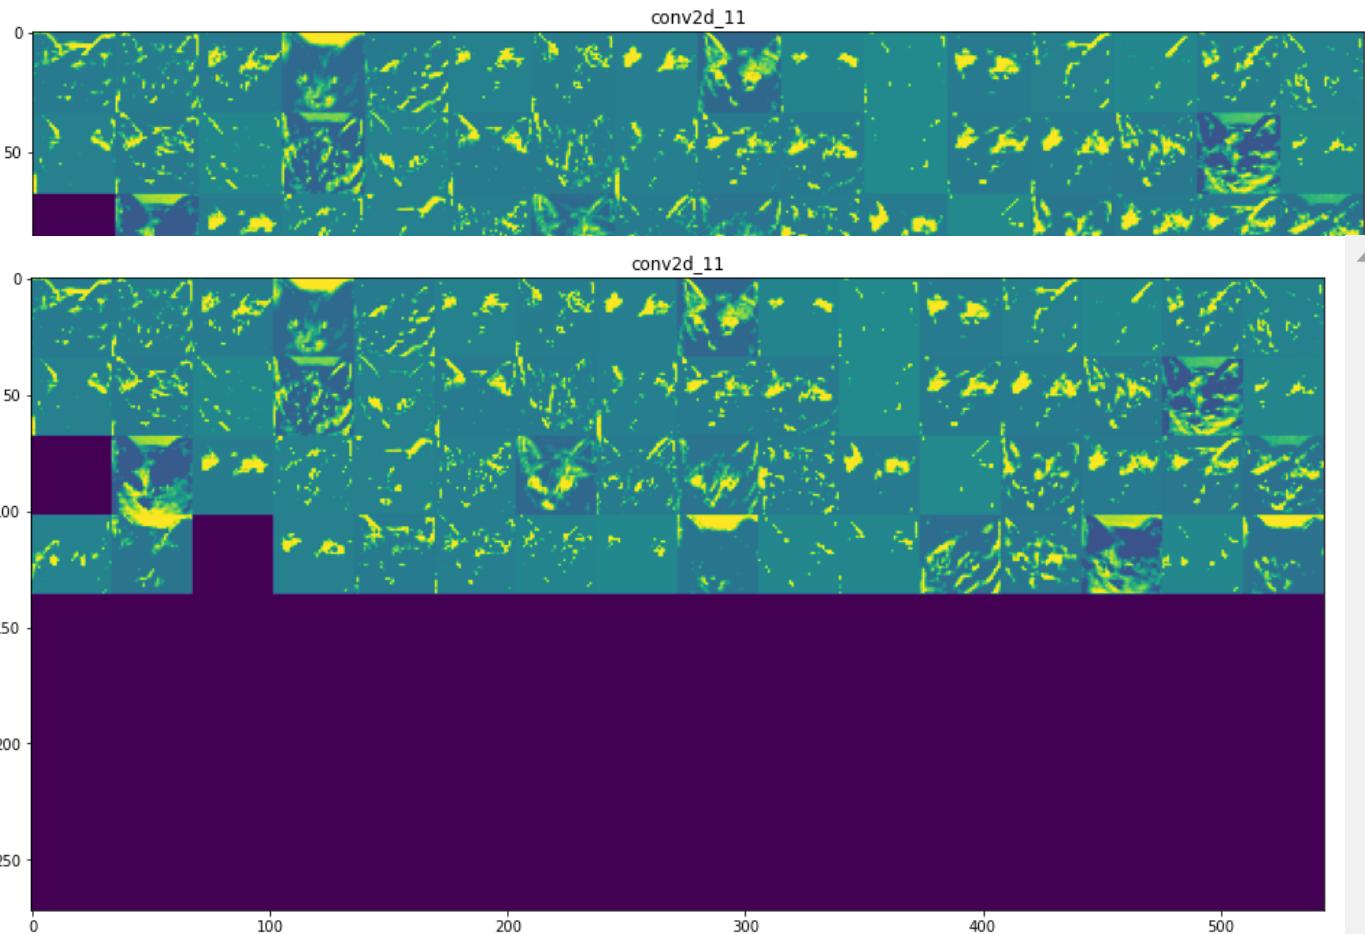
C:\Users\User\Anaconda3\lib\site-packages\ipykernel_launcher.py:33: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

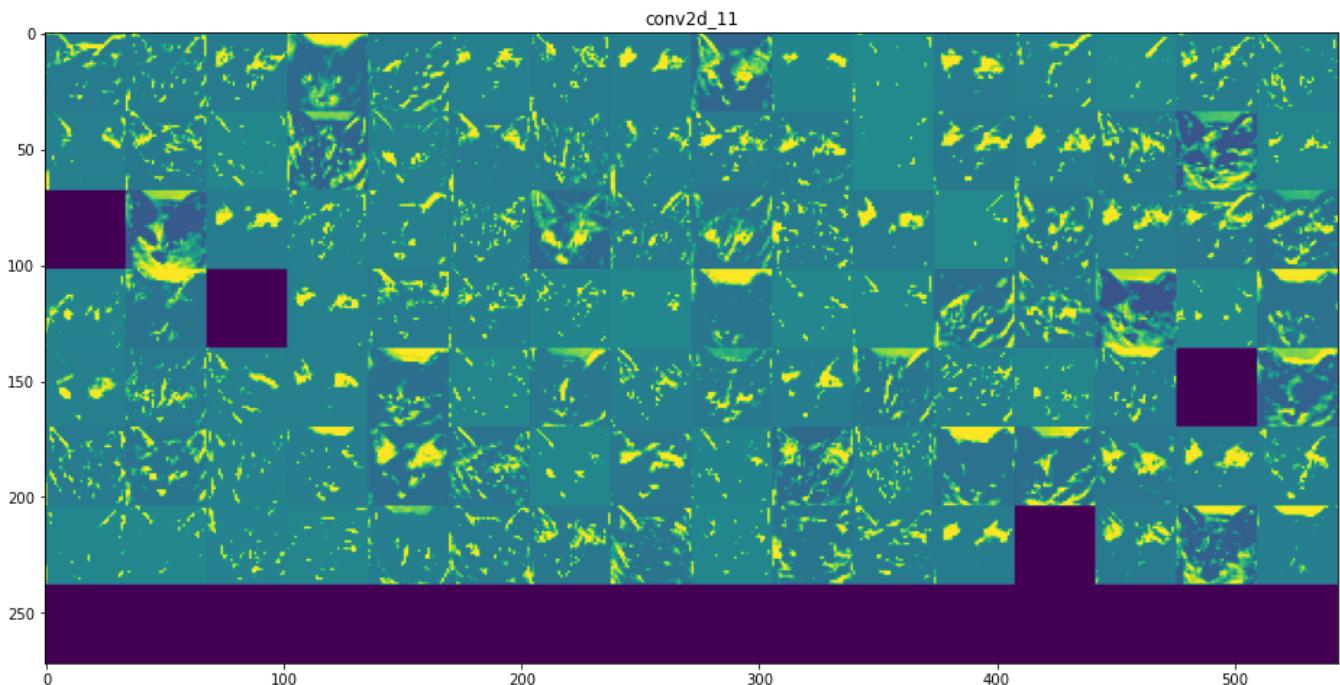
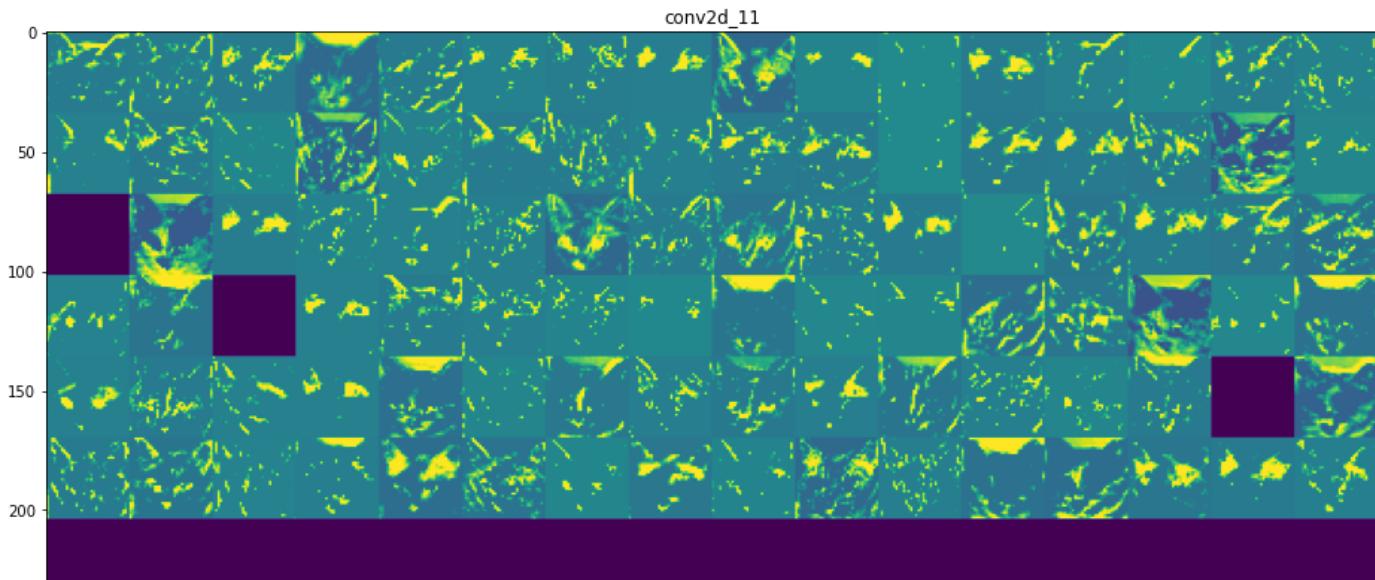






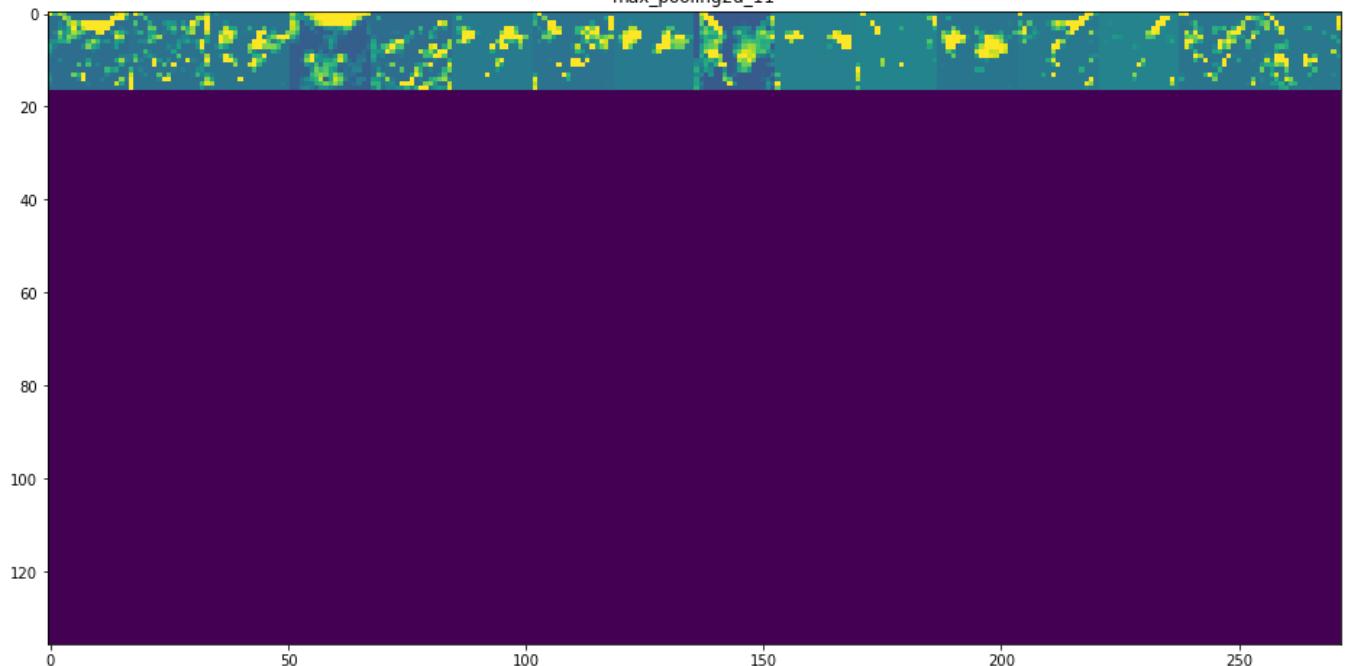


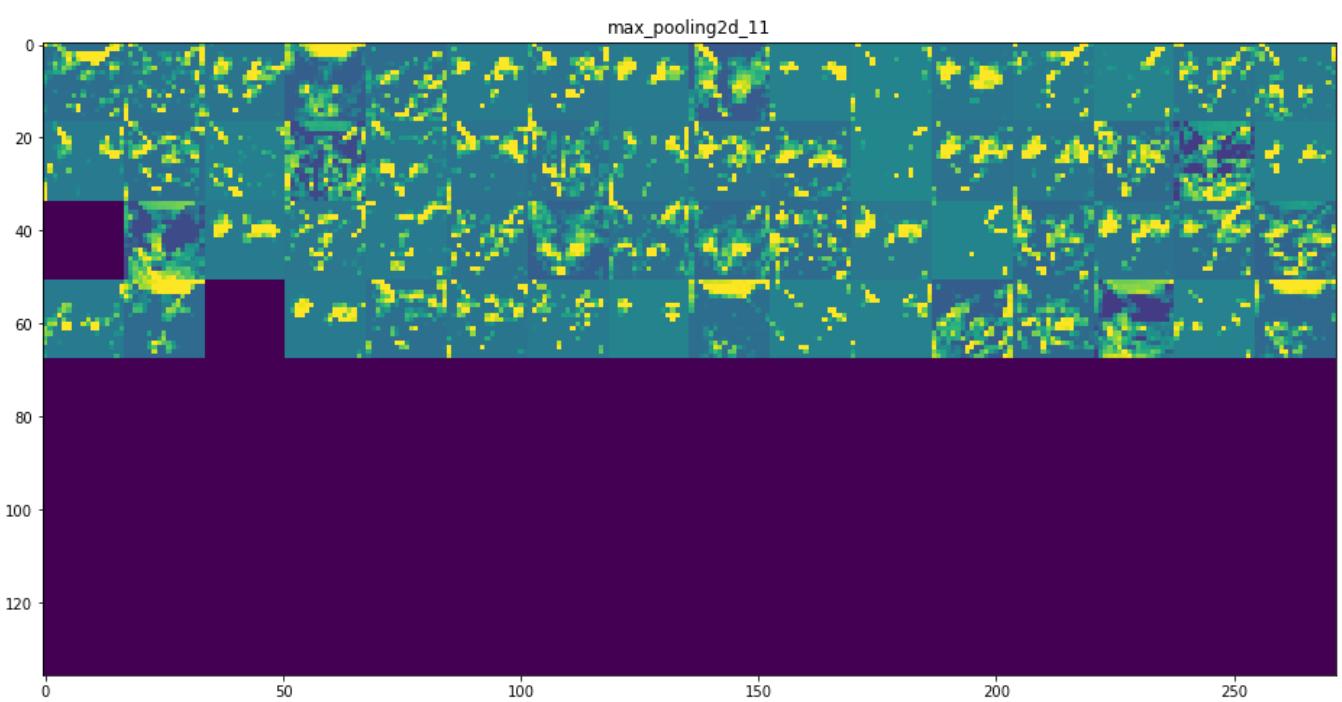
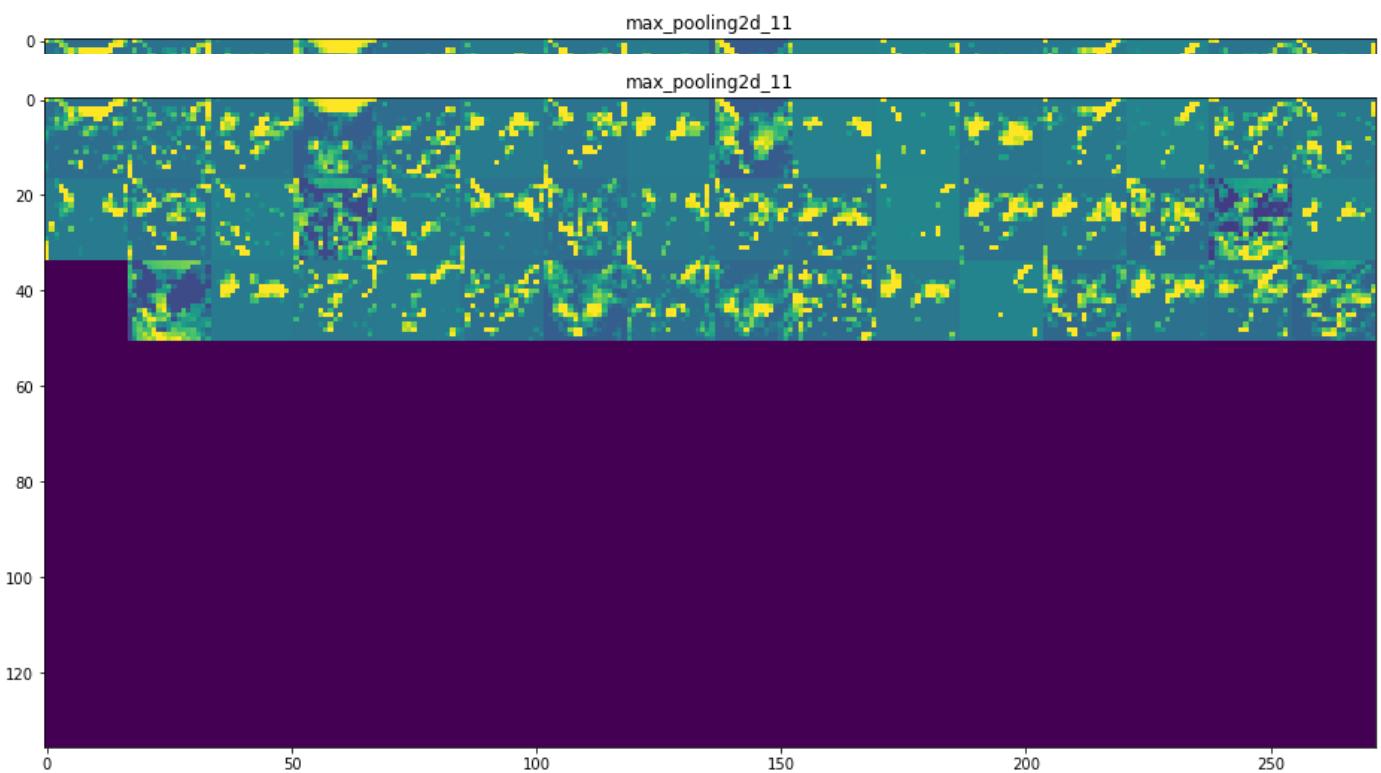


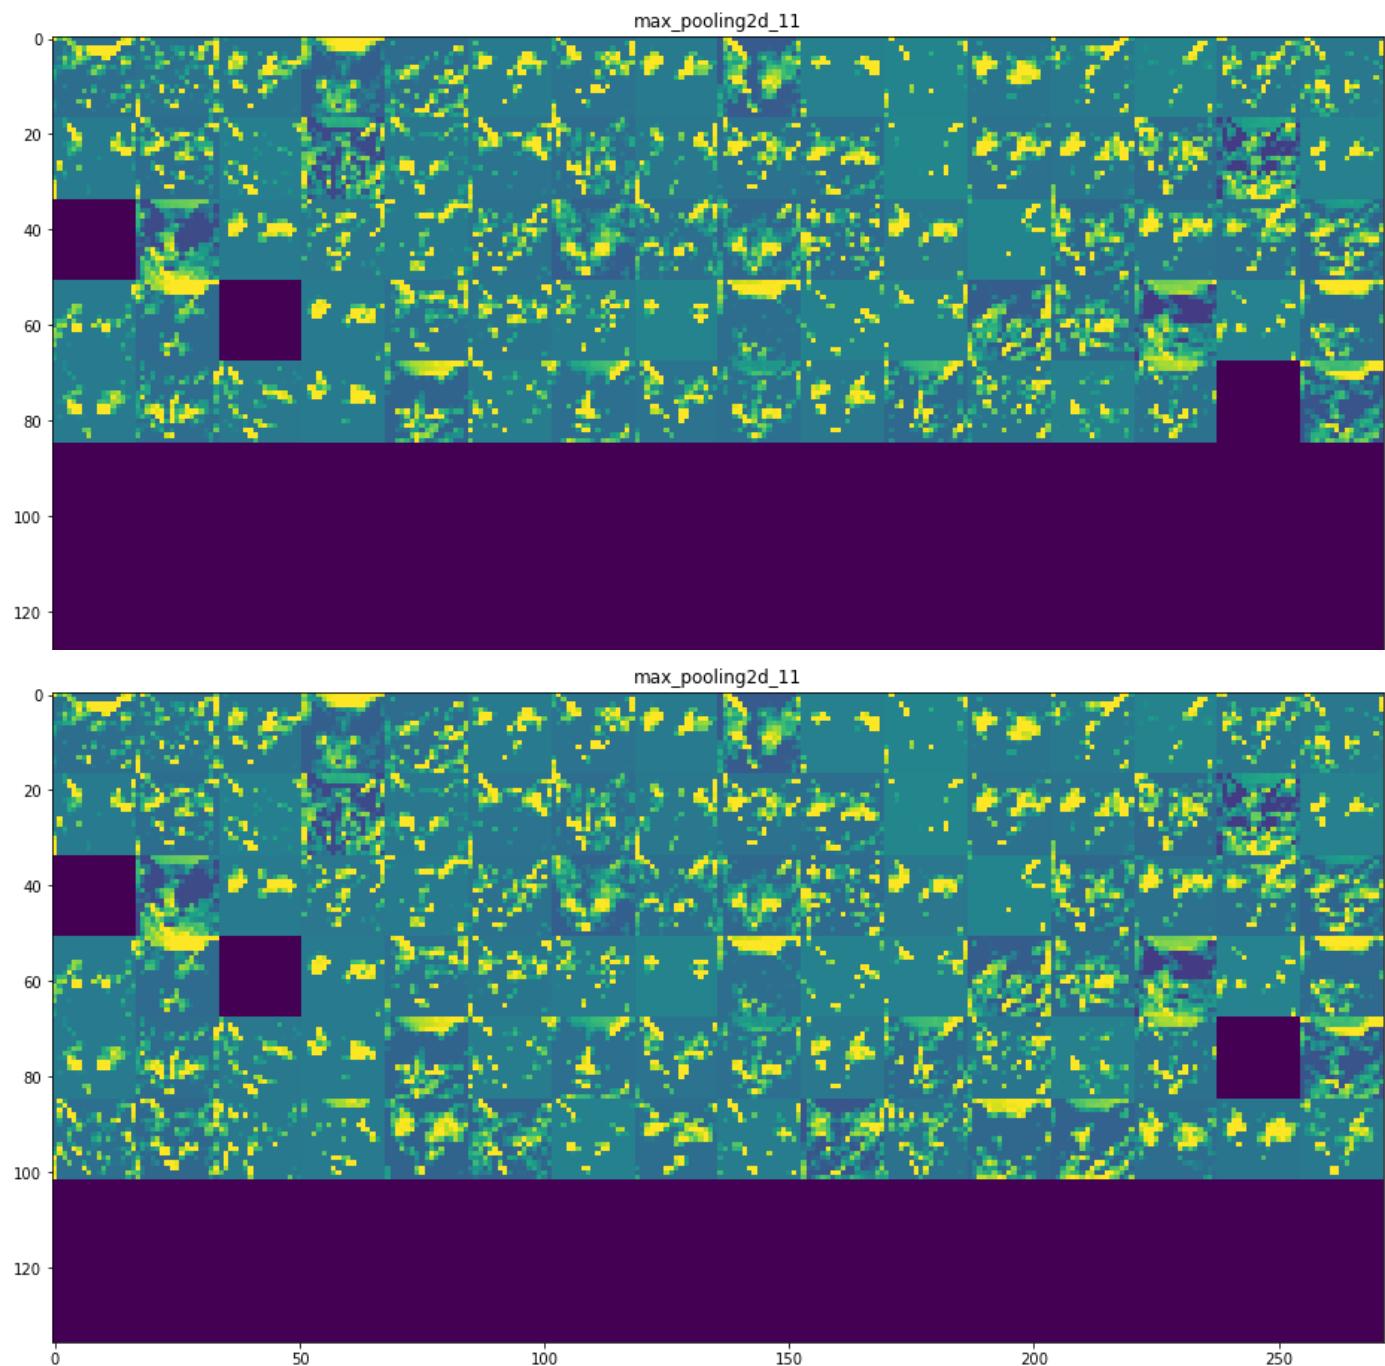


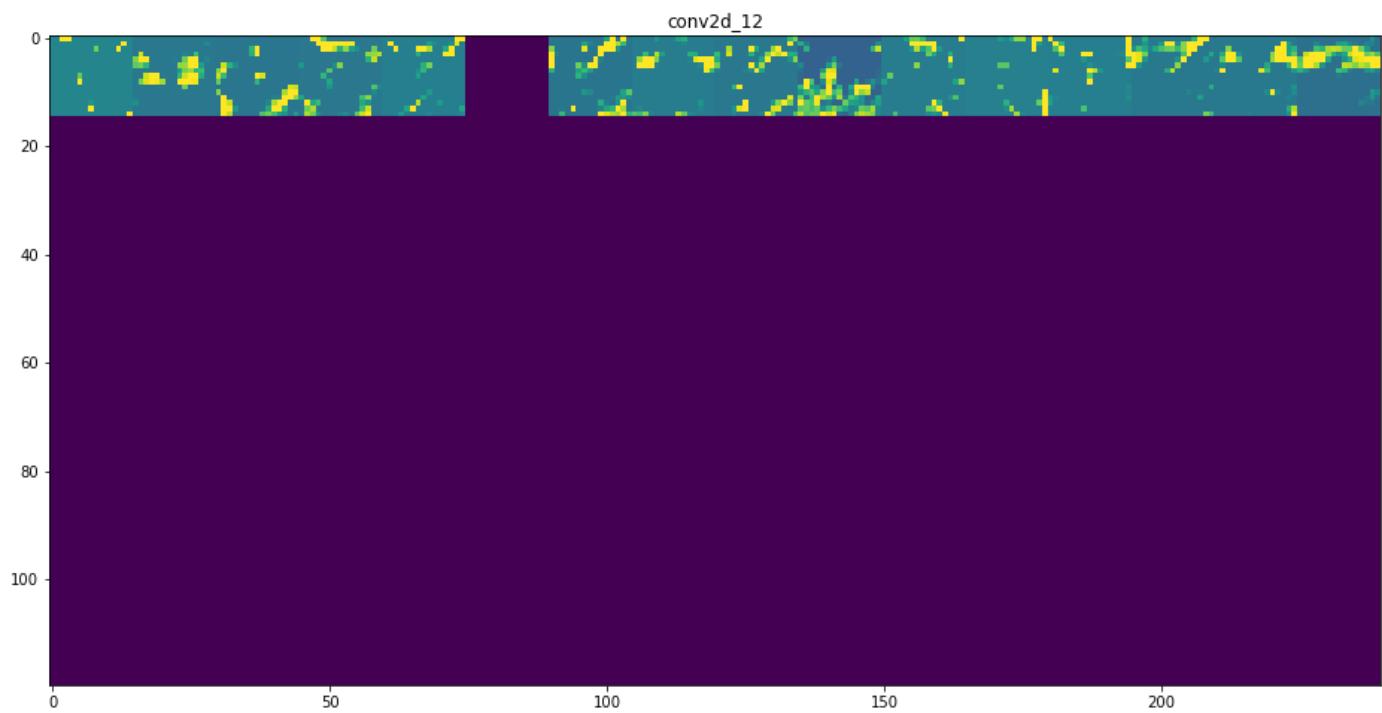
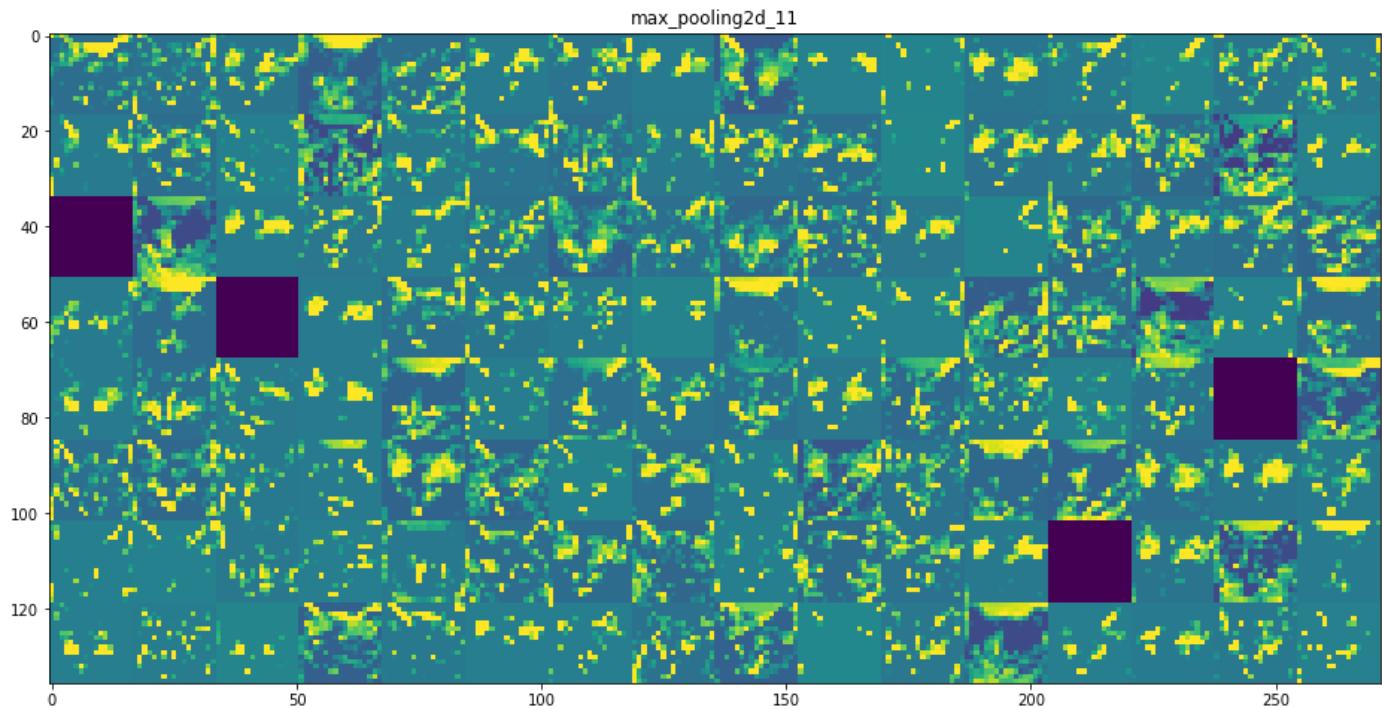
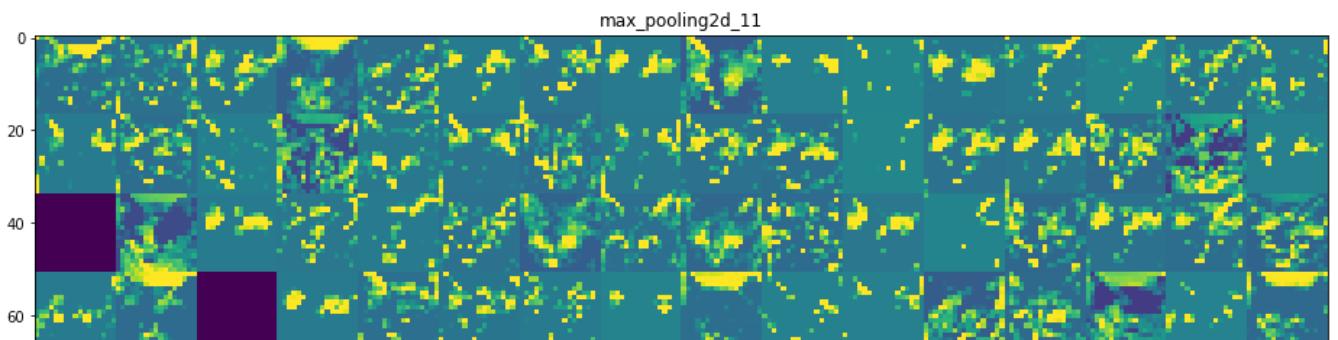
conv2d_11

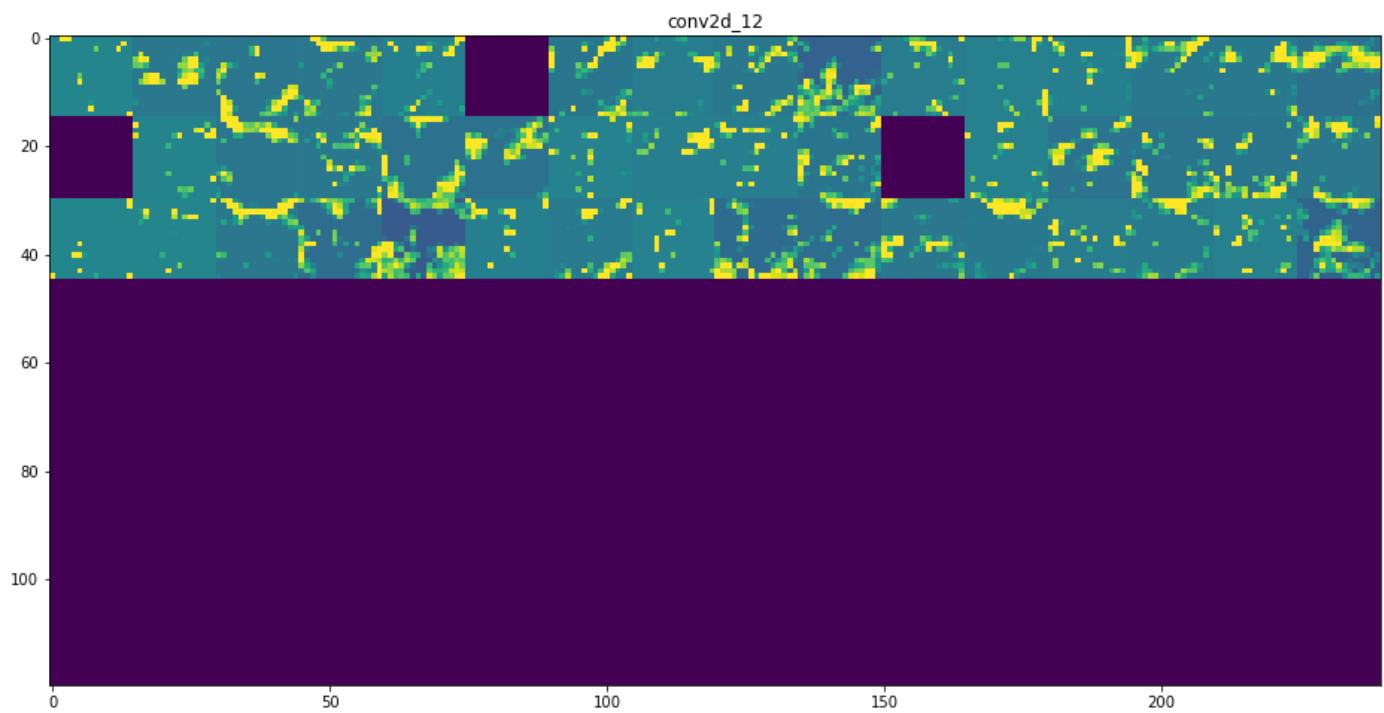
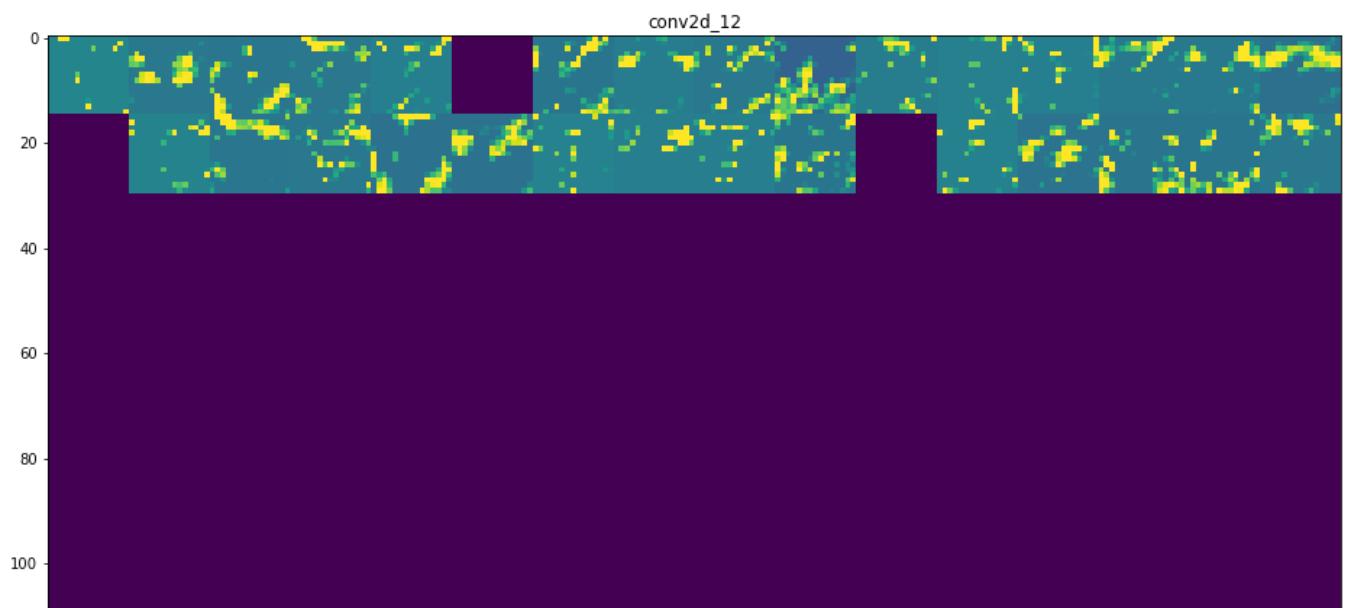
max_pooling2d_11

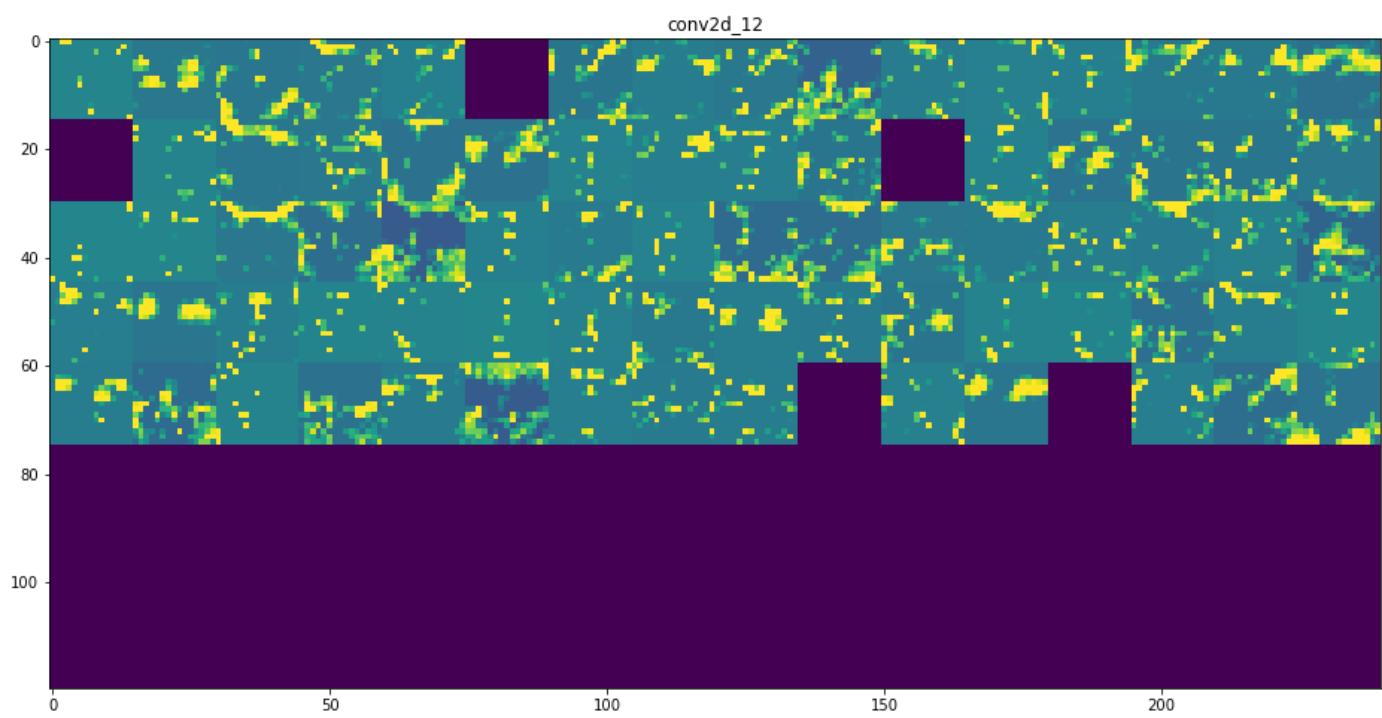
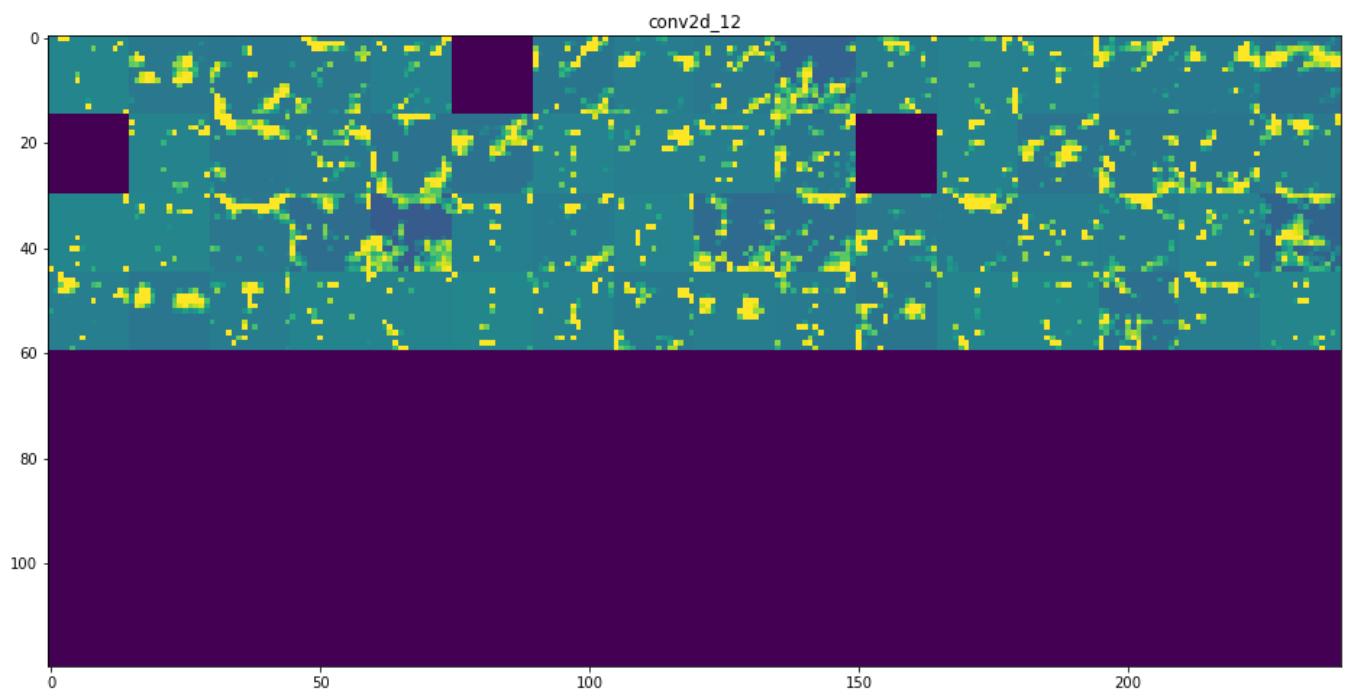


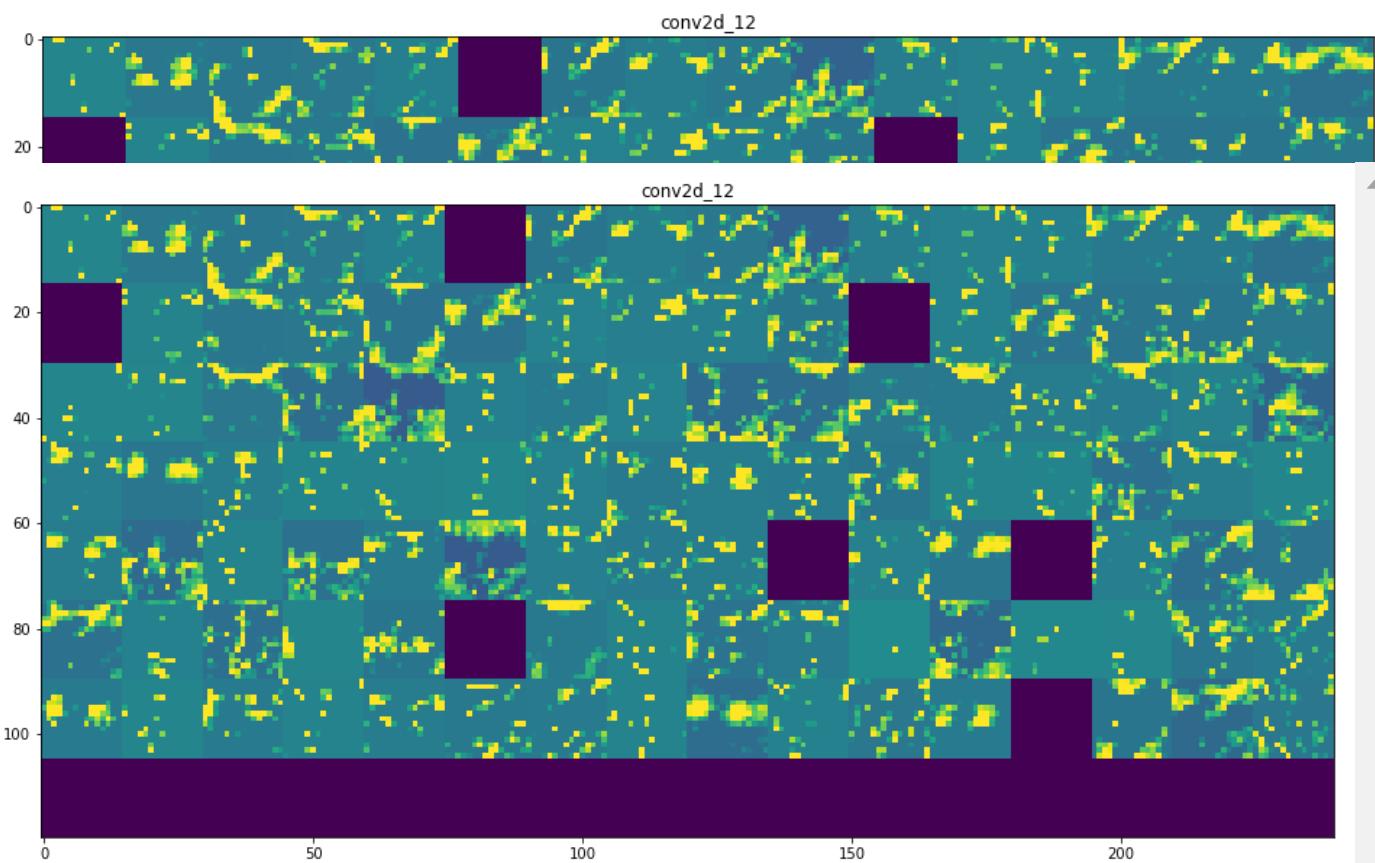


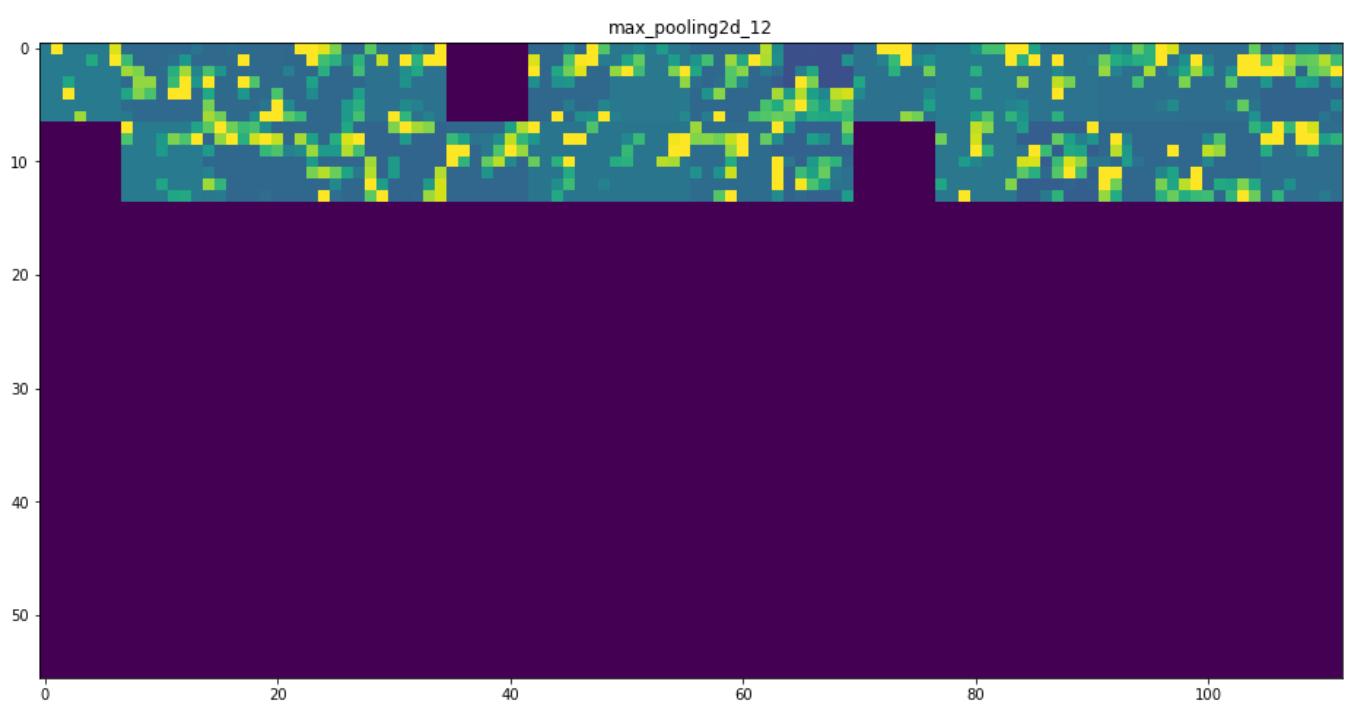
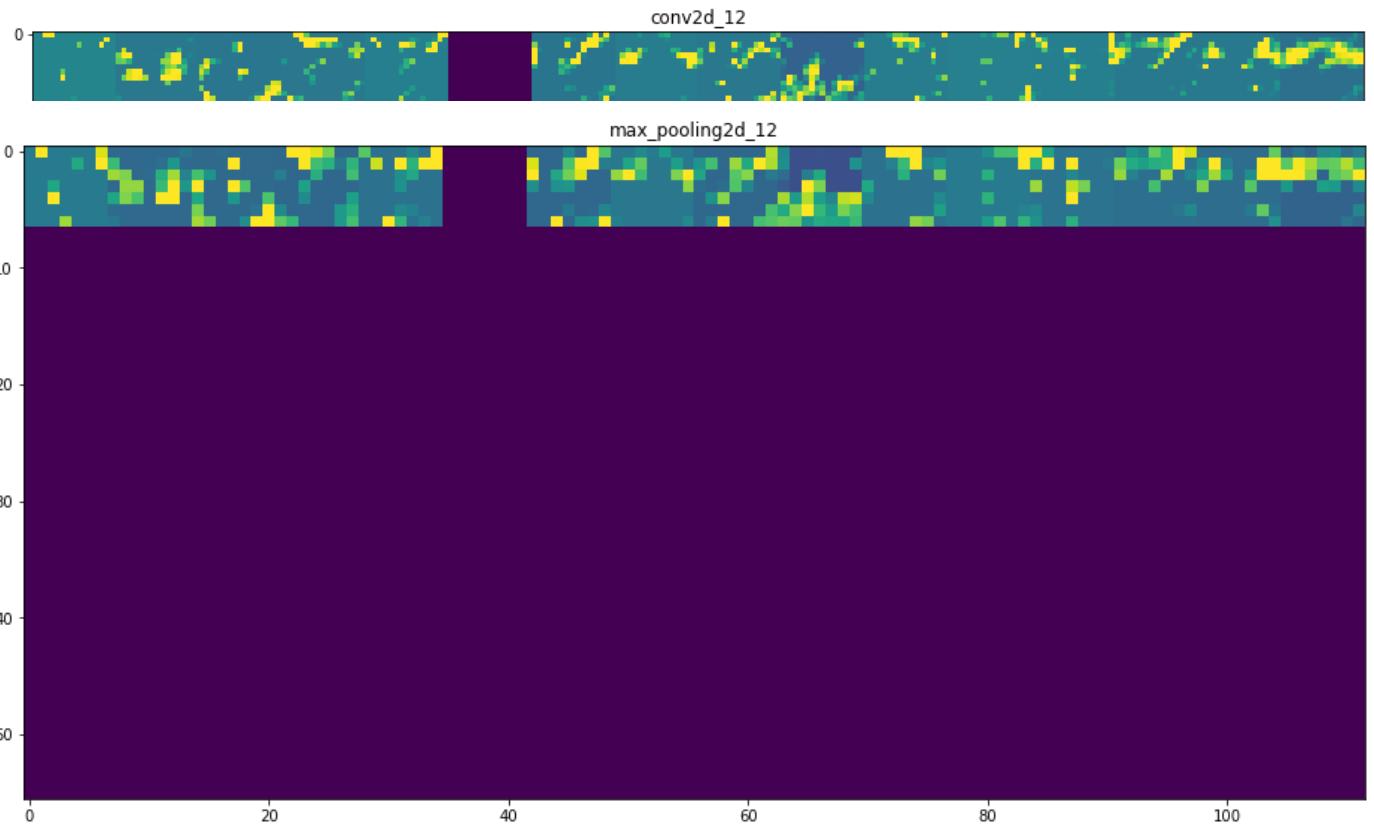


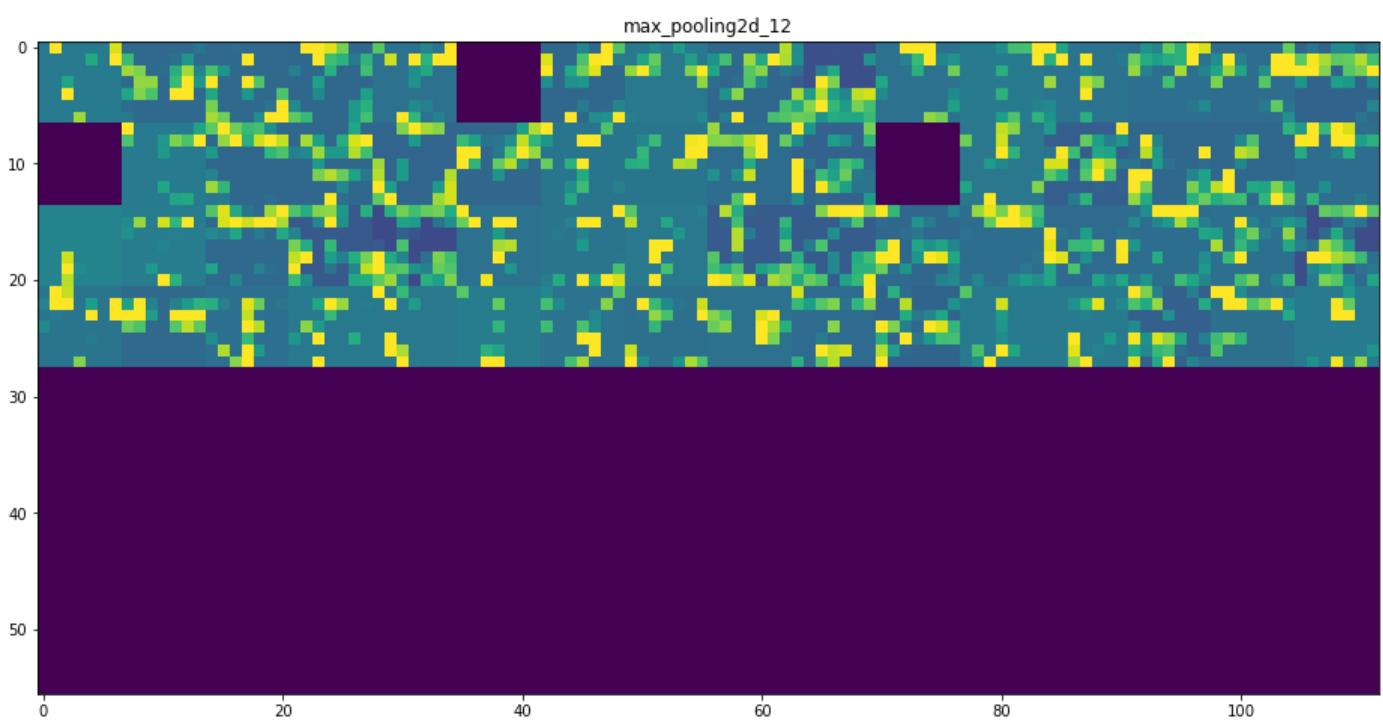
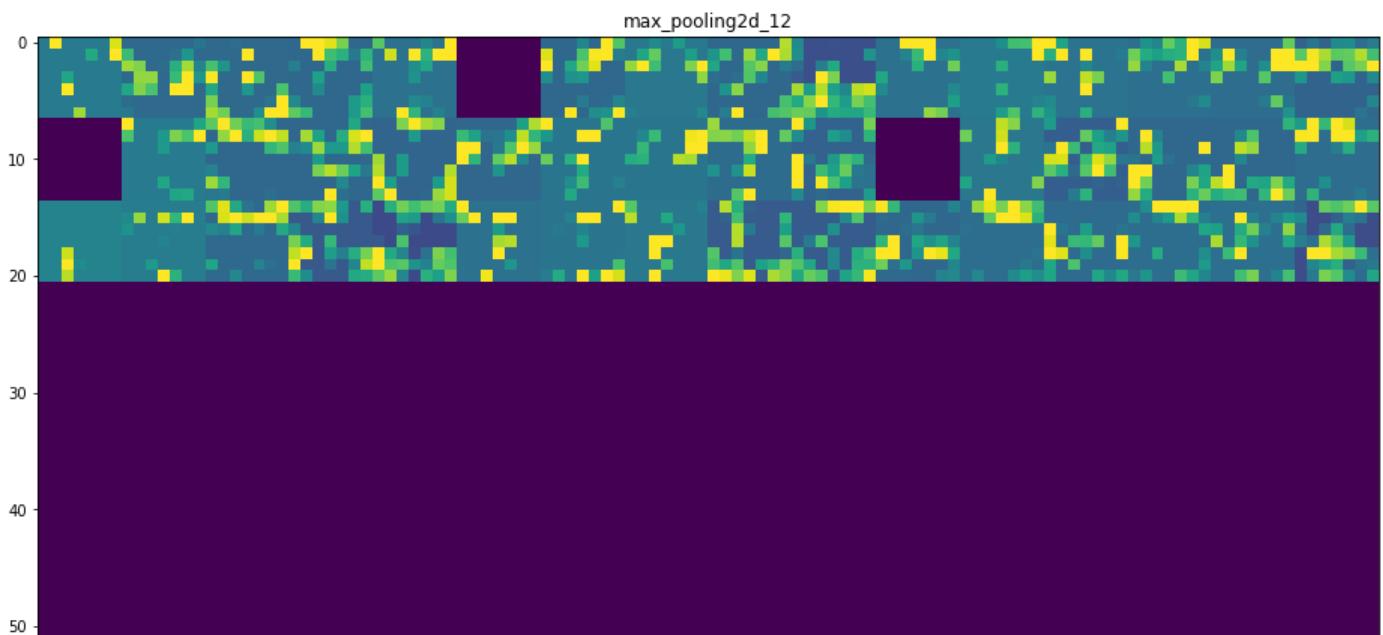


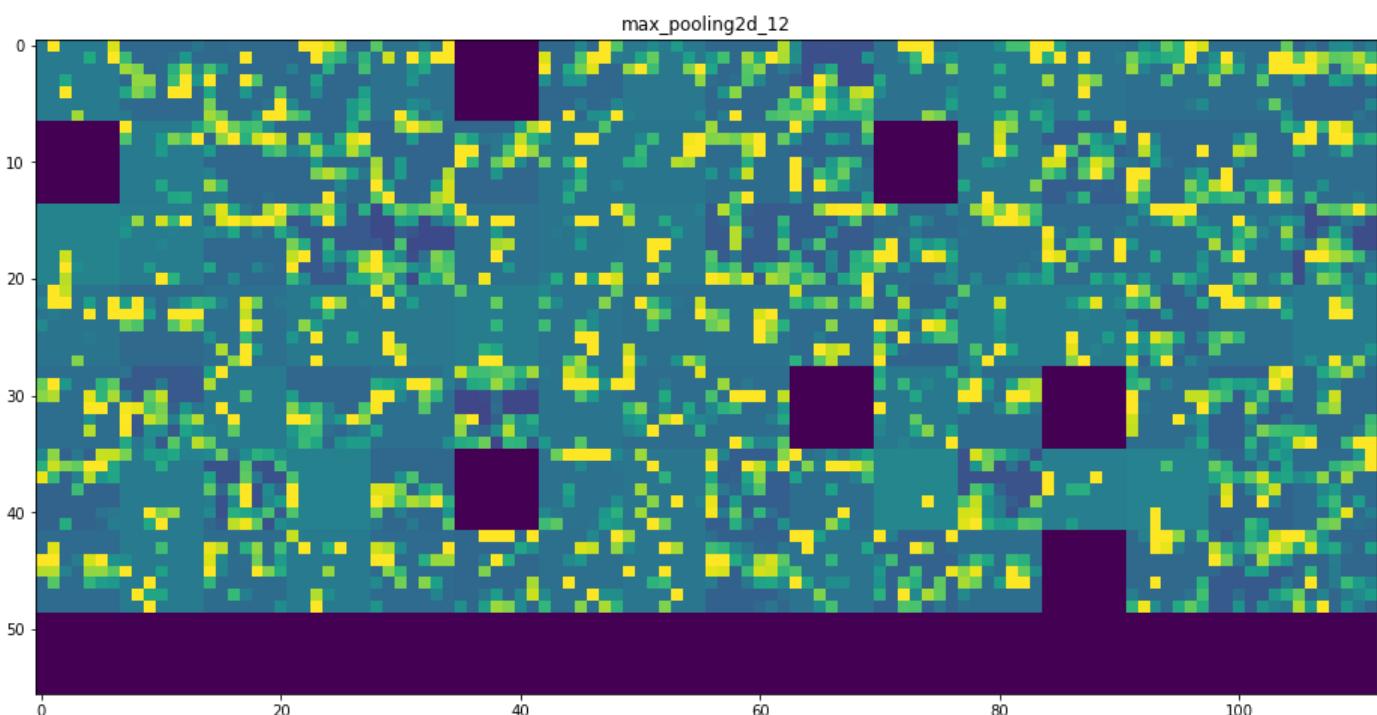
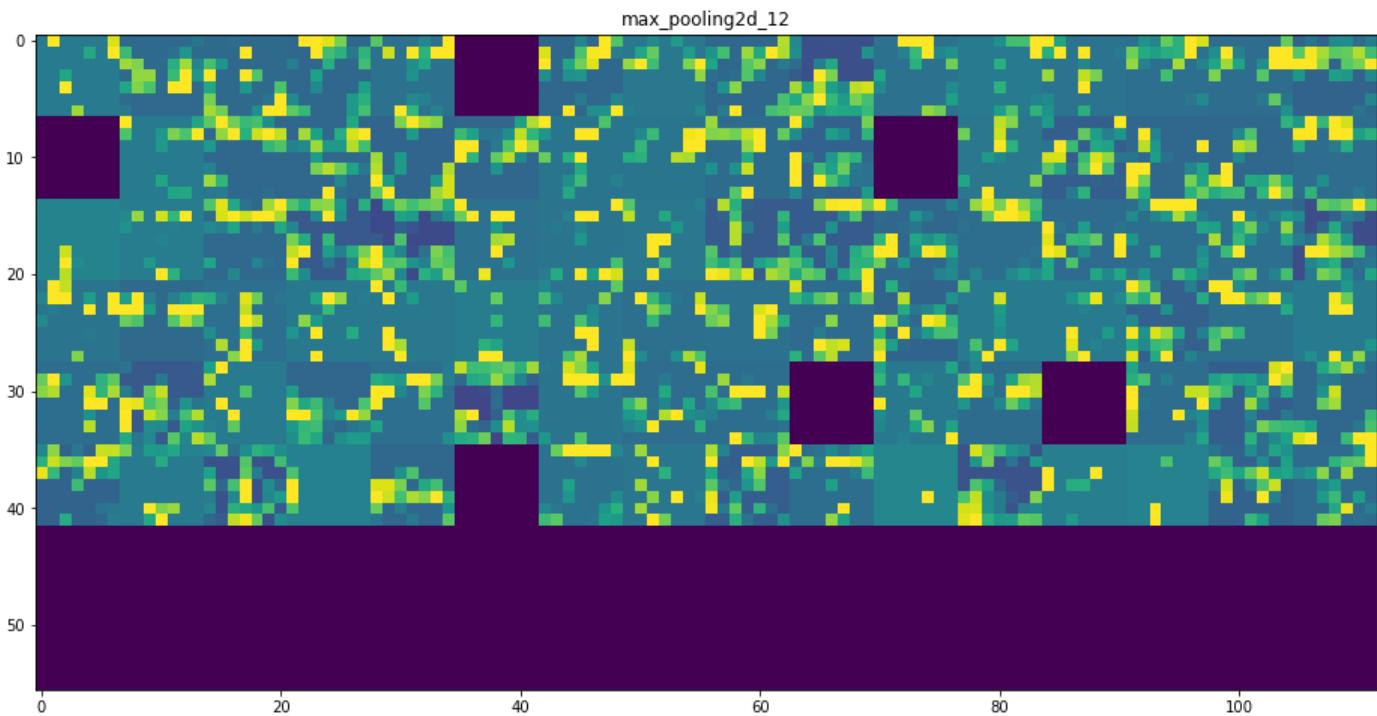
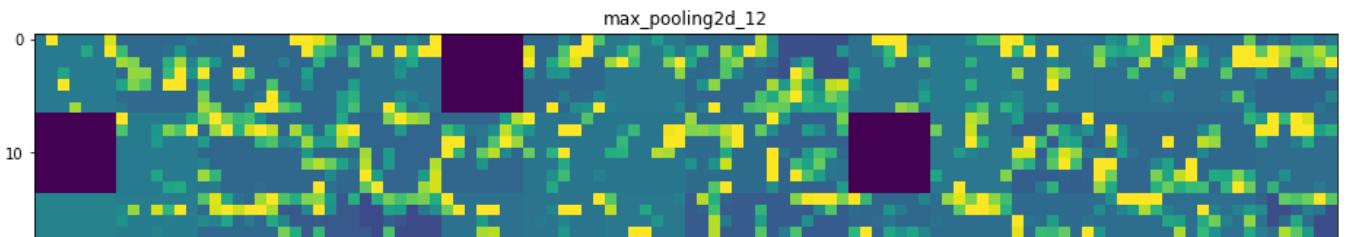


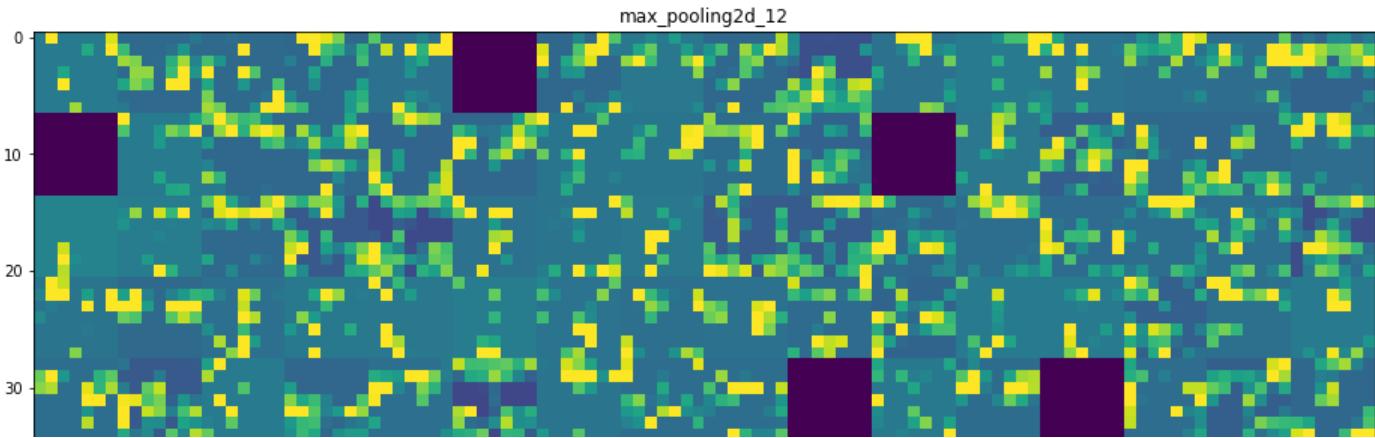












```
In [104]: # -----
# F
# Visualize convolutional filters: get the gradient of the loss with regard |
# to the input, apply stochastic gradient descent, include a code for filter |
# visualization and generate a grid of all filter response patterns in a layer|
# -----
```

```
In [103]: # -----
# DEFINING THE LOSS TENSOR FOR FILTER VISUALIZATION |
# -----
model = VGG16(weights='imagenet', include_top=False)
layer_name = 'block3_conv1'
filter_index = 0

layer_output = model.get_layer(layer_name).output
loss = K.mean(layer_output[:, :, :, filter_index])
```

```
In [105]: # -----
# OBTAINING THE GRADIENT OF THE LOSS WITH REGARD TO THE INPUT |
# -----
grads = K.gradients(loss, model.input)[0]
# the call to gradients returns a list of tensors (of size 1 in this case)
# keeps only the 1st elements which is a tensor
```

```
In [106]: # -----
# GRADIENT-NORMALIZATION TRICK |
# -----
grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)
# add 1e-5 before dividing to avoid accidentally dividing by 0
```

```
In [107]: # -----
# FETCHING NUMPY OUTPUT VALUES GIVEN NUMPY INPUT VALUES |
# -----
iterate = K.function([model.input], [loss, grads])

loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])
```

In [109]:

```
# -----
# LOSS MAXIMIZATION VIA STOCHASTIC GRADIENT DESCENT /
# -----  
  
input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.  
# starts from a grey image with some noise  
  
step = 1. # magnitude of each gradient update  
for i in range(40): # runs gradient ascent for 40 steps  
    loss_value, grads_value = iterate([input_img_data])  
    # computes the loss value and gradient value  
  
    input_img_data += grads_value * step  
    # adjusts the input image in the direction that maximizes the loss
```

In [110]:

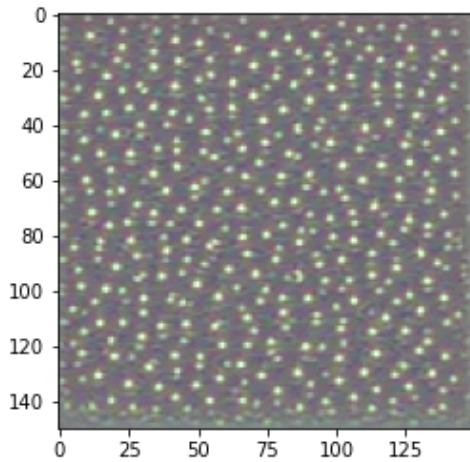
```
# -----
# UTILITY FUNCTION TO CONVERT A TENSOR INTO A VALID IMAGE /
# -----  
  
def deprocess_image(x):  
    # Normalizes the tensorL centers on 0, ensure that std is 0.1  
    x -= x.mean()  
    x /= (x.std() + 1e-5)  
    x *= 0.1  
  
    # Clips to [0, 1]  
    x += 0.5  
    x = np.clip(x, 0, 1)  
  
    # Converts to an RGB array  
    x *= 255  
    x = np.clip(x, 0, 255).astype('uint8')  
    return x
```

In [112]: # -----

```
# FUNCTION TO GENERATE FILTER VISUALIZATIONS /  
# -----  
  
def generate_pattern(layer_name, filter_index, size=150):  
    # Builds a loss function that maximizes that activation of the nth  
    # filter of the layer under consideration  
    layer_output = model.get_layer(layer_name).output  
    loss = K.mean(layer_output[:, :, :, filter_index])  
  
    # Computes the gradient of the input picture with regard to this loss  
    grads = K.gradients(loss, model.input)[0]  
  
    # Normalization trickL normalizes the gradient  
    grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)  
  
    # Returns the loss and grads given the input picture  
    iterate = K.function([model.input], [loss, grads])  
  
    # Starts from a grey image with some noise  
    input_img_data = np.random.random((1, size, size, 3)) * 20 + 128  
  
    # Runs gradient ascent for 40 steps  
    step = 1.  
    for i in range(40):  
        loss_value, grads_value = iterate([input_img_data])  
        input_img_data += grads_value * step  
  
    img = input_img_data[0]  
    return deprocess_image(img)
```

```
plt.imshow(generate_pattern('block3_conv1', 0))
```

Out[112]: <matplotlib.image.AxesImage at 0x264d06e9bc8>



In [114]:

```
# -----  
# GENERATING A GRID OF ALL FILTER RESPONSE PATTERNS IN A LAYER /  
# -----  
  
layer_name = 'block1_conv1'  
size = 64  
margin = 5  
  
# Empty (black) image to store results  
results = np.zeros((8 * size + 7 * margin, 8 * size + 7 * margin, 3))  
  
for i in range(8): # Iterate over the rows of the results grid  
    for j in range(8): # iterates over the columns of the results grid  
        # Generates the pattern for filter i+(j*8) in layer_name  
        filter_img = generate_pattern(layer_name, i + (j * 8), size=size)  
  
        horizontal_start = i * size + i * margin  
        horizontal_end = horizontal_start + size  
        vertical_start = j * size + j * margin  
        vertical_end = vertical_start + size  
        results[horizontal_start: horizontal_end, vertical_start: vertical_end, :] = filter_img  
  
# Displays the results grid  
plt.figure(figsize=(20, 20))  
plt.imshow(results)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

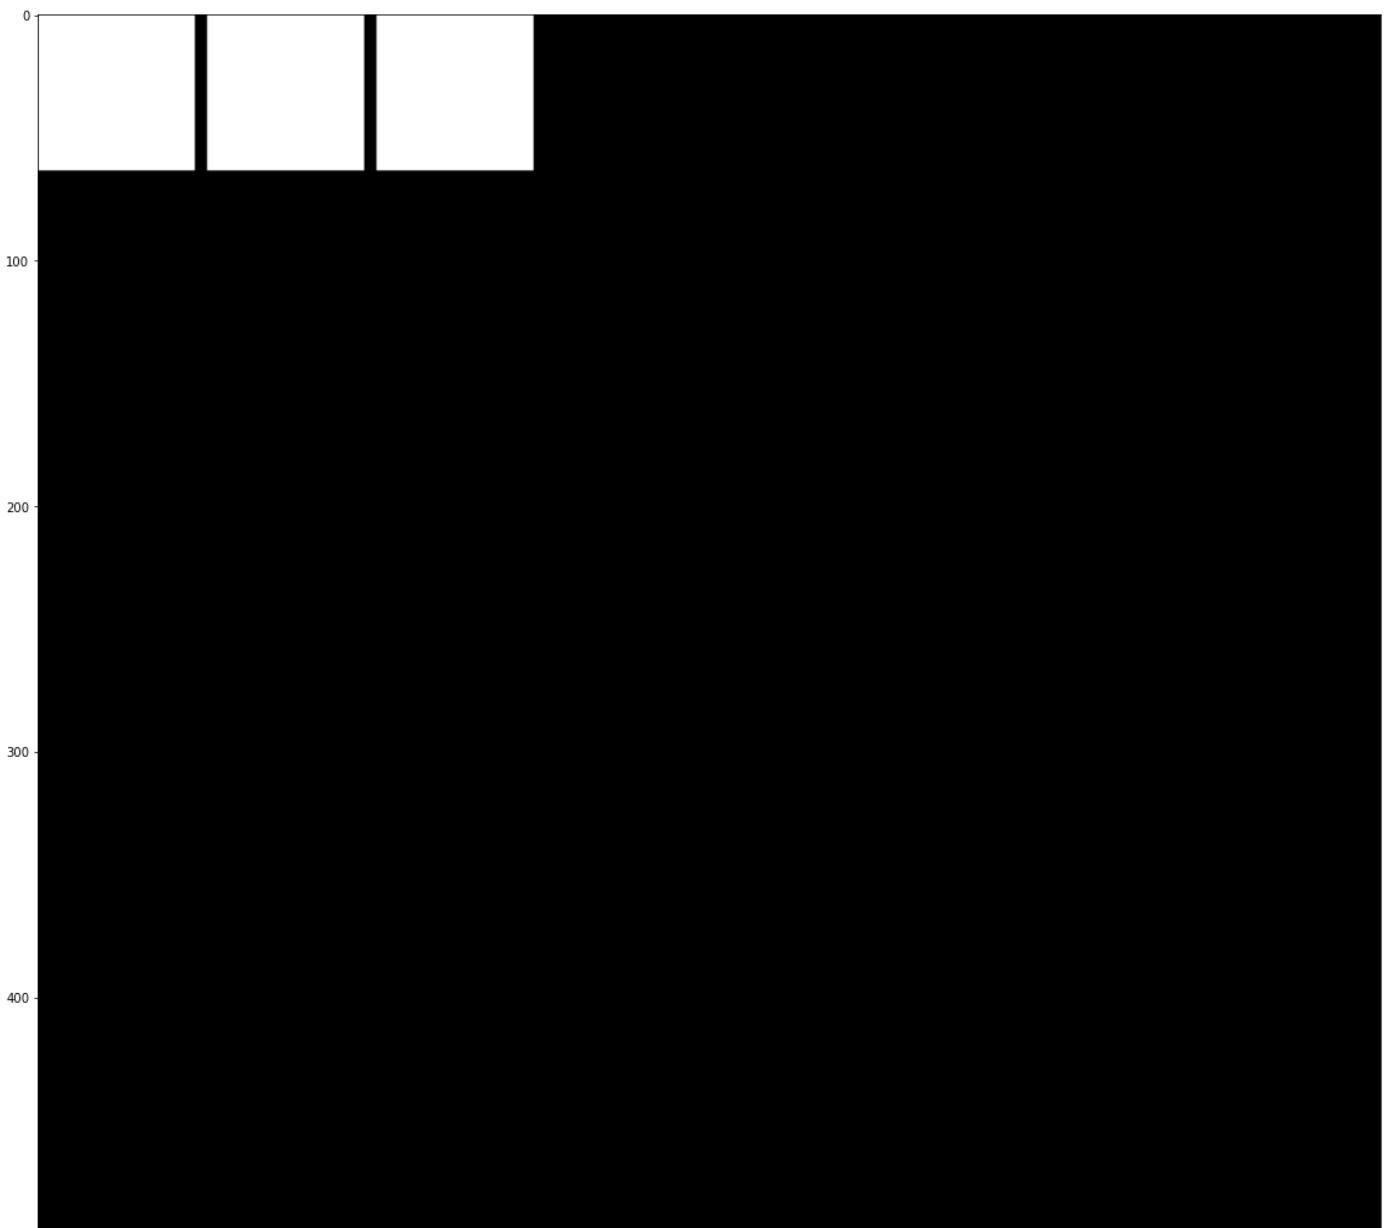
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

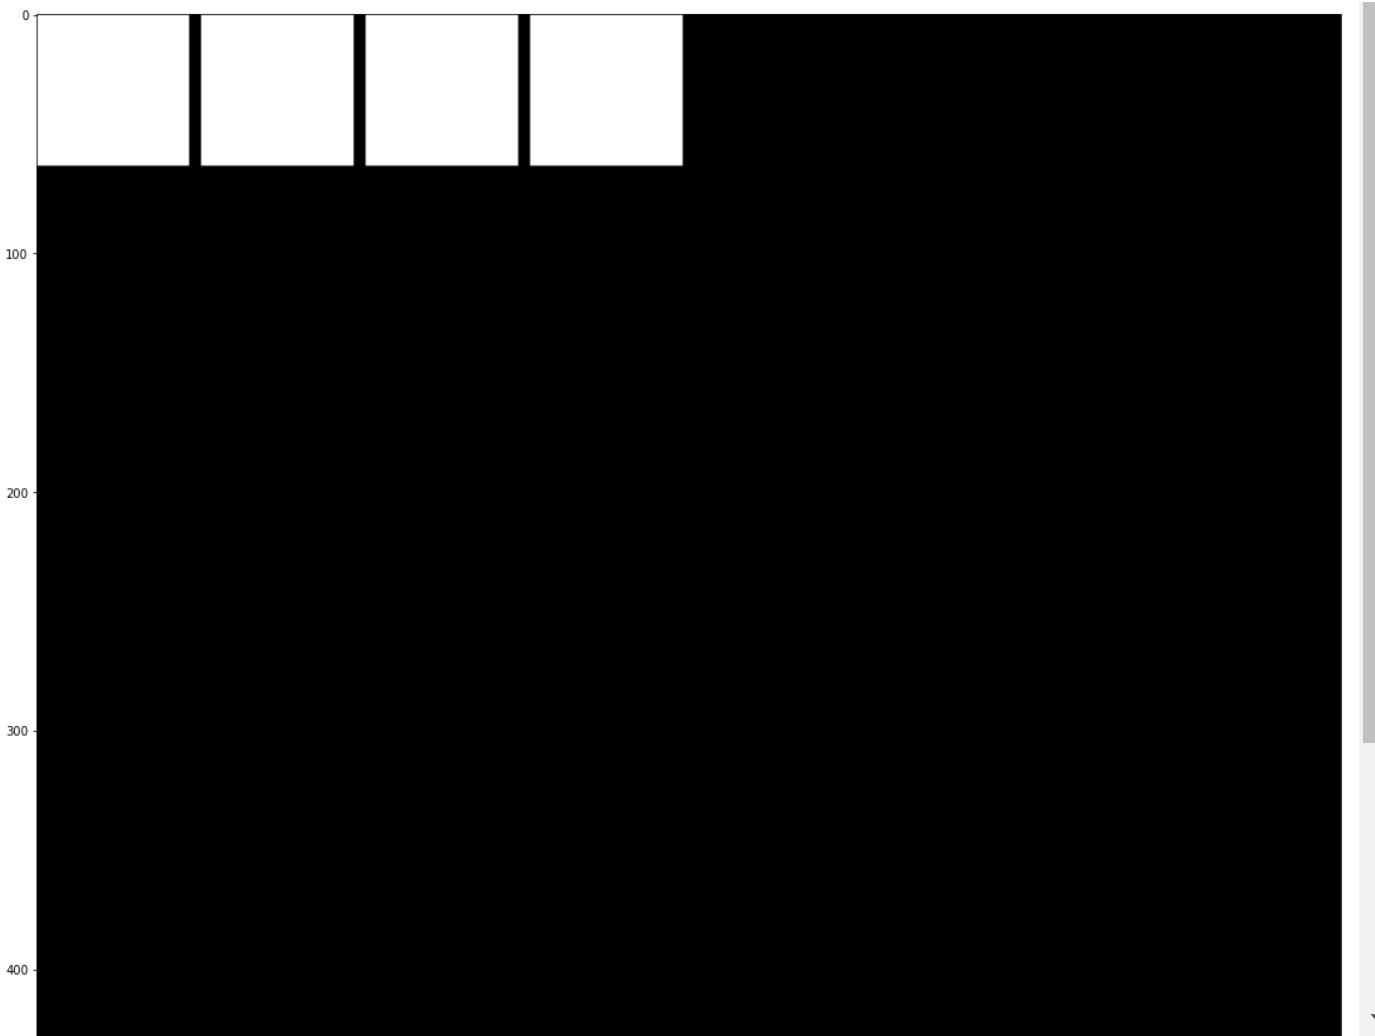
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



▼
▲

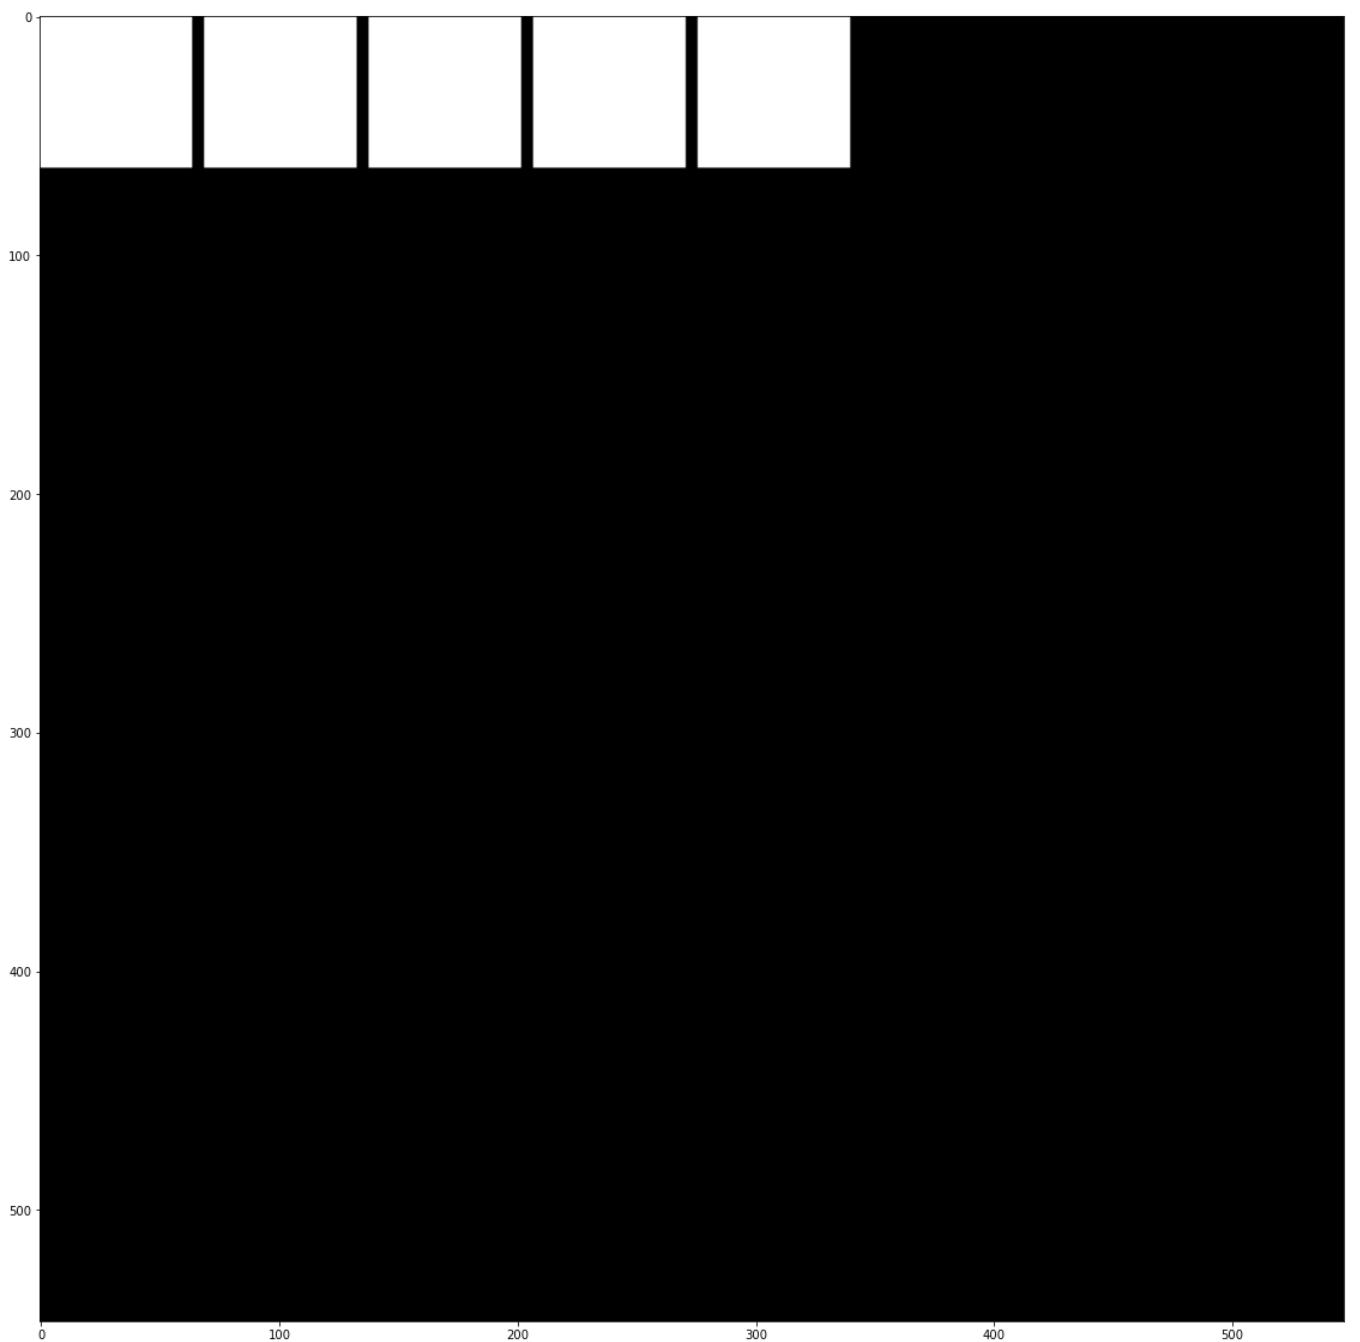




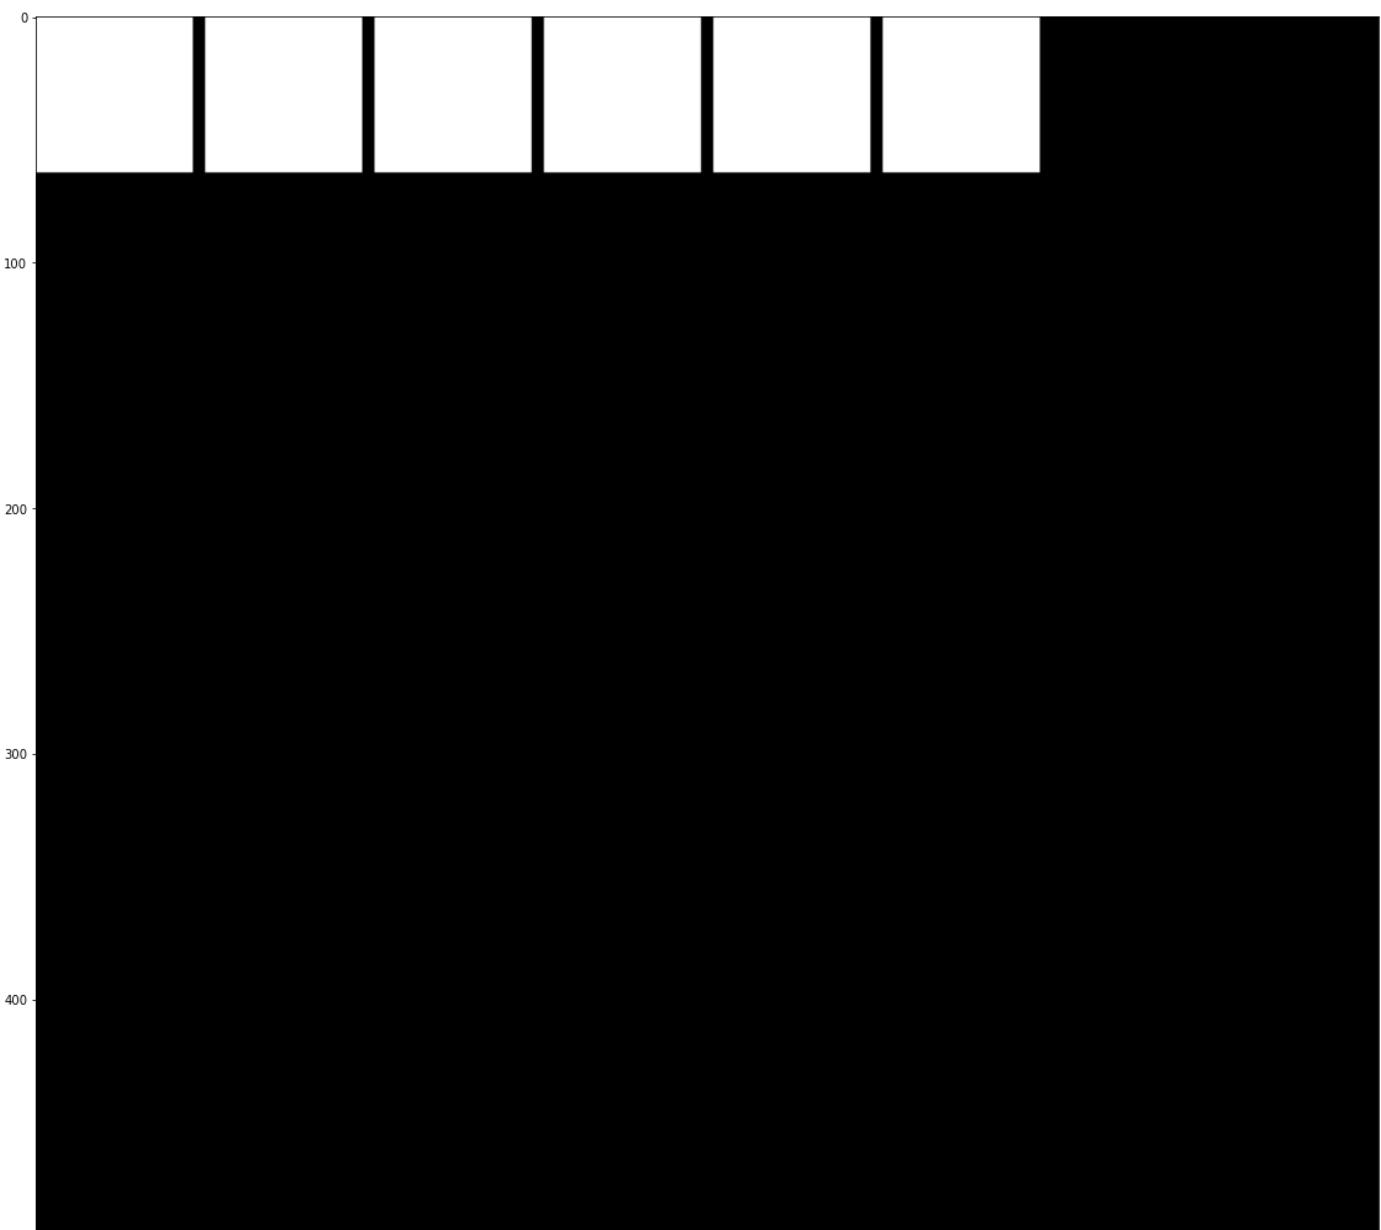


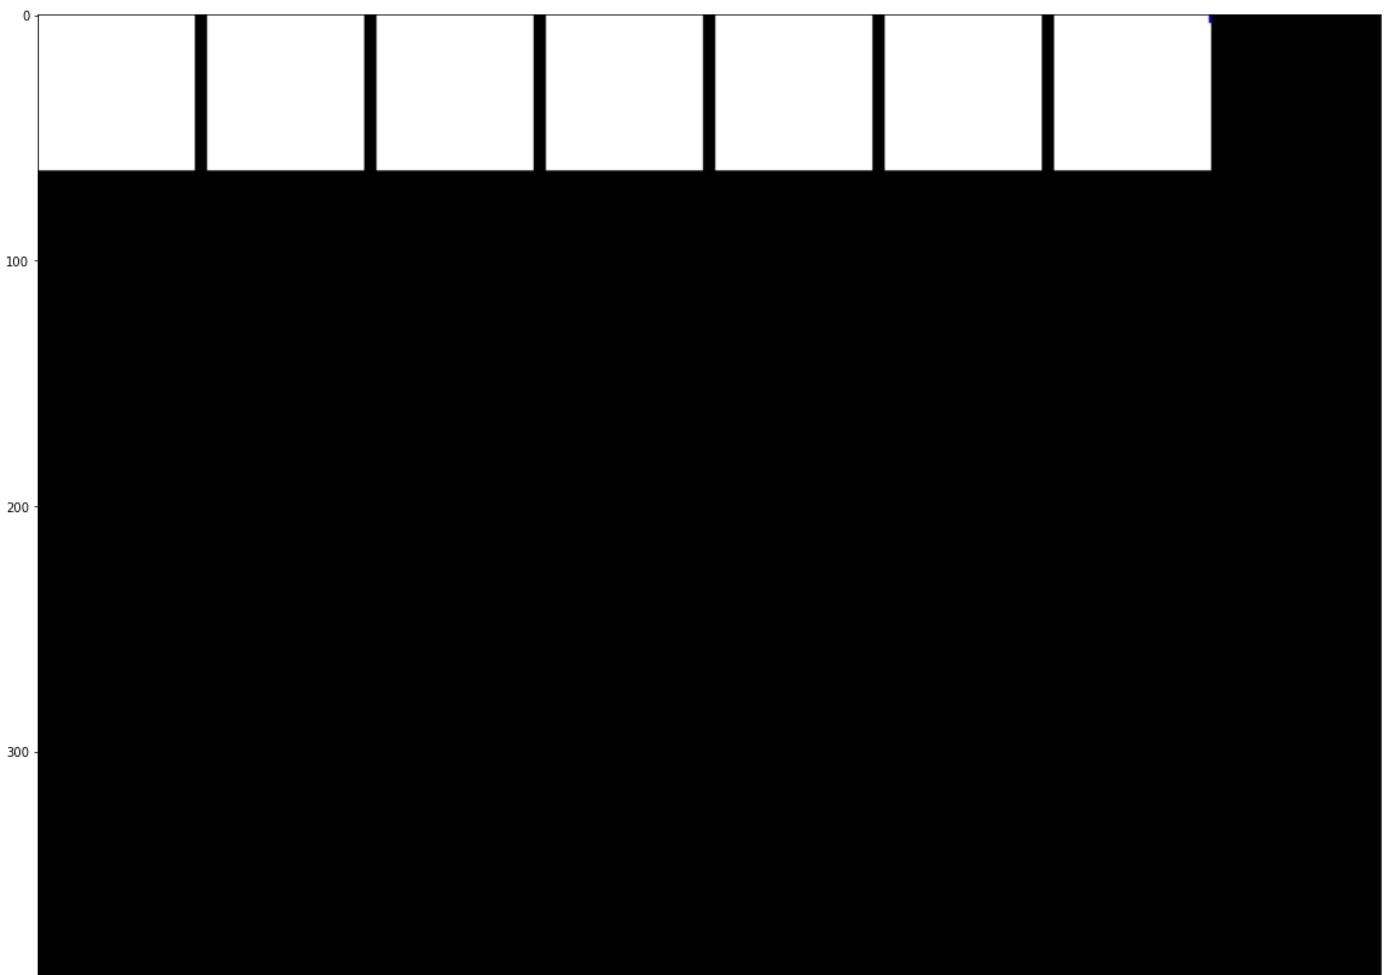
▼

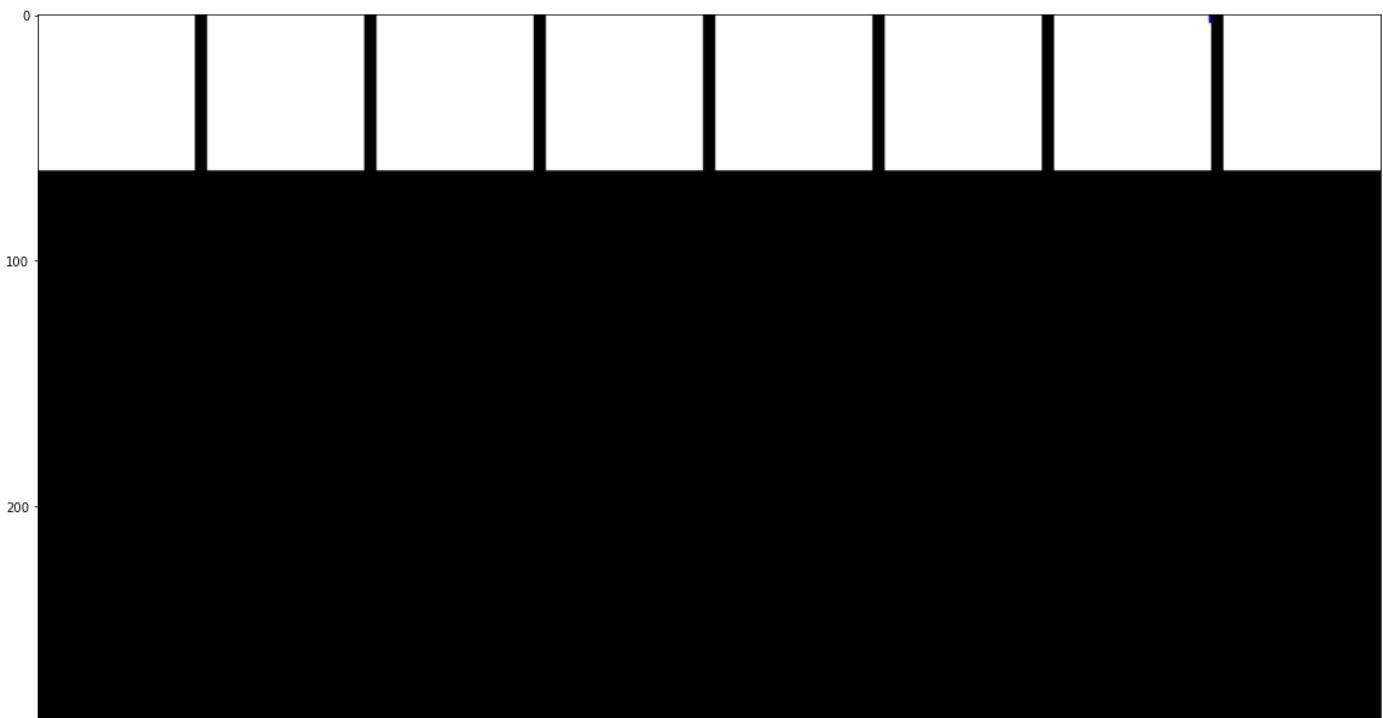
▲

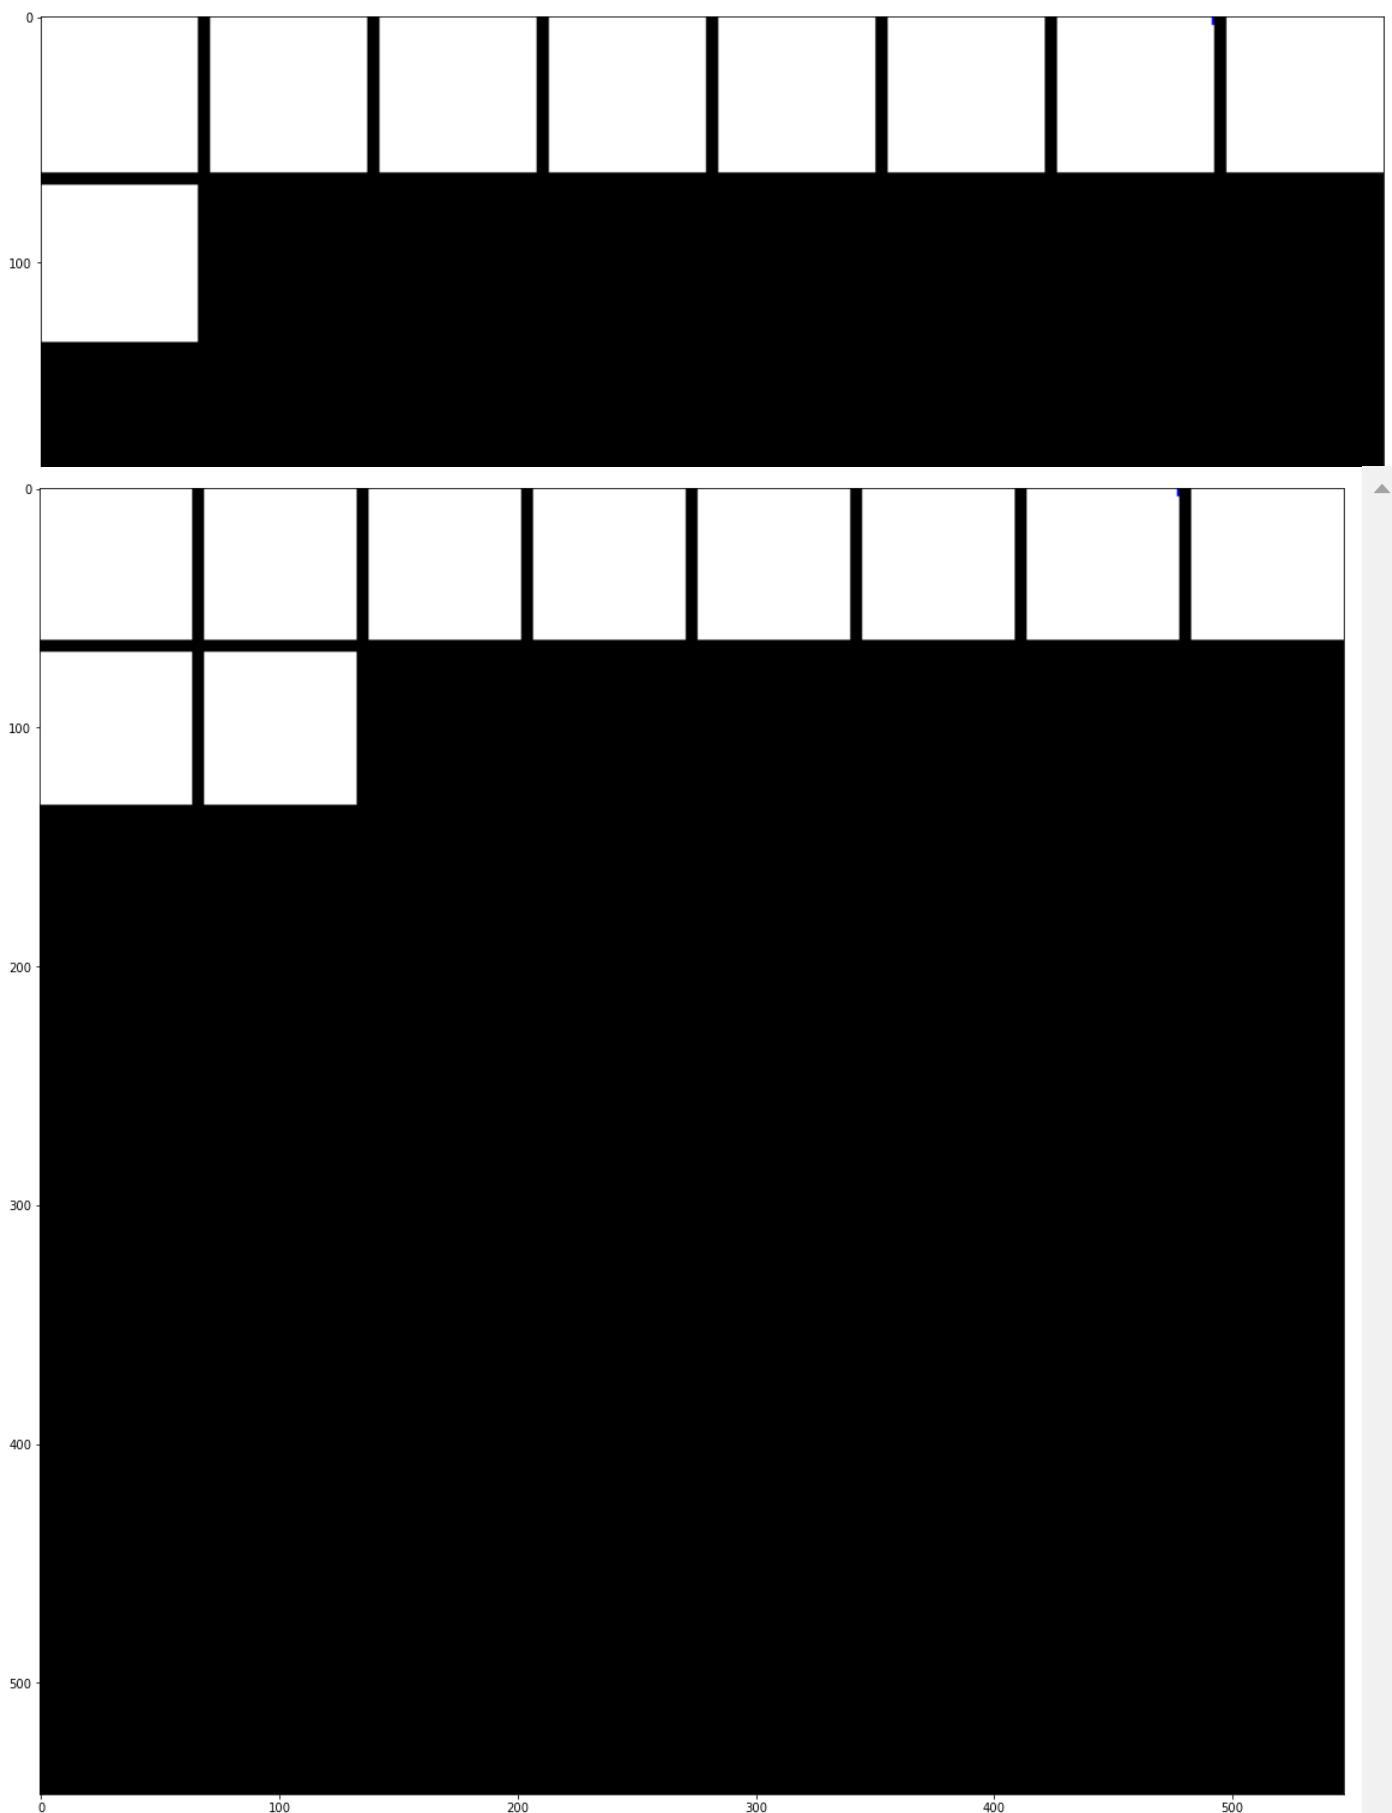


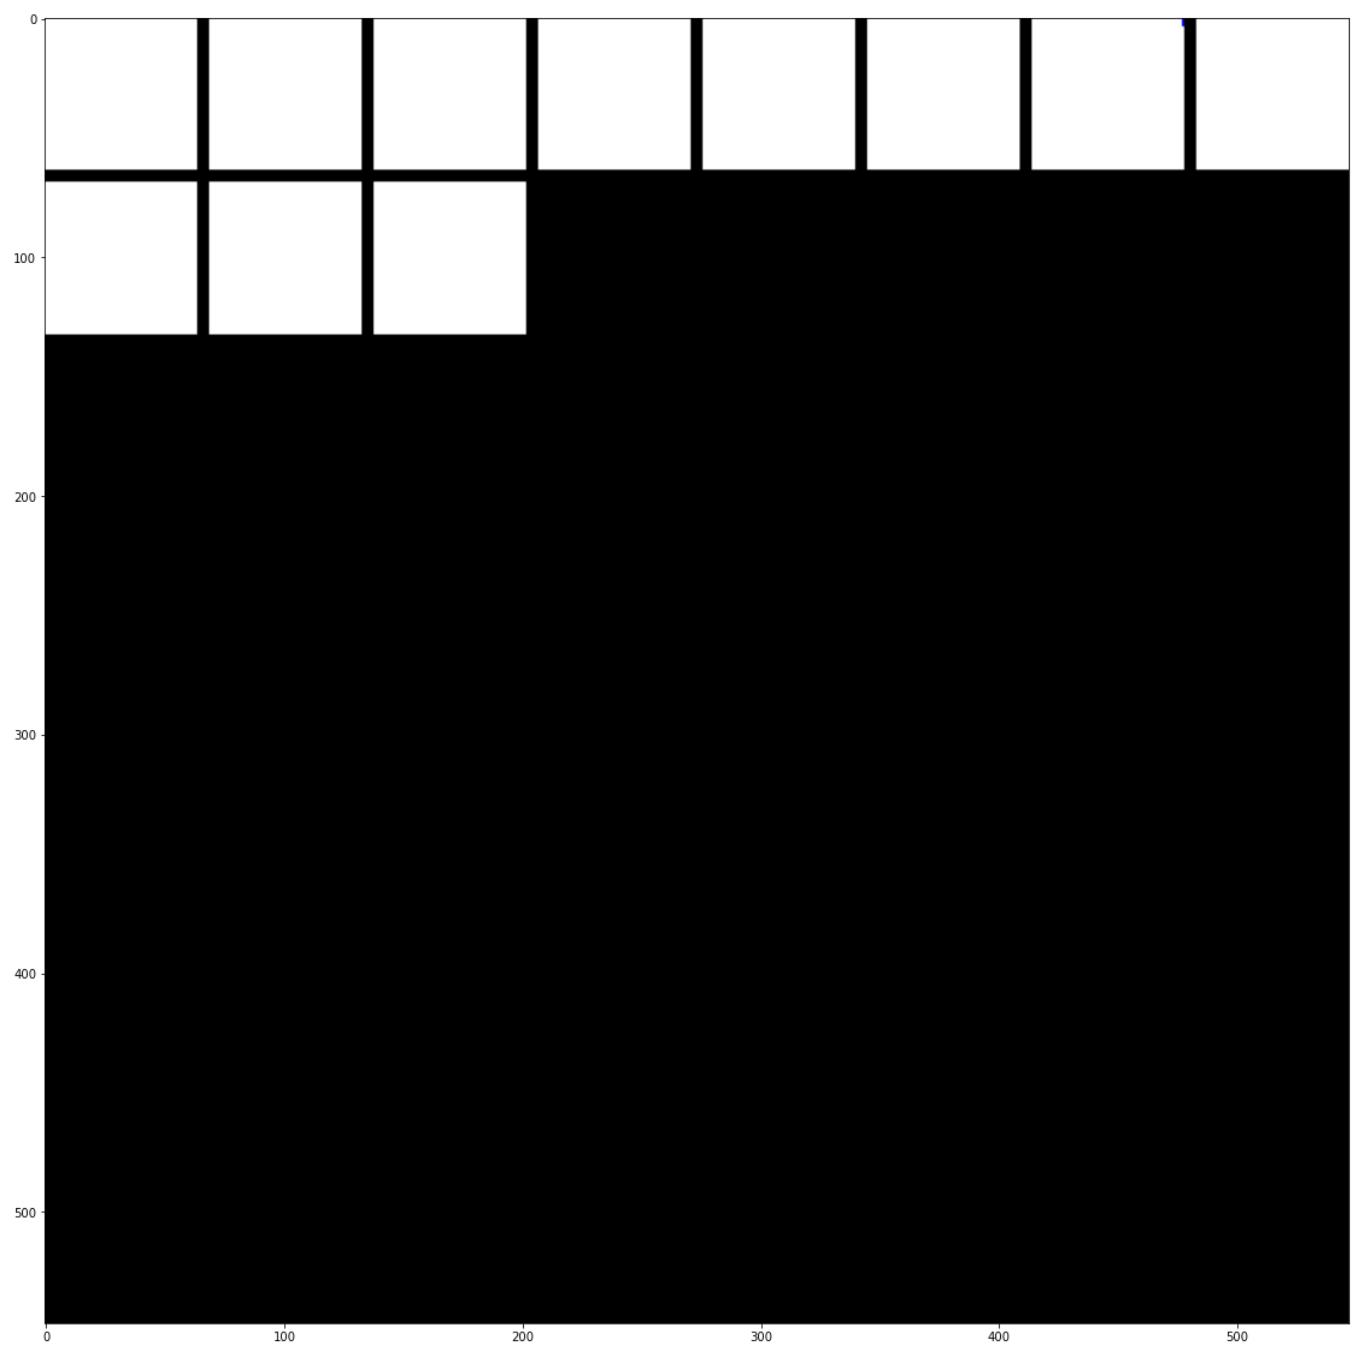
▼

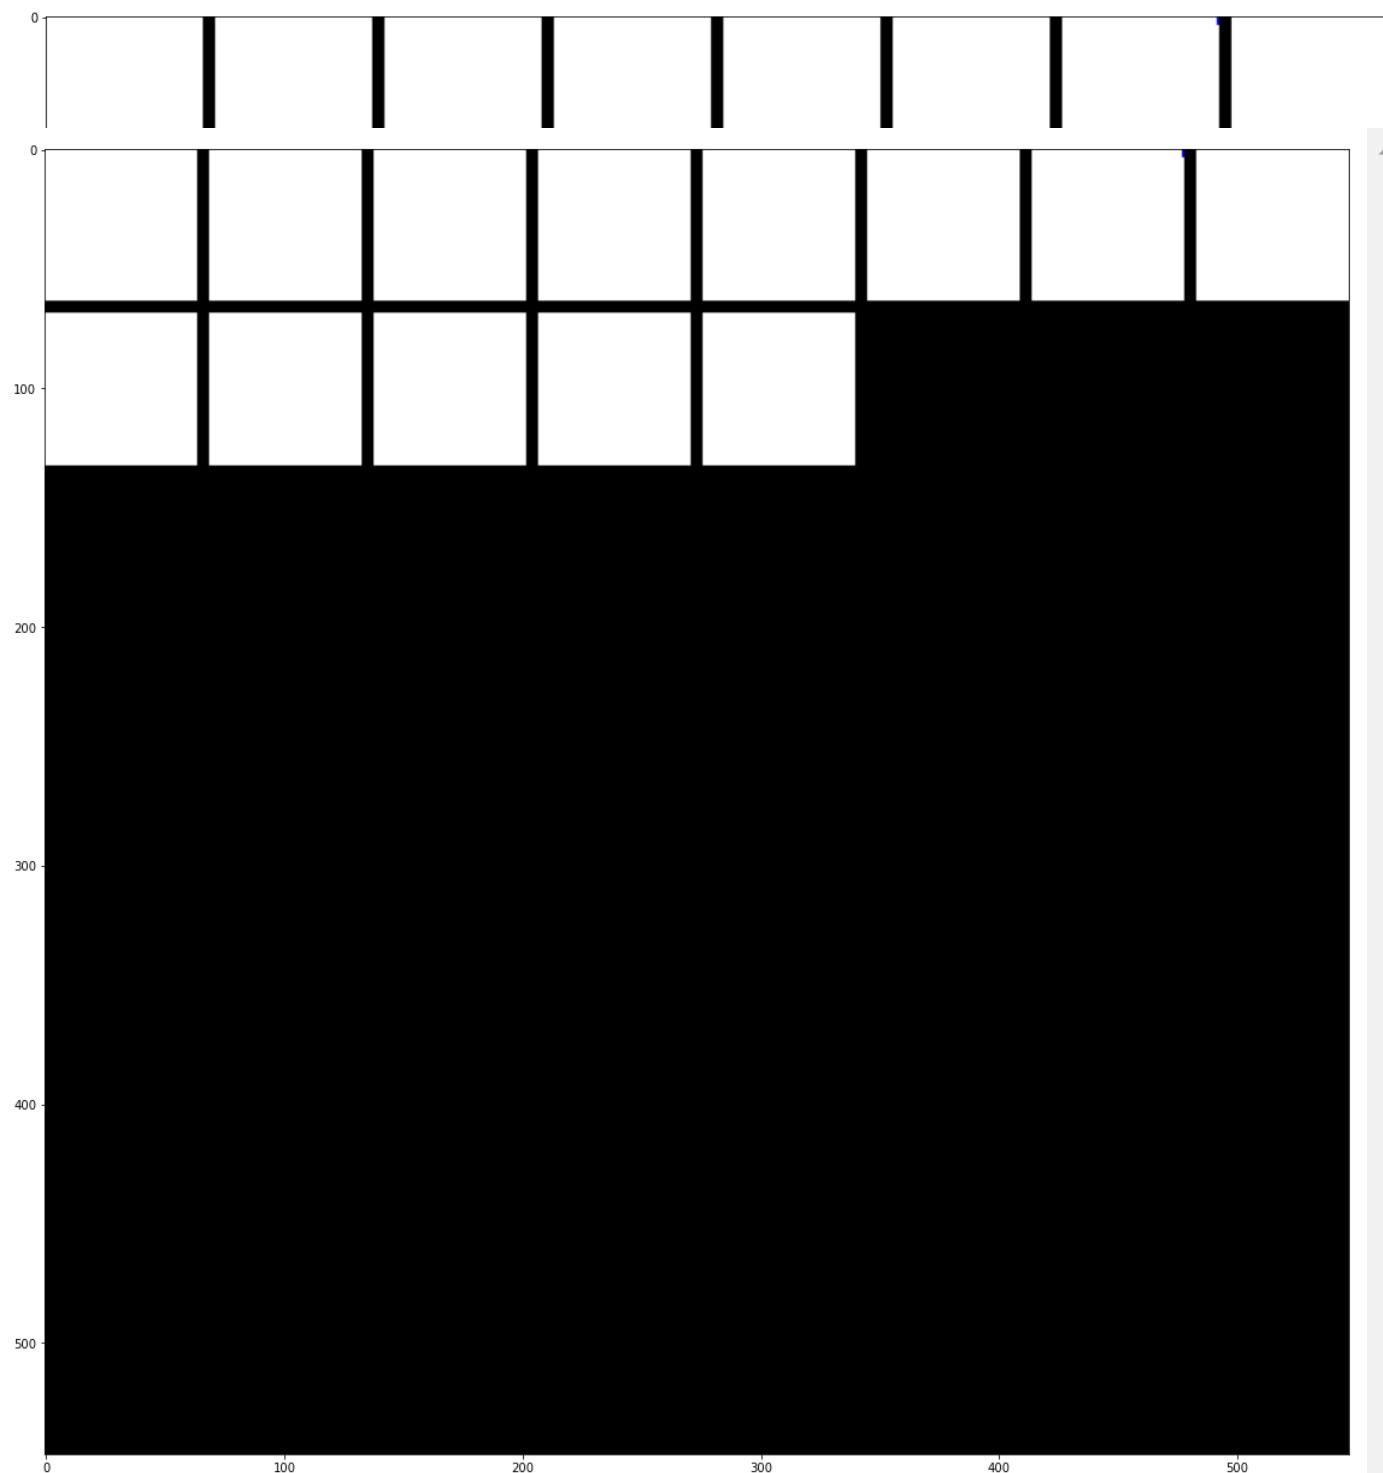


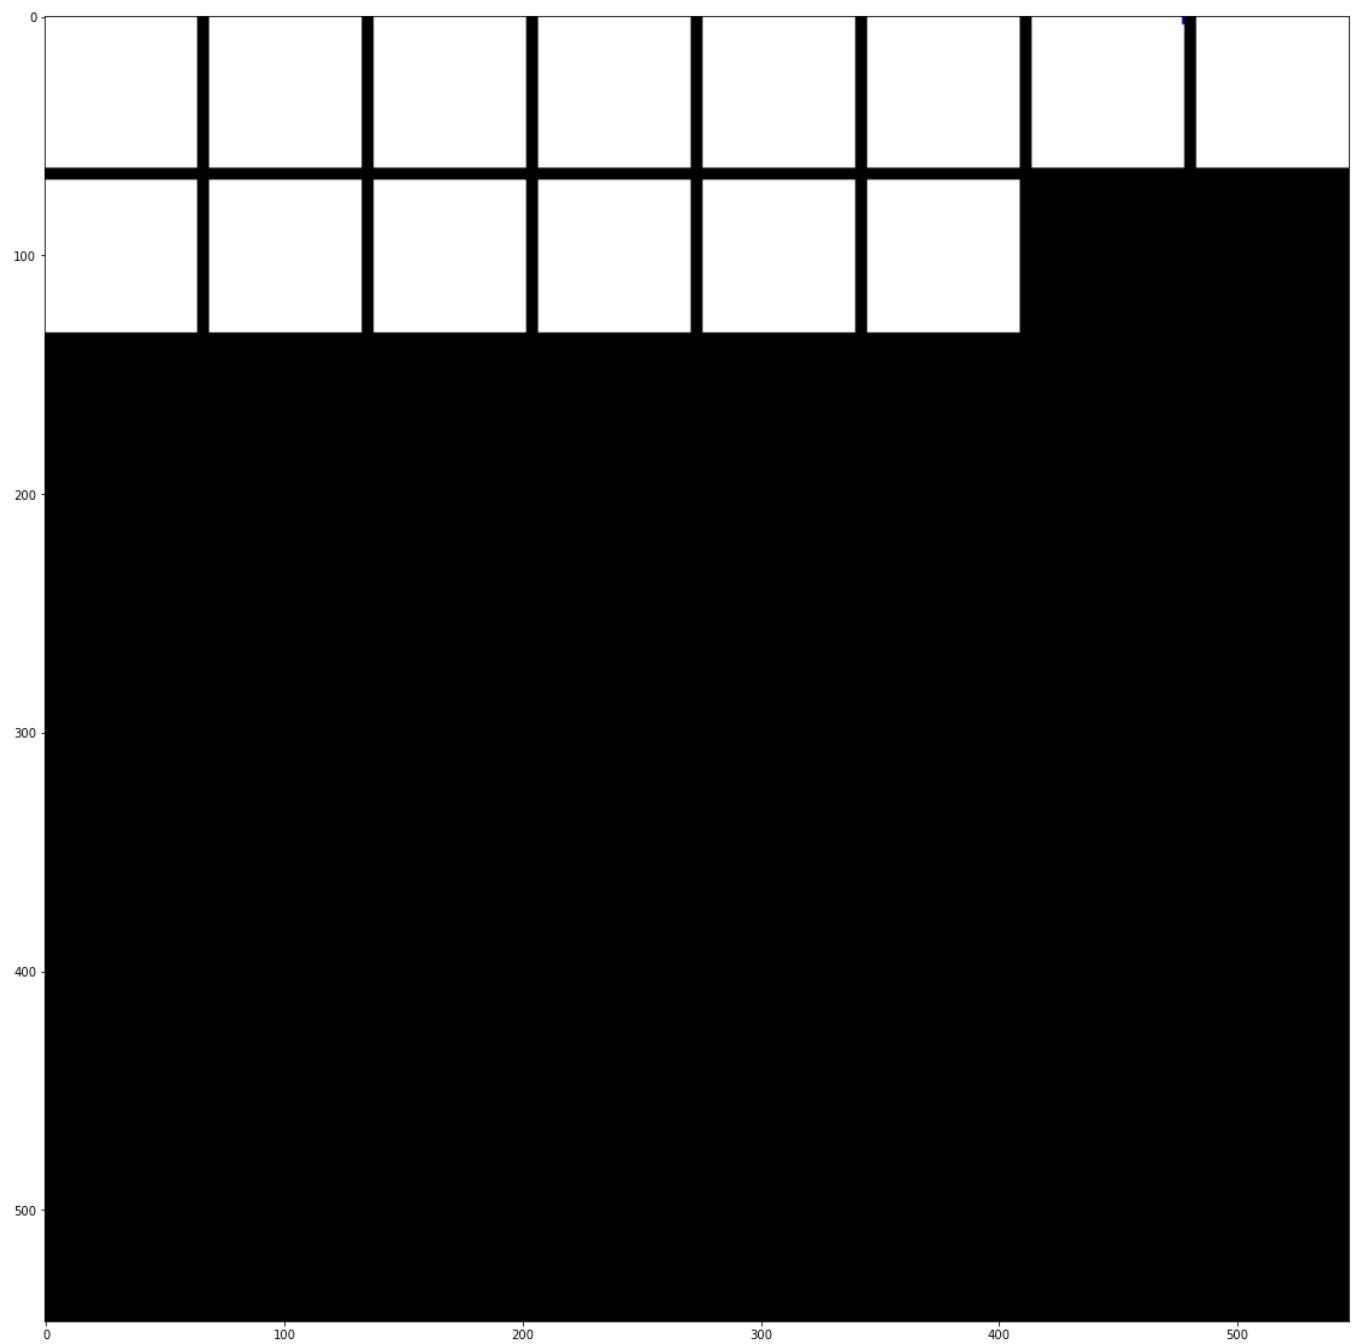


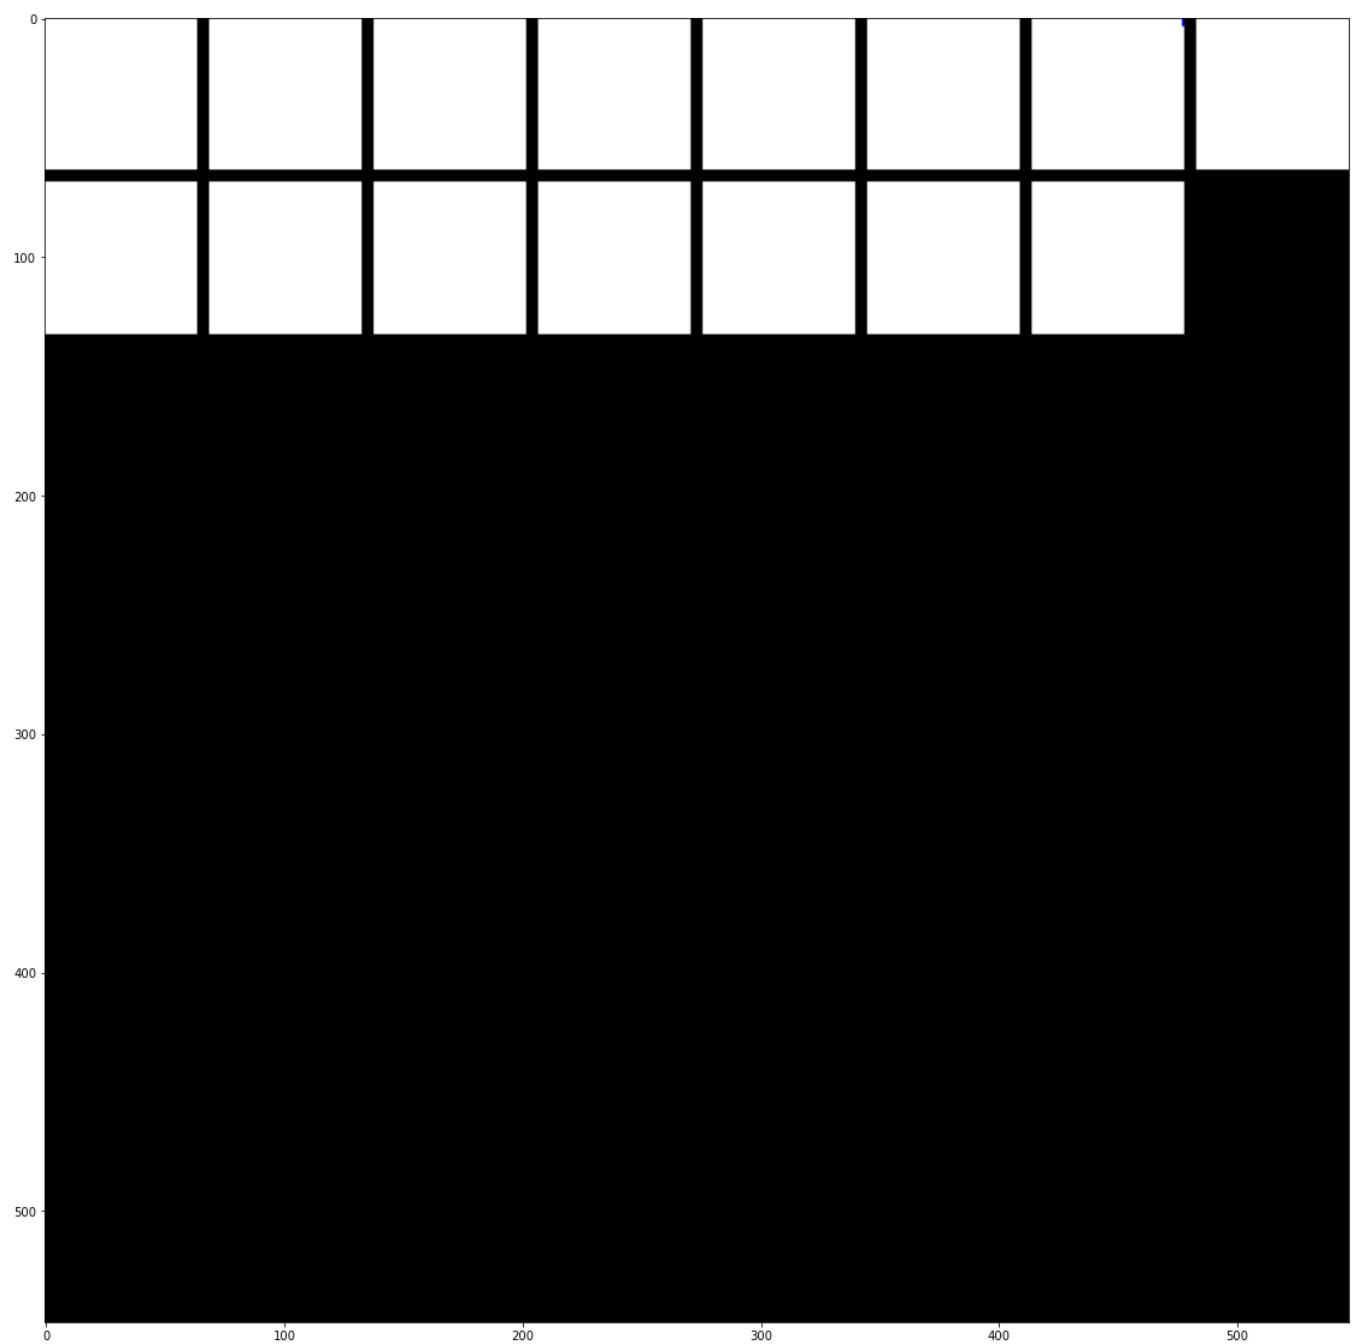


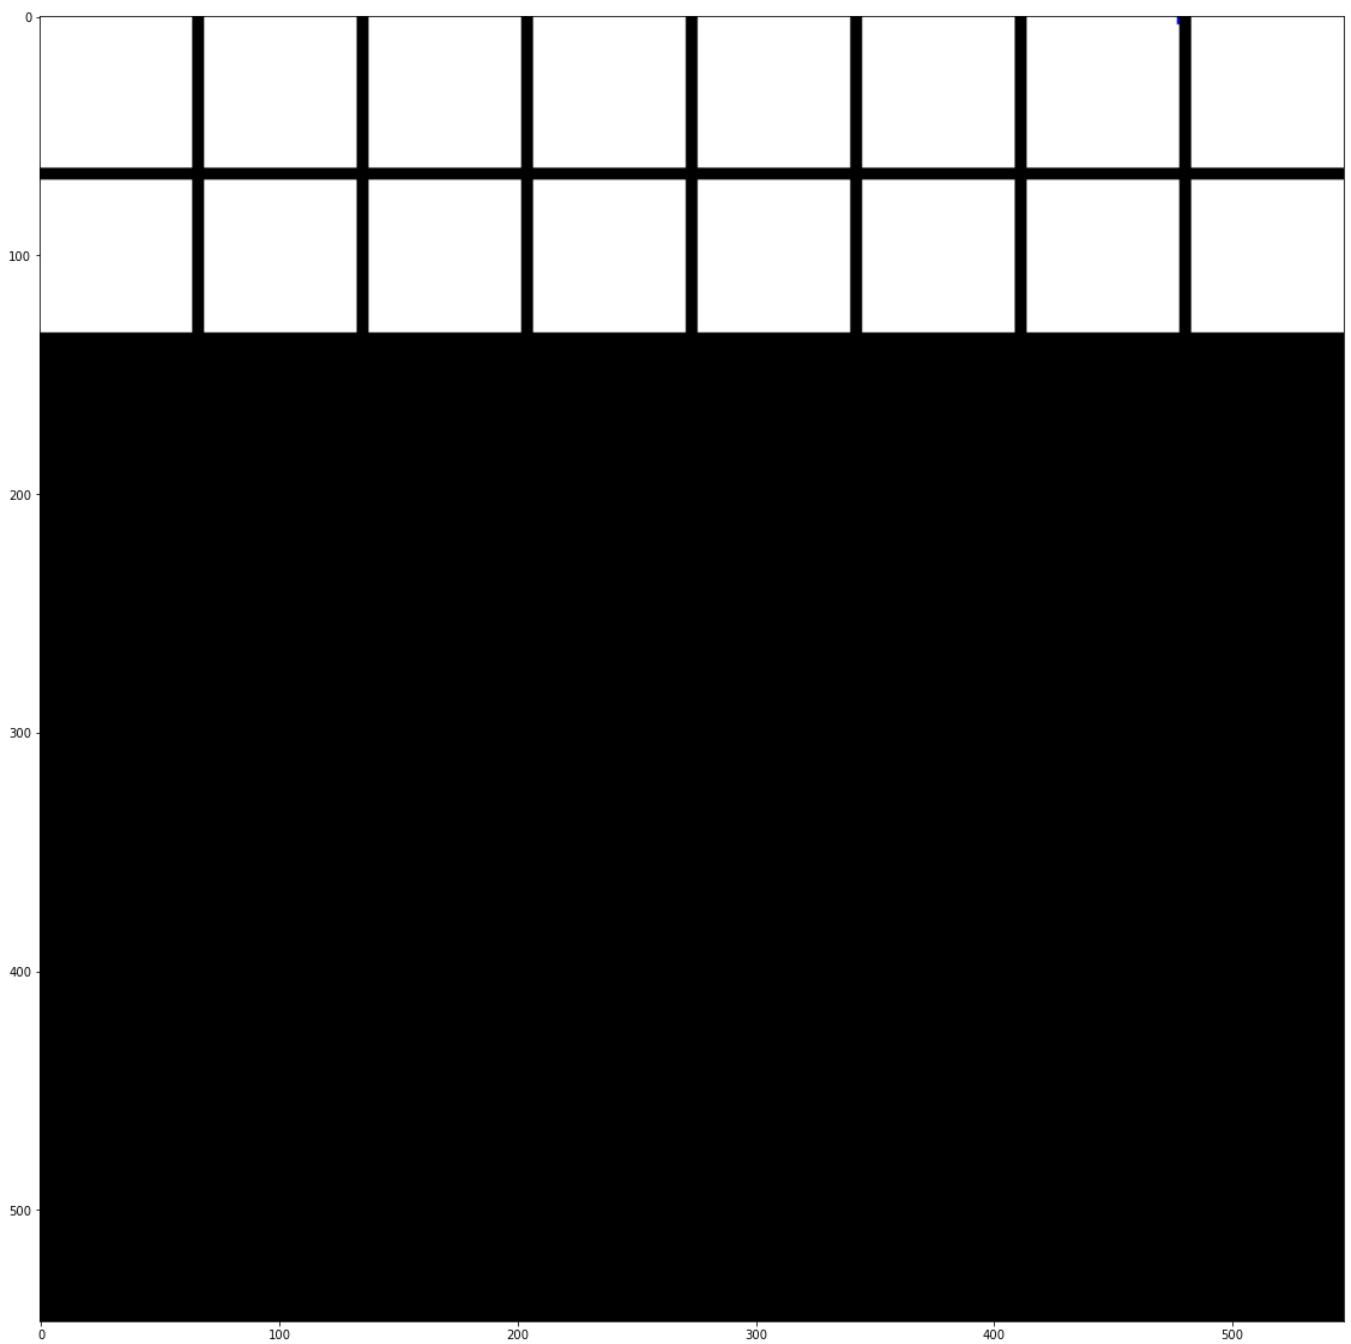


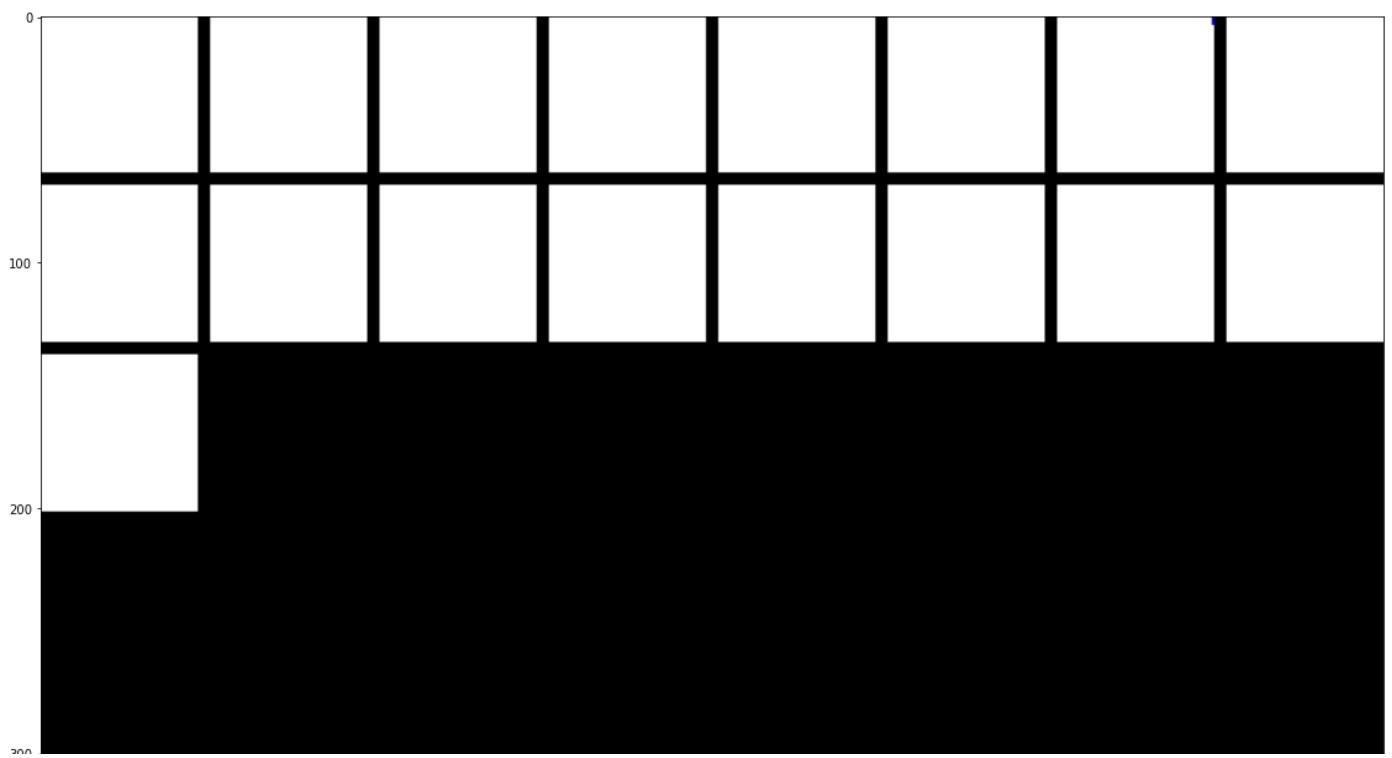


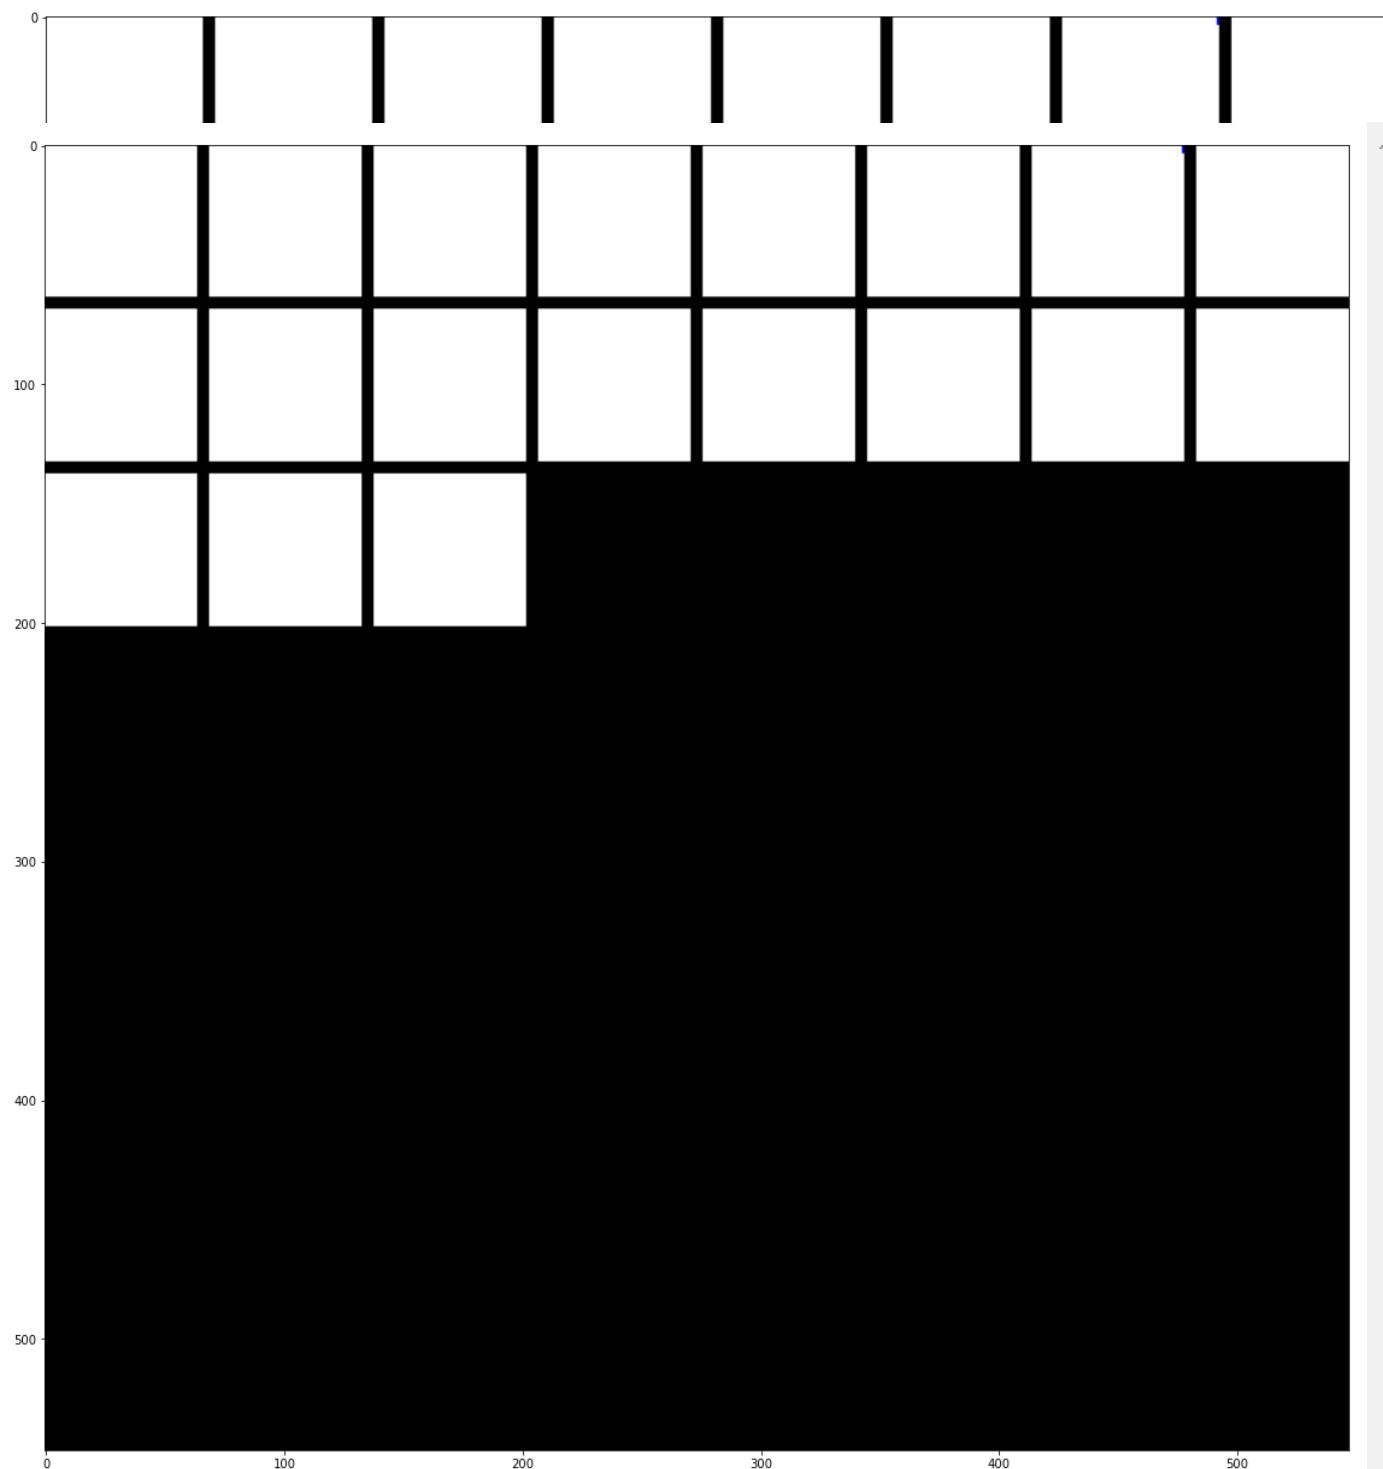


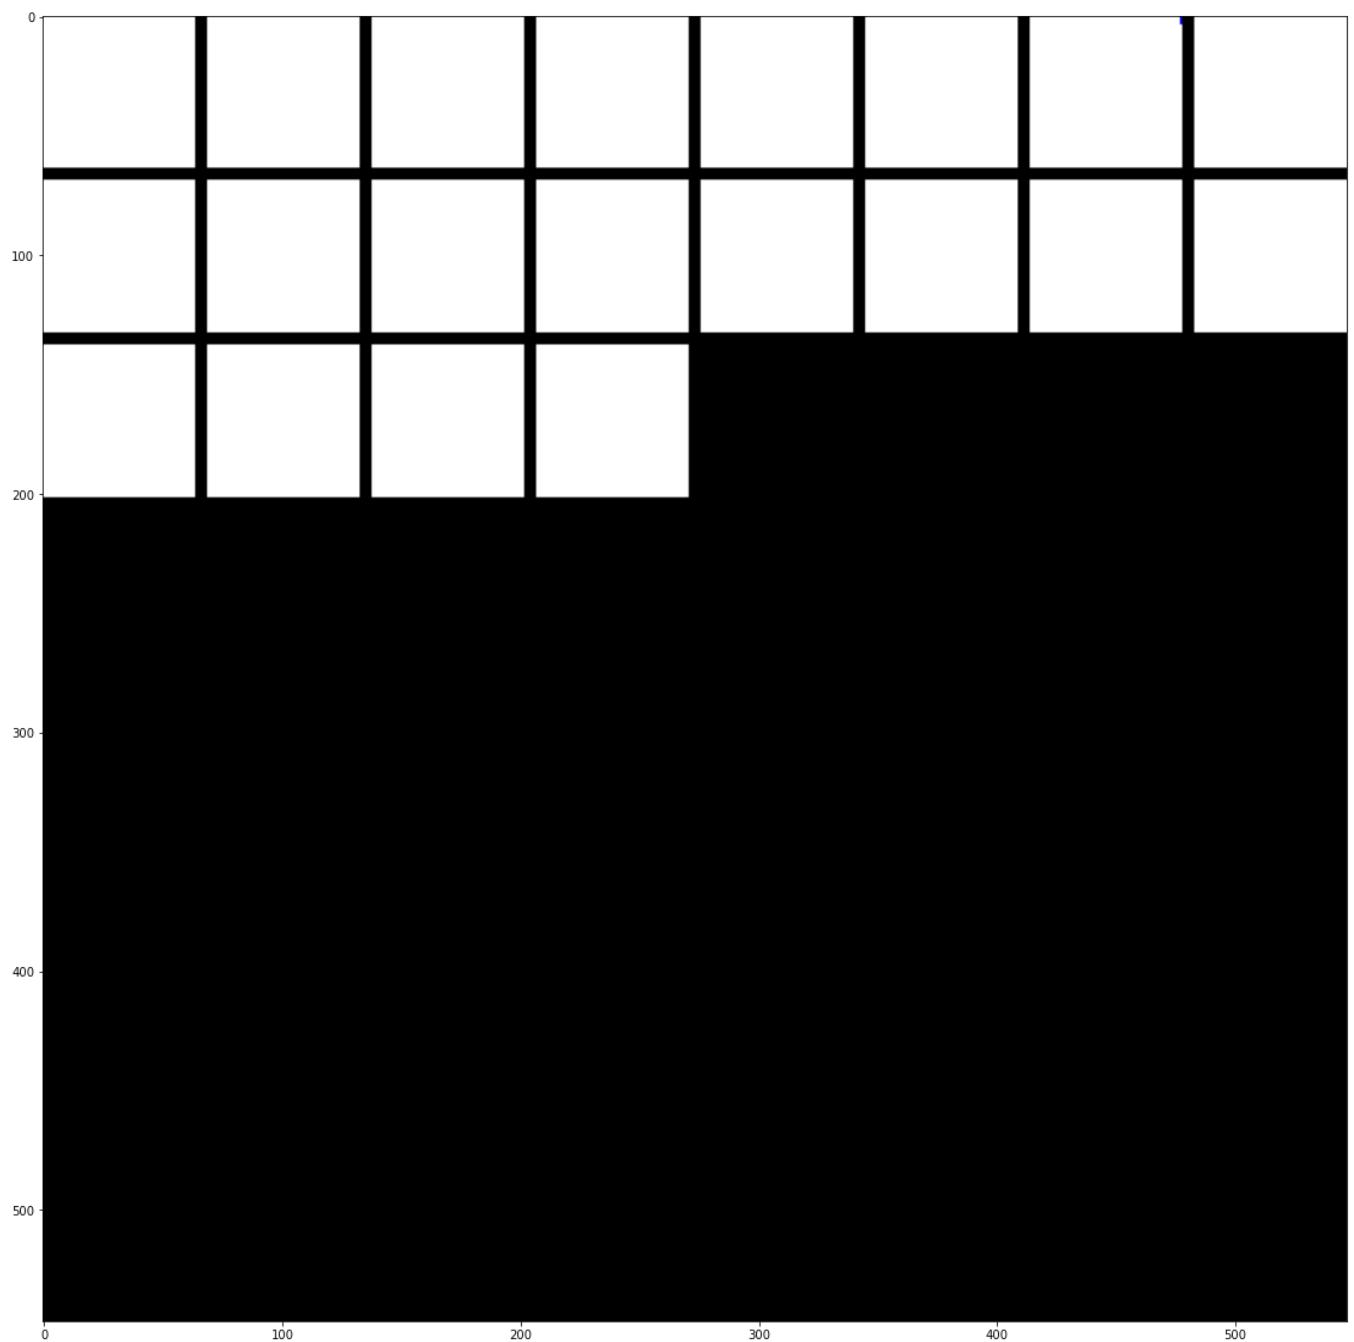


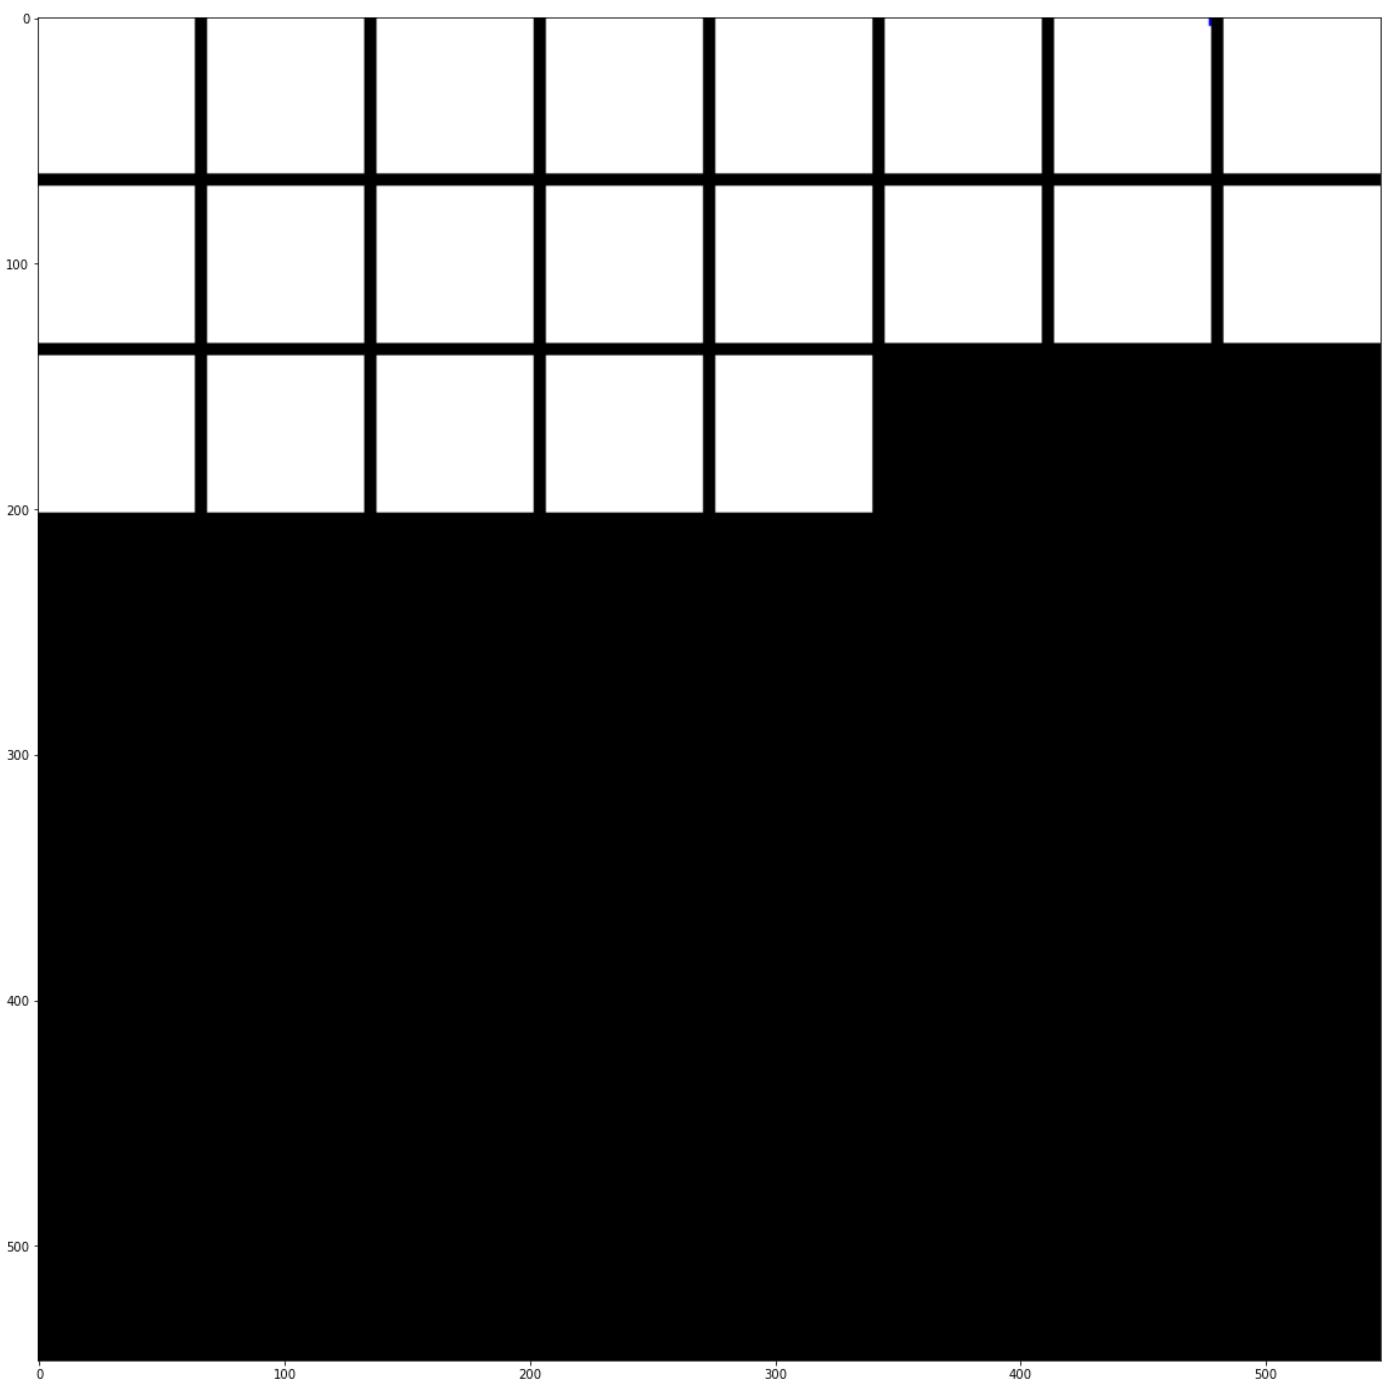


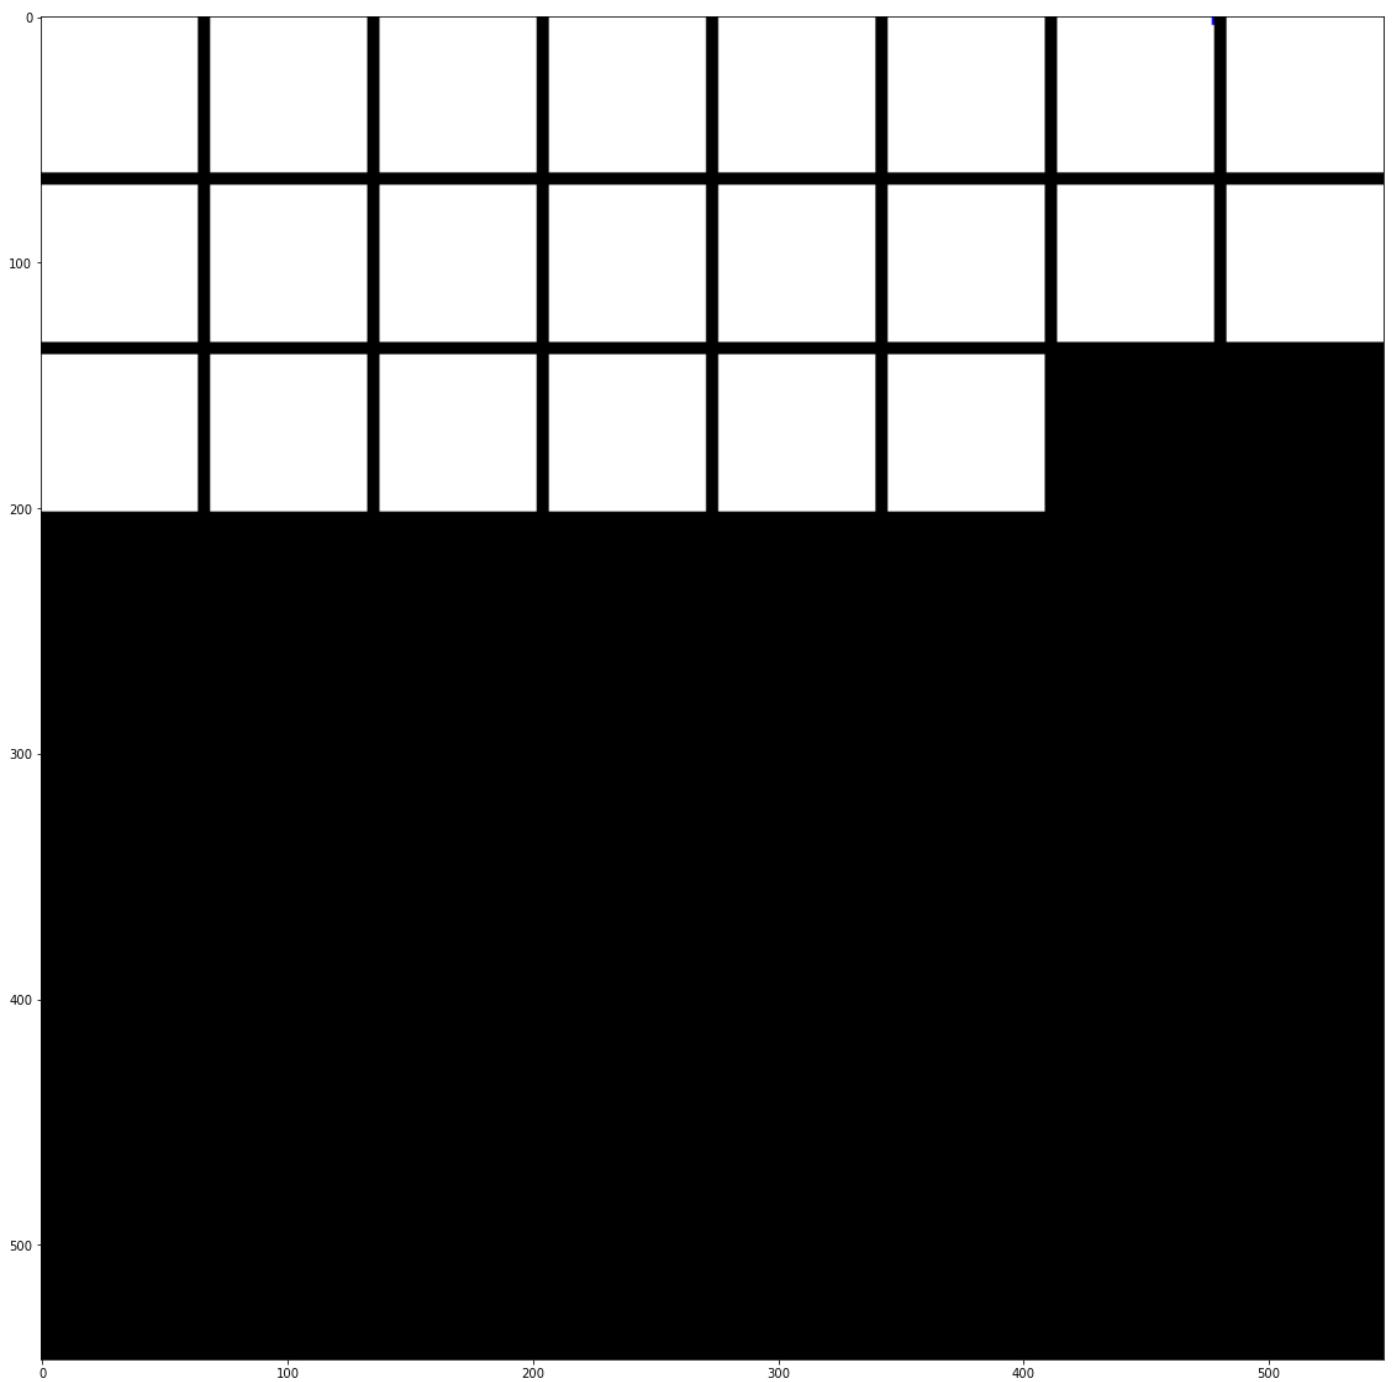


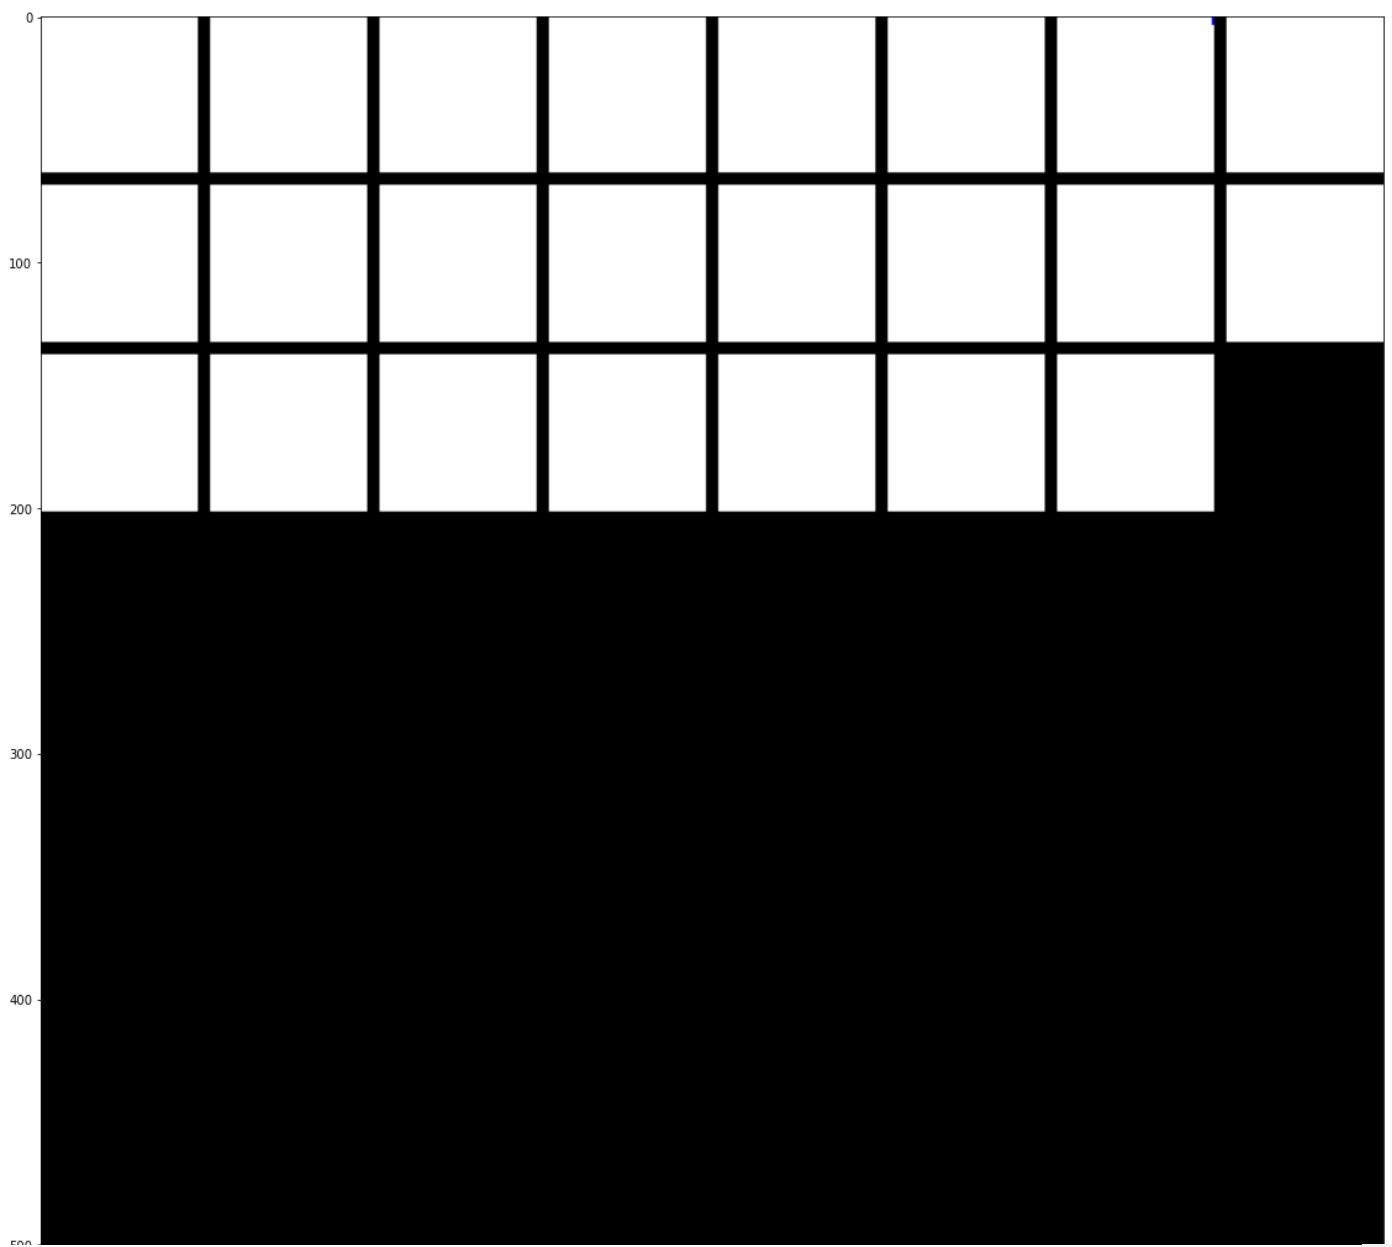


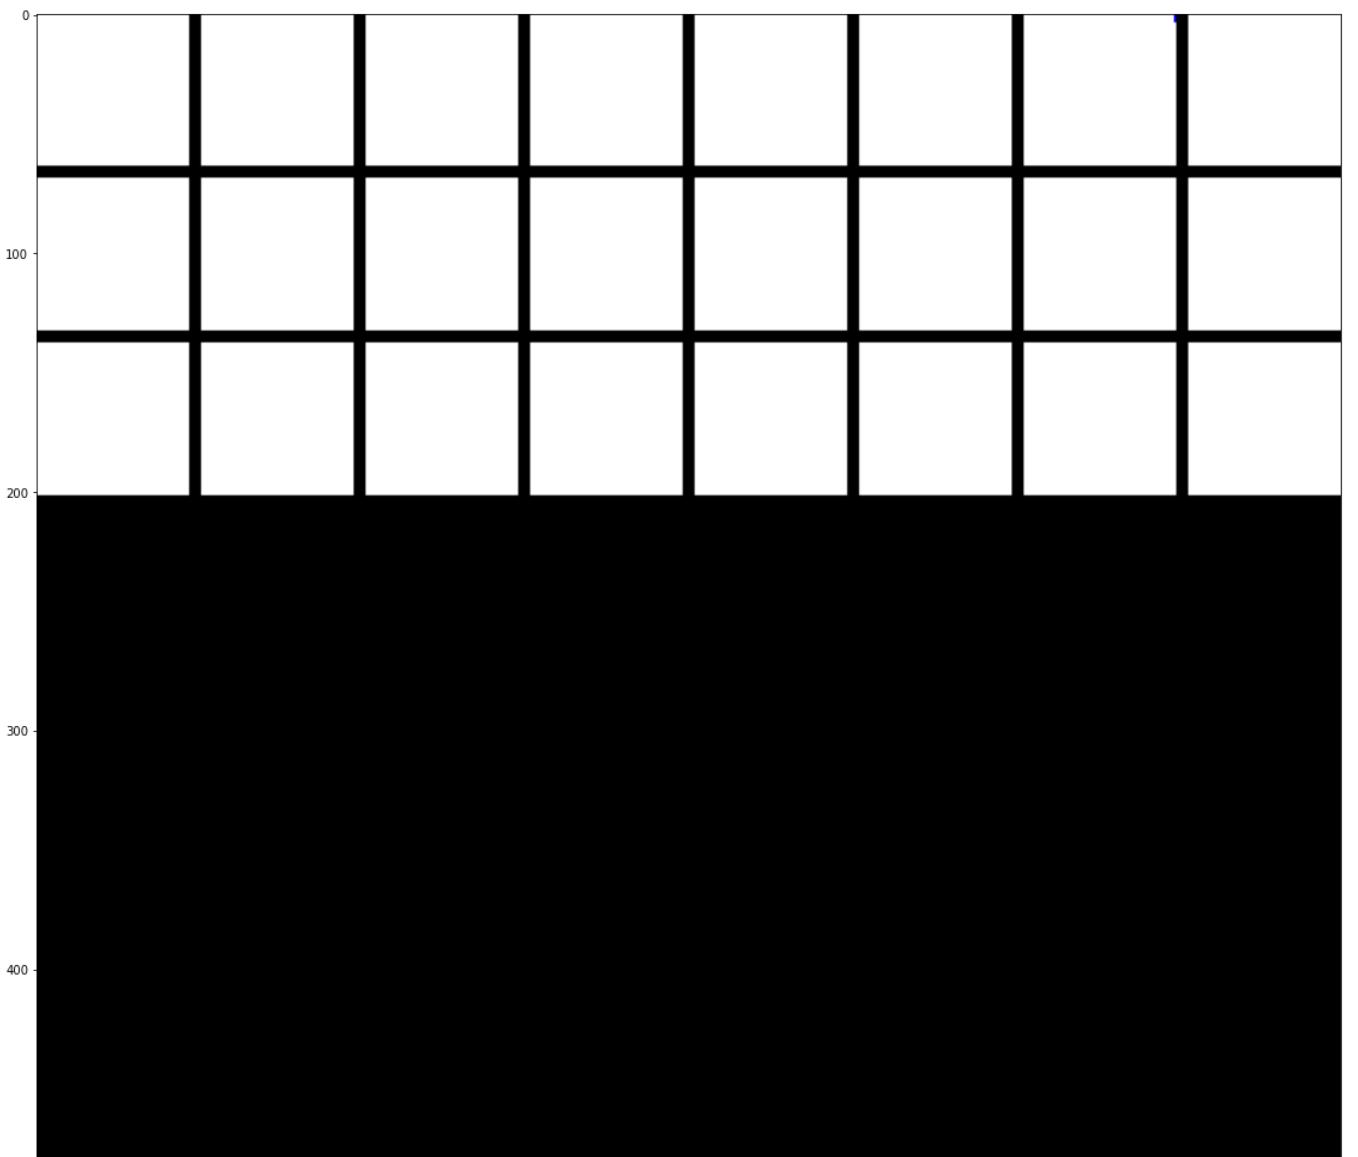


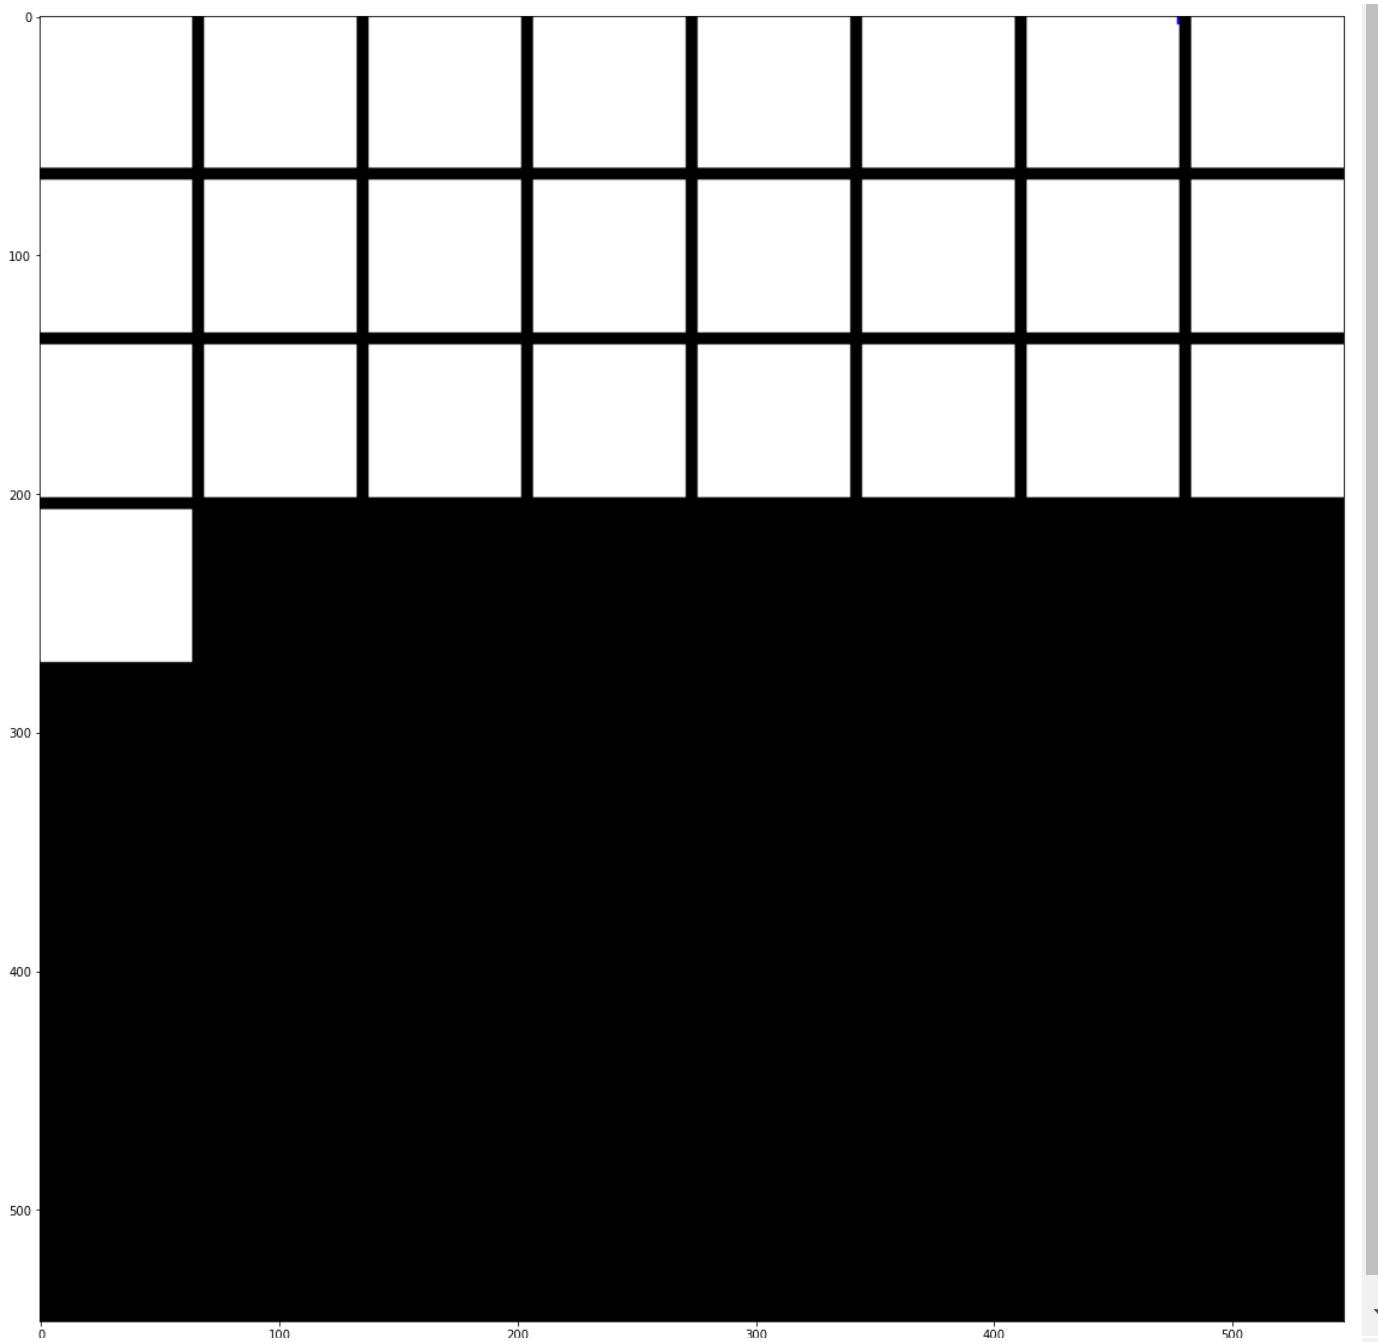






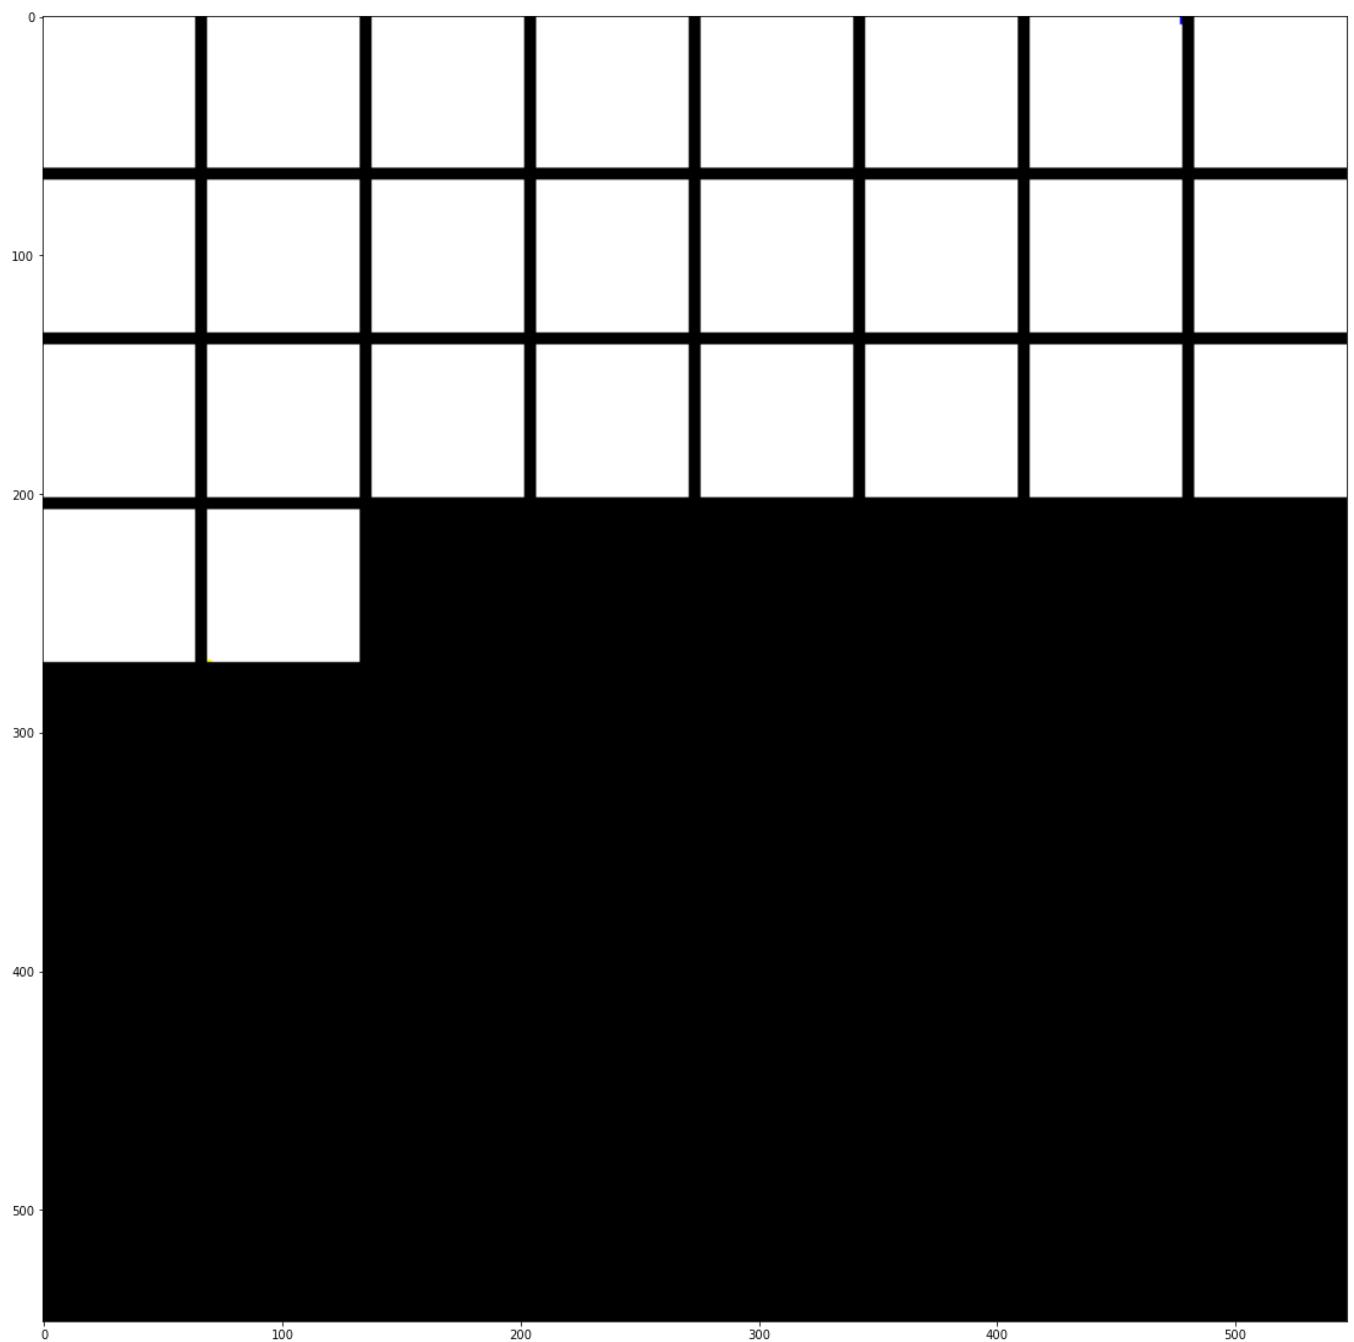


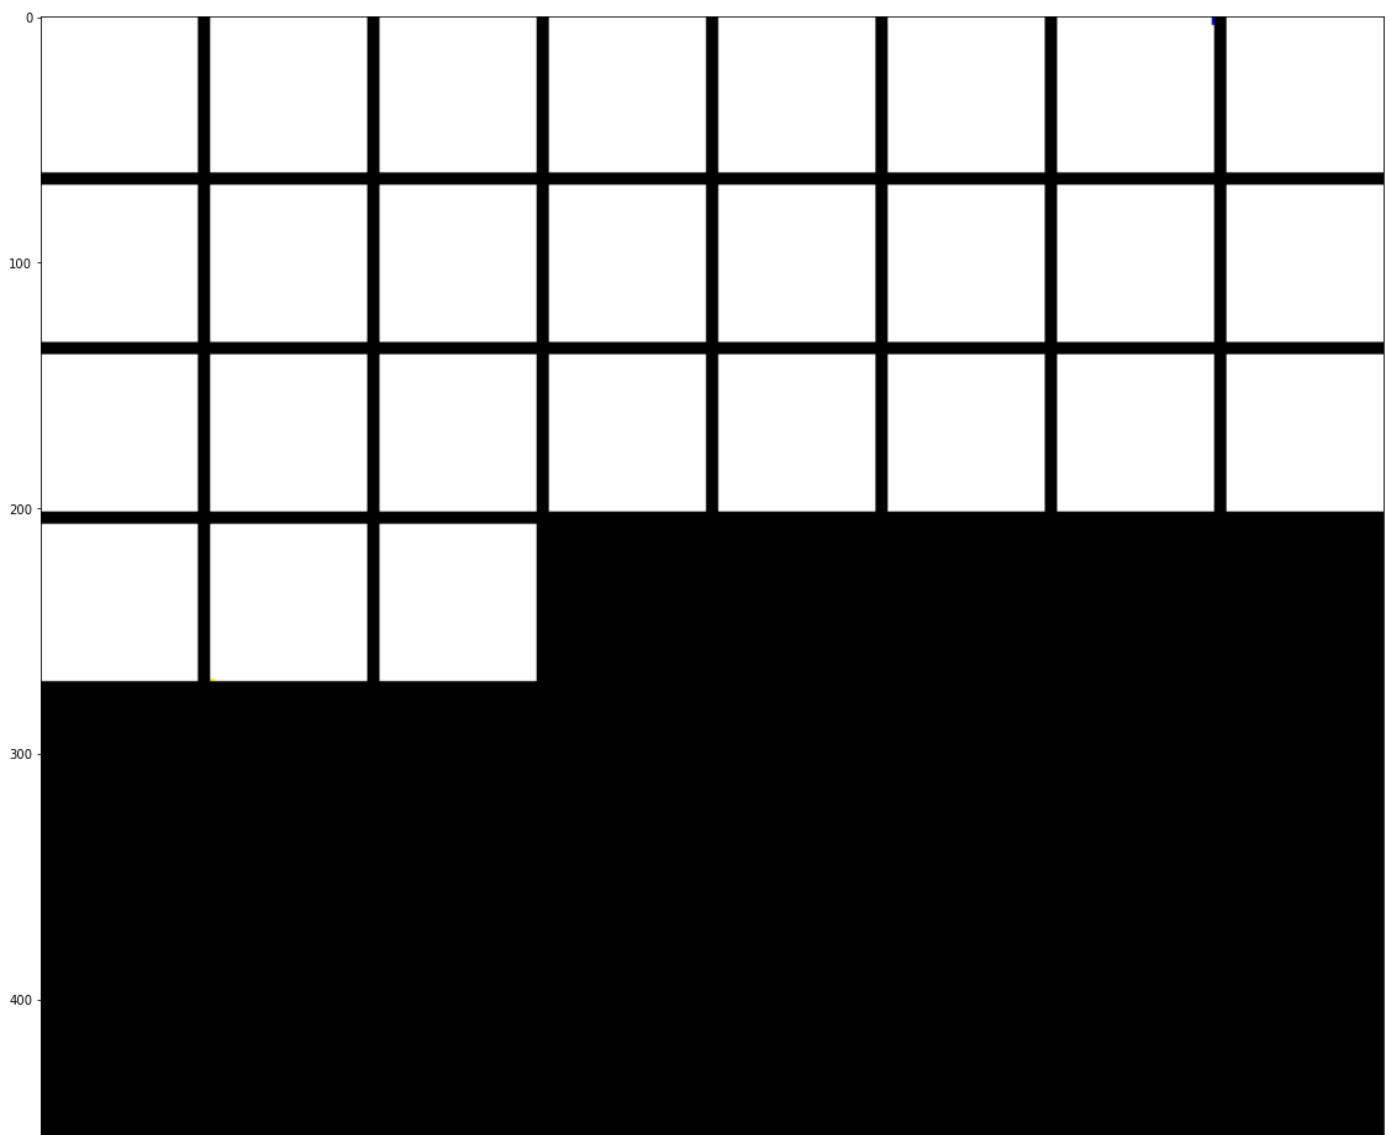




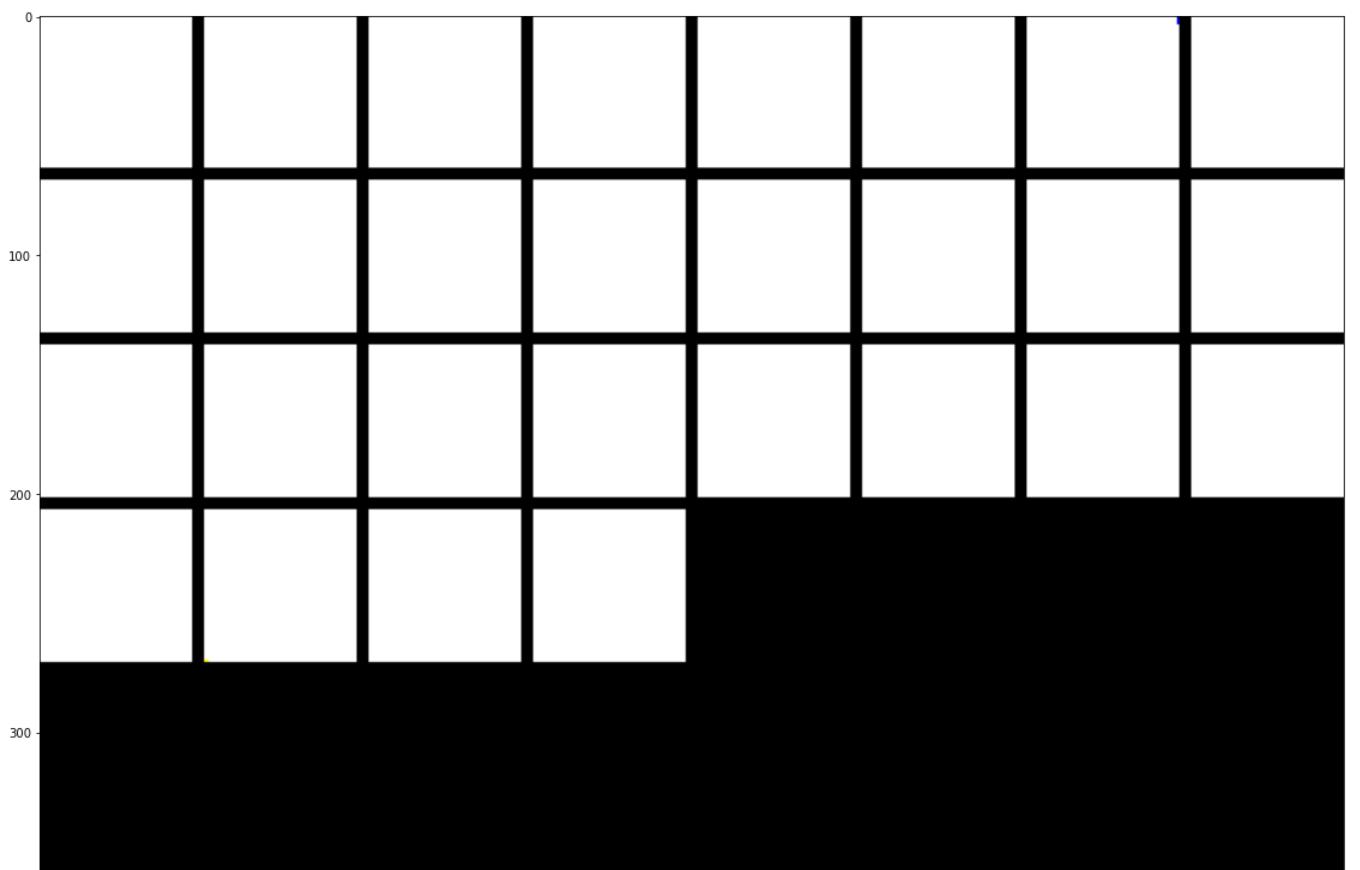
▼

▲





▲



▼
▲

